

# Does Fisher–Rao Birth–Death Help Adaptive Biasing Force?

## From Model Potentials and a Condensed-Phase Dimer to Molecular Torsions

August 16, 2026

### Abstract

We study whether augmenting the Adaptive Biasing Force (ABF) method with a Fisher–Rao (FR) birth–death correction improves the computation of a reaction-coordinate free energy, and what kind of improvement it can reasonably be expected to provide. We test this on five core case studies (then extended along the collective variable itself) spanning model potentials, a condensed-phase dimer, and molecular torsions: a two-dimensional double-well *metastability* model with a quadrature reference; a two-dimensional *entropic-bottleneck* model with a narrow transverse channel, for which both the reference free energy and the conditional law  $Y \mid X = x \sim \mathcal{N}(0, 1/(\beta\omega(x)^2))$  are analytic; a many-body condensed-phase *Weeks–Chandler–Andersen (WCA) dimer* in a dense solvent, evaluated against a high-accuracy thermodynamic-integration reference; and two united-atom alkanes whose slow coordinate is a signed dihedral — *butane* (one torsion, an easy control with an exact analytic reference) and *pentane* (two torsions coupled by the 1–5 “pentane effect”, so that the second dihedral is a hidden slow coordinate when only the first is biased). The alkanes use the correct nonlinear generalized mean force (including the geometric/curvature term) and periodic estimators. In each case the FR step reallocates walkers along the reaction coordinate by killing replicas in over-represented regions and cloning whole configurations from under-represented regions, nudging the empirical marginal  $\hat{p}_t$  toward a target  $q_t$ .

The central finding is not a single universal success story, but a regime map. In the metastability model, where plain ABF already samples reasonably well, estimated-target FR is a modest accelerator and the oracle diagnostic shows that target direction can matter. In the cold entropic bottleneck, ABF is starved and estimated-target FR gives a large gain (+51.3%, 20/20 seeds), but the uniform target nearly matches it and the oracle target is worse than ABF; here density correction and free-energy accuracy are not the same objective. A dedicated entropy-dominant stress test (transverse dimension  $m$ , total barrier held fixed) refines this further: as the barrier is made more entropic the *achievable* (oracle) gain grows specifically with the entropic share, yet the *deployable* estimated target does not capture it — an entropic-specific benefit that is real but target-limited. In the WCA dimer, tuned FR again improves the free energy reliably (+55.0%, 10/10 seeds), while the best accuracy occurs near the edge of an accuracy–diversity tradeoff. The molecular cases sharpen the picture: in *butane*, ABF already resolves the only slow coordinate, so mFR turns itself off (its birth–death rate falls to  $\approx 0.000$  and its free-energy error matches ABF, **equivalent** within a pre-registered  $\pm 10\%$  practical-equivalence margin) — the extra mechanism has no return where it is not needed, rather than failing. In *pentane*, where a second coupled dihedral makes the conditional law  $p(\varphi_2|\varphi_1)$  the binding difficulty, we separate genuine improvement (better  $F(\varphi_1)$  with faithful conditional) from false improvement (flatter marginal but corrupted conditional), and report the matched-seed outcome the data

support rather than assuming the coupled case must be positive. A final study varies the *collective variable itself* to stress-test the mechanism: replacing the easy torsion with a harder chemically meaningful scalar — the pentane end-to-end distance  $R_{15}$  — makes ABF genuinely sample-starved *along the biased coordinate* at low temperature (the first starved coordinate in the study), yet even here mFR gives no improvement (equivalent at the deployable rate, monotonically worse and lineage-collapsing as it is pushed), because that starvation is a *dynamical* bottleneck — rare torsional-barrier crossings — that marginal reallocation can only amplify, not create; and promoting the hidden torsion into a genuine two-dimensional CV ( $\varphi_1, \varphi_2$ ), with an exact vector generalized mean force and a periodic FFT Poisson projection, removes the starvation entirely (ABF resolves both barriers directly), so mFR is again neutral, now for a well-understood reason.

Two all-atom peptides extend the negative to real molecules and one constructed model supplies the missing positive control. Vacuum *alanine* and *valine* dipeptides were both selected as systems where mFR ought to help and both return **equivalent**: valine’s slow coordinate turns out to be the backbone  $\varphi$  rather than the side-chain rotamer it was chosen for — its  $11\text{--}18 k_B T$   $\chi_1$  barrier is a backbone-*clamped* conditional barrier, worth  $1.1\text{--}7.4 k_B T$  once the backbone is free — and in both peptides ABF discovers *and* establishes every state on its own, so reallocation has no deficit to repair. Reading these together with the alkane results gives regimes rather than a spectrum: *ABF-sufficient* (found early, populated early; mFR neutral), *discovery-limited* (never found; mFR cannot act), and *establishment-limited* (found early, populated slowly), which is the only regime mFR serves and the one that never arose on its own — together with the *conditional-equilibration-limited* regime described below, which is checked first because it decides whether the other three verdicts mean anything. A two-state reduction explains the separation among the last three: discovery is a first-passage event among  $N$  walkers and takes  $\sim 1/(Nk_{01})$ , whereas establishment is a relaxation at rate  $k_{01} + k_{10}$  with no factor of  $N$ , so more walkers buy the first and not the second and the gap between them widens with  $N$ . We therefore build that regime deliberately: an entropic-gateway model with analytic free energy, screened by ABF alone over a 128-cell *finite-budget establishment map* whose classification and anchor cell were frozen and committed *before* any Fisher–Rao arm was run. On 32 fresh seeds at rates frozen in a separate calibration pass, mFR improves the integrated free-energy error by  $-12.48\%$   $[-16.30, -9.84]$  on 31/32 seeds, and beats *its own matched sham* — the same number of clone/delete events at the same times, direction randomised — directly by  $-14.95\%$   $[-16.22, -12.34]$ . That run is a replicate on a corrected batch structure and is the one the preregistration’s amendment nominates for quoting; the first pass gave  $-12.12\%$  and  $-11.79\%$  and is neither withdrawn nor overwritten. An estimator-independent frozen-bias endpoint, which discards the online accumulators entirely, agrees at  $-12.40\%$ .

The same control is then carried to the many-body system. At the accepted WCA-dimer configuration, on 16 fresh seeds with no hyperparameter reopened, practical mFR improves the integrated error by  $-17.97\%$   $[-22.73, -17.16]$  on 16/16 seeds and beats its own matched sham by  $-20.96\%$   $[-25.04, -11.54]$ , while the sham itself is *not resolvably* better than ABF ( $-3.19\%$ , interval spanning zero on 9/16 seeds). These are scored against a corrected high-precision reference; the cached reference used in the first pass was later found wrong by up to  $67\sigma$  inside the transition region, which moved this headline from  $-22.83\%$  to  $-17.97\%$  without changing any preregistered verdict. Paired on the same seeds, round trips move by  $+1.09\%$   $[+0.77, +1.50]$  while the error moves by  $-17.97\%$ , so in this setup the mechanism is consistent with redistributing population among replicas already crossing rather than with buying barrier crossings — which is what a two-state reduction of the kinetics predicts. These are the first controls in the study that separate Fisher–Rao steering from resampling of the same intensity and timing, and they attribute the gain to the steering in both a designed model and a condensed-phase molecular system.

Attribution to *directed* selection is not the same as attribution to *this* direction, and the final experiment separates them. Run against the two directed rules already in use — the Chapter 6 Laplacian rule and simple count balancing — at matched turnover on 16 paired seeds with nothing retuned, mFR is **tied** by count balancing ( $+0.66\%$ , TOST interval  $[-3.96, +2.55]$  inside

a  $\pm 5\%$  margin), which recovers the entire gain on its own ( $-18.23\%$  against ABF), while the Laplacian rule is inert ( $+0.73\%$ ). The establishment-limited gain is therefore real, replicated and attributable to direction rather than turnover, but the Fisher–Rao geometry is not what delivers it, and this report claims a principled rule consistent with the mechanism rather than a win over prior art.

A structural result explains why a tie was the likely outcome and reorganises the map. For *any* score depending on the reaction coordinate alone — which includes all three rules above — the mean-field selection flow satisfies  $\partial_t p_t(y \mid \xi) = 0$ : it rescales the marginal and leaves every conditional exactly invariant. Selection therefore cannot repair a conditional that has not equilibrated, which is a fourth regime rather than a tuning failure, and the map accordingly leads with a validity gate on the ABF baseline itself. Deca-alanine is the cautionary instance: its discovery and establishment gates both pass, yet independently prepared orthogonal ensembles disagree on the local mean force by 0.61 of  $|F'|$  (against 0.042 for WCA and 0.036 for the gateway), so the classification that would have made it a target for the method is inadmissible.

Across all cases, FR is best understood as a marginal-reallocation mechanism whose value depends on how starved the ABF estimator is *along the biased coordinate*, whether that starvation is a support deficit reallocation can fill or a dynamical bottleneck it cannot, whether the conditional ensemble has equilibrated at all, how quickly cloned orthogonal degrees of freedom re-equilibrate, and how aggressively birth–death erodes ancestor diversity. The oracle target is therefore a diagnostic control, not a deployable method and not a universal upper bound.

## 1 Executive synthesis

This section states the v1 result in full before any of the machinery. Everything asserted here is traceable to Tables 31 and 32 and, through them, to the artifact-derived inventory `results/closure/v1_results_inventory.csv`.

**What marginal Fisher–Rao changes.** ABF learns a bias from a mean force estimated on whatever configurations its walkers happen to occupy. The marginal Fisher–Rao (mFR) step adds one thing to that: at fixed population size it kills replicas where the empirical reaction-coordinate marginal  $\hat{p}_t$  exceeds a target  $q_t$  and clones replicas where it falls short. It moves *existing* walkers between regions of the biased coordinate. It does not change the potential, the estimator, or the dynamics between events, and it cannot put a walker anywhere that no walker has already reached.

**Four finite-time regimes.** Which of those properties binds is decided by three timescales measured against the run length: how long a relevant state takes to be **found**, how long its orthogonal degrees of freedom take to **equilibrate** at fixed reaction coordinate, and how long it then takes to be adequately **populated**. Every benchmark in this report falls into one of the resulting regimes (Table 31), classified by the *first* gate it fails.

- **Conditional-equilibration-limited** — the orthogonal ensemble has not equilibrated, so the ABF mean force is an average over a non-equilibrium conditional and the baseline itself is invalid. Deca-alanine and the  $\beta = 2$  cells of  $R_{15}$ . By Proposition 2 no marginal selection rule can help; more importantly, no discovery or establishment verdict computed on such a baseline is admissible. Deca-alanine passes both of those gates and would have been filed as a target for the method, which is why this check now runs first.
- **ABF-sufficient** — found early, populated early. Butane and pentane on  $\varphi_1$ , pentane on the torsion torus  $(\varphi_1, \varphi_2)$ , both dipeptides, the metastability model, and the entropic bottleneck at  $\beta \leq 4$ . There is no marginal imbalance left for reallocation to correct.

- **Discovery-limited** — not found. Pentane’s end-to-end distance  $R_{15}$  at  $\beta = 2$ , the only genuinely starved biased coordinate in the study. Selection amplifies a configuration that exists; it cannot manufacture a torsional crossing.
- **Establishment-limited** — found early, populated slowly. The constructed entropic gateway, the condensed-phase WCA dimer, and the entropic bottleneck at  $\beta = 8$ . This is the regime mFR serves.

|                                                                                                                                                                                                                                               |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p>ABF-sufficient <math>\implies</math> mFR is unnecessary or neutral,<br/> discovery-limited <math>\implies</math> mFR cannot act before discovery,<br/> establishment-limited <math>\implies</math> directed mFR reallocation can help.</p> |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

Figure 1: The v1 conclusion. The implications run one way: the regime is diagnosed from discovery and establishment evidence that does not mention mFR, and the mFR outcome is then a prediction rather than the definition. Table 31 states the evidence for each benchmark separately from the outcome for exactly this reason.

**The two positives, and the control that earns them.** Both establishment-limited systems were tested against a *matched-turnover sham*: the same number of kill-and-clone events, at the same times, with the direction randomised. This is the control that separates Fisher–Rao steering from resampling of the same intensity, and until it existed “mFR helps” and “any turnover of this magnitude helps” were indistinguishable.

On the constructed gateway, at a cell classified and frozen before any Fisher–Rao arm was written, practical mFR improves the time-integrated free-energy error by  $-12.48\%$  on 31/32 fresh seeds in the quoted replicate ( $-12.12\%$  on 31/32 in the first pass), and beats its own sham directly in both. On the WCA dimer, with no hyperparameter reopened and 16 fresh seeds, it improves the same endpoint by  $-17.97\%$  on 16/16 seeds and beats its own sham by  $-20.96\%$  on 15/16. Neither sham is resolvably better than ABF: the gateway’s lands on the adverse side of zero and WCA’s spans it ( $-3.19\%$  on 9/16 seeds). (The WCA figures are scored against a corrected high-precision reference; the cached one used first was wrong inside the transition region and had put this headline at  $-22.83\%$ .)

That gain is attributable to *directed* selection, but not to *this* direction. At matched turnover against the two directed rules already in use, mFR is tied by simple count balancing ( $+0.66\%$ , TOST interval  $[-3.96, +2.55]$  inside  $\pm 5\%$ ), which reaches  $-18.23\%$  against ABF on its own, while the Chapter 6 Laplacian rule is inert ( $+0.73\%$ ). The claim this report can support is a principled rule consistent with the establishment mechanism, not a performance win over prior art.

Barrier round trips in the WCA test differ by  $+1.09\%$  between mFR and ABF while the free-energy error differs by  $-17.97\%$  — roughly a sixteen-fold separation between the transport change and the accuracy change, both measured on the same runs against the same reference. In this setup the mechanism is therefore consistent with redistributing population among replicas that were already crossing, rather than with buying additional crossings.

**The boundary cases, which are not failed benchmarks.** Four molecular systems and two model potentials return neutral or adverse results, and they are load-bearing rather than embarrassing: they are what makes the positive claim narrow enough to be useful. Butane, pentane on

$\varphi_1$ , the torsion torus, and both dipeptides establish that when ABF already resolves the slow coordinate along the biased CV, directed reallocation has no work left to do — in the alanine dipeptide even the *oracle* target, built from the reference free energy itself, changes nothing across a  $22\times$  range of realised event rate. Pentane’s  $R_{15}$  establishes the harder point: making the biased coordinate genuinely starved is *not* sufficient, because that starvation is a dynamical bottleneck rather than a support deficit, and pushing mFR hard enough to repair the nominal support makes the free energy substantially worse while ancestor diversity collapses and round trips fall. Amplification is not discovery.

**Principal limitations.** The positive rests on two systems, one of them constructed for the purpose. The gateway’s sham *equivalence* test passes in the first confirmatory pass and misses the preregistered margin in the replicate, so equivalence is not established there and the claim rests on the direct arm-versus-sham contrast instead; on the WCA dimer that direct contrast was preregistered, and the sham fails equivalence in the adverse direction. No practical (estimated-target) mFR arm was ever run on alanine, and no mFR arm at all was run on valine, whose screen stopped the study before one was justified. The deployable estimated target remains the binding constraint wherever an oracle target shows headroom. And every conclusion is conditioned on the specific model, coordinate, budget, and implementation tested here.

Finally, a parallel audit line in this repository — forked from the same commit, closed separately, and *not* merged — reaches more conservative conclusions on the WCA system: it finds the cached thermodynamic-integration reference inflates the measured improvement by about  $2\times$  on a related endpoint, it rejects variance reduction as the mechanism, and it reports the Fisher–Rao score as no better than plain count balancing. All three have since been checked directly and all three hold: the reference was rebuilt and corrected, the variance-reduction language has been withdrawn in favour of coverage equalisation, and the count-balancing claim now has a preregistered turnover-matched test behind it. What remains unreconciled is only that the two lines use different endpoints, so their effect sizes are not directly comparable; Remark 19 states the current position.

## 2 Introduction

### 2.1 Free-energy computation along a reaction coordinate

A central task in computational statistical mechanics is to compute the free energy associated with a low-dimensional *reaction coordinate* of a high-dimensional system in thermal equilibrium [11]. For most of the development we take a configuration to be a point  $q = (x, y) \in \mathbb{R}^{21}$  drawn from the canonical (Boltzmann–Gibbs) density

$$\pi(x, y) \propto e^{-\beta V(x, y)}, \tag{1}$$

where  $V$  is the potential energy and  $\beta = 1/(k_B T)$  is the inverse temperature. Fixing the reaction coordinate to the first coordinate,  $\xi(x, y) = x$ , the free energy  $F(x)$  and the local mean force  $F'(x)$  along  $\xi$  are the objects of interest. The difficulty is metastability: when  $V$  has well-separated basins along  $x$ , unbiased Langevin dynamics crosses the barrier rarely, so a direct estimate of  $F$  converges intolerably slowly.

---

<sup>1</sup>We write a configuration as  $q = (x, y)$  in the two-dimensional cases; the many-body case of Section 7 uses  $q \in \mathbb{R}^{2n}$  with a scalar collective variable  $z(q)$  in place of  $x$  (Remark 1). The symbol  $q_t(\cdot)$ , always carrying a time subscript and a scalar argument, denotes instead the Fisher–Rao *target marginal* of Section 4; the two are never used in the same role.

## 2.2 Adaptive Biasing Force

The Adaptive Biasing Force (ABF) method [5, 6, 8] removes this metastability adaptively. It maintains a running estimate  $\hat{F}'_t(x)$  of the mean force and applies the antiderivative  $B_t(x) = \int^x \hat{F}'_t$  as a bias that depends only on  $x$ . As  $B_t \rightarrow F$ , the effective dynamics along  $x$  becomes diffusive and the reaction coordinate is explored freely, while — crucially — the bias does not change the conditional law of the orthogonal degrees of freedom given  $x$ . ABF is mathematically well understood, with long-time convergence guarantees [10], and is a natural baseline against which to measure any proposed enhancement.

## 2.3 Why add a Fisher–Rao correction?

ABF accelerates exploration through the *force*, but it does not directly control *where the particles are*. Early in a simulation the estimated bias is poor, so the empirical reaction-coordinate marginal  $\hat{p}_t(x)$  can remain concentrated in the initial basin, starving other values of  $x$  of the samples ABF needs to estimate  $F'(x)$  there. A Fisher–Rao (FR) birth–death correction [12] addresses exactly this: interpreting the evolution of  $\hat{p}_t$  as a gradient flow of relative entropy in the Fisher–Rao geometry [1], one reallocates probability mass by *killing* particles in over-represented regions and *cloning* particles in under-represented regions, driving  $\hat{p}_t$  toward a chosen target marginal  $q_t(x)$  without integrating any additional dynamics. Coupling such a correction to ABF is appealing because the two act on complementary objects (force versus density), but it is not automatically beneficial: a birth–death step resamples *whole* configurations  $(x, y)$ , and if it distorts the conditional law  $Y \mid X = x$  it will corrupt precisely the quantity ABF estimates.

## 2.4 Central question and contributions

This report asks how FR changes ABF across five core controlled systems — and, in a final study, along the collective variable itself — not whether one slogan can describe all of them:

1. **When does marginal reallocation help ABF, and when can it hurt?** (Q1)
2. **What target and rate choices make FR useful in each regime?** (Q2–Q3)
3. **What limitations and failure modes appear when the system becomes more realistic?** (Q4)
4. **On a genuinely molecular reaction coordinate (a signed dihedral), does marginal FR help, is it superfluous, or can it produce a “false improvement” that flattens the biased marginal while corrupting a hidden conditional law?** (Q5)
5. **Is the gain, where it exists, attributable to Fisher–Rao steering at all — or would resampling of the same intensity and schedule have done as much?** (Q6)

Q6 was added late and is the sharpest of the six, because until it was asked every arm in this study shared the same replacement direction and differed only in how the target was built. Answering it requires a regime where mFR demonstrably helps *and* a control that strips the direction while holding the intensity fixed; Section 15 constructs both. The systems are (i) a two-dimensional double-well *metastability* model with a quadrature reference (Section 5); (ii) a two-dimensional *entropic-bottleneck* model with a narrow transverse channel, for which both the reference free energy *and* the conditional law  $Y \mid X = x$  are analytic (Section 6); (iii) a many-body *Weeks–Chandler–Andersen*

(WCA) *dimer* in a dense solvent with a high-accuracy thermodynamic-integration reference (Section 7); and two united-atom alkanes whose slow coordinate is a signed dihedral (Section 8) — (iv) *butane*, an easy control with an exact analytic torsional reference (Section 9), and (v) *pentane*, whose two coupled dihedrals make the second a hidden slow coordinate when only the first is biased (Section 10). Together they span an easy baseline where target direction can matter, a cold analytic bottleneck where density correction and free-energy accuracy can come apart, a many-body system where improved coverage must be balanced against lineage diversity, and two molecular torsions that test the mechanism where the collective variable is a nonlinear function of a three-dimensional configuration.

Two later additions complete the picture. Two all-atom dipeptides — vacuum *alanine* and *valine* (Section 14) — were selected as systems where mFR ought to help and instead serve as *neutrality controls*, because ABF discovers and populates every metastable state on its own. Their failure is systematic rather than accidental: it is what happens whenever the slow coordinate lies inside the collective variable, which is a property of the coordinate rather than of the method, and which no further choice of molecule would fix. The remaining regime — a state found early but populated slowly — is therefore constructed deliberately as an *entropic gateway* (Section 15), screened with ABF alone and with its anchor cell frozen before any Fisher–Rao arm ran.

Our contributions are: (i) a consistent ABF–FR notation that covers the analytic two-dimensional models, the many-body WCA collective variable, and the periodic molecular dihedral with its correct nonlinear generalized mean force (Sections 3, 4 and 8.2); (ii) matched-seed production studies that quantify both the practical estimated-target method and diagnostic uniform/oracle targets, including a pre-registered practical-equivalence test for the easy butane control and configuration-resolved hidden-conditional diagnostics for pentane (Sections 5 to 7, 9 and 10); (iii) a pre-registered screening protocol that decides, *before* any Fisher–Rao arm is run, whether a system can test the method at all — separating states that are never found from states that are found and stay under-populated, against a *bias-aware* target that does not mistake correct flattening for starvation (Section 14.1); (iv) a matched sham-resampling control that holds the replacement schedule and count fixed while randomising the direction, which is what makes the positive result attributable (Section 15.4); and (v) a cross-case synthesis (Section 17) that resolves the examples into three regimes and explains why the conclusions are intentionally different across them. A recurring methodological point is the careful interpretation of an *oracle* target that uses the reference free energy: it is a diagnostic control, never a deployable method and not a universal upper bound. Its behaviour is itself part of the evidence: best in the easy metastability setting, approximately tied with the practical target in WCA, and worse than ABF in the cold entropic bottleneck.

### 3 Background: ABF for a reaction-coordinate free energy

#### 3.1 Free energy and the unbiased marginal

With reaction coordinate  $\xi(x, y) = x$ , the unbiased reaction-coordinate marginal of the canonical density (1) is

$$p_{\text{ref}}(x) = \frac{\int e^{-\beta V(x, y)} dy}{\iint e^{-\beta V(x, y)} dx dy}, \quad (2)$$

and the free energy along  $\xi$  is, up to an additive constant,

$$F_{\text{ref}}(x) = -\frac{1}{\beta} \log p_{\text{ref}}(x) + C. \quad (3)$$

Because  $F$  is defined only up to the constant  $C$ , every free-energy profile in this report — both estimates and the reference — is *centred* (its mean over the interior evaluation window is subtracted) before any  $L^2$  error is formed; otherwise a meaningless constant offset would dominate the comparison. The reference quantities are computed by a one-dimensional  $y$ -quadrature on a fine grid with a log-sum-exp shift for numerical stability, which is exact up to discretisation and provides the ground truth  $F_{\text{ref}}, F'_{\text{ref}}, p_{\text{ref}}$ .

### 3.2 The mean force is the object ABF estimates

Differentiating (3) and using  $\partial_x p_{\text{ref}}$  gives the *mean force*

$$F'_{\text{ref}}(x) = \frac{\int \partial_x V(x, y) e^{-\beta V(x, y)} dy}{\int e^{-\beta V(x, y)} dy} = \mathbb{E}_{\pi}[\partial_x V(X, Y) | X = x]. \quad (4)$$

Equation (4) is the cornerstone of ABF: the gradient of the free energy equals the *conditional expectation of the instantaneous force*  $\partial_x V$  given the reaction coordinate. It is a local, sample-based quantity — no integral of  $V$  is required — so it can be estimated on the fly from the particles currently at  $X \approx x$ , and the free energy is then recovered by quadrature,  $\hat{F}(x) = \int^x \hat{F}'$ .

### 3.3 Biased dynamics

ABF simulates  $N$  overdamped Langevin replicas that share one running mean-force estimate  $\hat{F}'_t(x) \approx F'_{\text{ref}}(x)$  with bias potential  $B_t(x) = \int^x \hat{F}'_t$ . Each replica  $(X_t, Y_t)$  evolves as

$$\begin{aligned} dX_t &= [-\partial_x V(X_t, Y_t) + \hat{F}'_t(X_t)] dt + \sqrt{2/\beta} dW_t^x, \\ dY_t &= -\partial_y V(X_t, Y_t) dt + \sqrt{2/\beta} dW_t^y, \end{aligned} \quad (5)$$

with independent Brownian motions  $W^x, W^y$ . The added force  $+\hat{F}'_t(X_t)$  acts *only* in the  $x$ -direction; the  $y$ -dynamics is unbiased. When  $\hat{F}'_t \rightarrow F'_{\text{ref}}$  the net drift along  $x$  becomes  $-\partial_x(V - F)$ , whose  $x$ -marginal equilibrium is flat, so the reaction coordinate diffuses freely across the barrier.

### 3.4 Why the bias does not spoil the mean-force estimate

A bias that depends only on  $x$  leaves the conditional law of  $Y$  given  $X = x$  invariant. This is what makes ABF self-consistent.

**Proposition 1** (Conditional invariance under an  $x$ -only bias). *Let  $\nu$  be the canonical measure  $\nu \propto e^{-\beta V(x, y)}$  and let  $\nu_B \propto e^{-\beta [V(x, y) - B(x)]}$  be the measure biased by any function  $B$  of  $x$  alone. Then for every  $x$  the conditionals coincide,  $\nu_B(y | X = x) = \nu(y | X = x)$ , and consequently*

$$\mathbb{E}_{\nu_B}[\partial_x V(X, Y) | X = x] = \mathbb{E}_{\nu}[\partial_x V(X, Y) | X = x] = F'_{\text{ref}}(x). \quad (6)$$



*Proof.* At fixed  $X = x$  the bias factor  $e^{\beta B(x)}$  is a constant in  $y$  and cancels between numerator and normaliser of the conditional density:

$$\nu_B(y \mid X = x) = \frac{e^{-\beta[V(x,y)-B(x)]}}{\int e^{-\beta[V(x,y')-B(x)]} dy'} = \frac{e^{-\beta V(x,y)}}{\int e^{-\beta V(x,y')} dy'} = \nu(y \mid X = x).$$

Taking the conditional expectation of  $\partial_x V$  under this common conditional gives (6), the last equality being (4).  $\square$

Proposition 1 is the reason ABF works: even though the biased ensemble (5) has a different  $x$ -marginal than the canonical one, the conditional samples  $Y \mid X = x$  remain canonical, so the empirical conditional average of  $\partial_x V$  is an unbiased target for  $F'_{\text{ref}}(x)$ . Any modification of ABF must respect this property; in particular, a birth–death correction that resamples whole configurations  $(x, y)$  is only admissible if it too leaves  $Y \mid X = x$  (approximately) intact — a hypothesis we test directly in each of the three case studies (Sections 5 to 7).

*Remark 1* (Generalisation to many-body systems). Two of our three systems are two-dimensional with  $\xi = x$ , but the WCA dimer of Section 7 is a many-body system: a configuration is the full vector  $q \in \mathbb{R}^{2n}$  of all particle coordinates, and the reaction coordinate is a scalar *collective variable*  $z(q)$  (a scaled bond length) rather than a single Cartesian component. Proposition 1 holds verbatim under the substitution  $x \mapsto z(q)$  and  $y \mapsto$  the orthogonal complement of  $z$  on the constant- $z$  manifold (here the solvent coordinates and the transverse dimer modes): a bias  $B(z)$  that depends on the collective variable alone is constant on each level set of  $z$ , so it cancels in the conditional density and leaves the law of all orthogonal degrees of freedom invariant. Consequently the mean force along  $z$  — which for a nonlinear  $z(q)$  also carries a geometric term (Section 7) — is still estimated without bias. The only requirement the FR step must meet is the same one as in two dimensions: a clone must copy the *entire* configuration  $q \in \mathbb{R}^{2n}$ , so that the orthogonal coordinates travel with  $z$  and the conditional law is preserved.

### 3.5 The ABF mean-force estimator

In the kernel form used by the reference implementation,  $\hat{F}'$  is a Nadaraya–Watson running average of the instantaneous force,

$$\hat{F}'_n(x) = \frac{\sum_{k \leq n} \sum_i K_h(x - X_k^i) \partial_x V(X_k^i, Y_k^i)}{\sum_{k \leq n} \sum_i K_h(x - X_k^i)}, \quad (7)$$

with a Gaussian kernel  $K_h$  of bandwidth  $h$ , accumulated over all replicas  $i$  and past steps  $k$ ; the free-energy estimate is its pinned trapezoidal antiderivative,  $\hat{F}_n(x) = \int_0^x \hat{F}'_n$ , and we use  $B_t \equiv \hat{F}$  as the bias. The GPU backend used for the production study replaces the  $O(n_{\text{grid}} N)$  kernel sums by an  $O(N + n_{\text{grid}})$  *binned-smoothed* estimator: particles are scattered onto the  $x$ -grid and the count and force histograms are Gaussian-smoothed, so that  $\hat{F}'(x) = \text{smooth}(\text{force})/\text{smooth}(\text{count})$  is a Riemann-sum approximation of (7) that converges to it as the grid spacing vanishes. The two estimators are validated to agree to well under 1% on static tests (Section 5.1).

## 4 The Fisher–Rao marginal correction

The Fisher–Rao (FR) correction is a birth–death / particle-reallocation step that acts on the *reaction-coordinate marginal* of the ensemble. It is interleaved with the ABF dynamics (5) and is designed to drive the empirical marginal  $\hat{p}_t(x)$  toward a target marginal  $q_t(x)$ , without adding any drift to the  $(x, y)$  dynamics.

## 4.1 Empirical and target marginals

At step  $n$ , after the Langevin proposal of all  $N$  replicas, the empirical reaction-coordinate marginal is estimated by a (reflected, edge-corrected) kernel density / binned-smoothed estimate of bandwidth  $\eta$ ,

$$\hat{p}_n(x) = \frac{1}{N} \sum_{i=1}^N K_\eta(x - X_n^i), \quad (8)$$

renormalised on the  $x$ -grid. The target  $q_n(x)$  is one of three choices (Section 4.7); all are nonnegative and normalised,  $\int q_n = 1$ .

## 4.2 Fisher–Rao score

For each particle the FR *score* is the pointwise log-ratio of the empirical to the target marginal, with its ensemble mean removed:

$$S_i = \log \frac{\hat{p}_n(X_n^i) + \varepsilon}{q_n(X_n^i) + \varepsilon} - \underbrace{\int \hat{p}_n(z) \log \frac{\hat{p}_n(z)}{q_n(z)} dz}_{= \text{KL}(\hat{p}_n \| q_n)}, \quad (9)$$

where  $\varepsilon > 0$  is a small regulariser and the integral baseline equals the Kullback–Leibler divergence  $\text{KL}(\hat{p}_n \| q_n)$  (a grid quadrature; it coincides with the Monte-Carlo baseline  $\frac{1}{N} \sum_k \log[\hat{p}_n(X^k)/q_n(X^k)]$  up to discretisation). Subtracting the baseline makes the score mean-zero, so the birth–death step below conserves the expected particle number. The score is clipped to  $[-S_{\max}, S_{\max}]$  for stability. Its interpretation is direct:

- $S_i > 0$ :  $X^i$  lies where  $\hat{p}_n > q_n$ , an *over-represented* region (too many particles);
- $S_i < 0$ :  $X^i$  lies where  $\hat{p}_n < q_n$ , an *under-represented* region (too few particles).

## 4.3 Birth–death resampling

The FR step approximates the Fisher–Rao gradient flow of  $\text{KL}(\hat{p} \| q)$  as a fixed- $N$  particle reallocation. Over a sub-step of length  $\Delta t$  and at effective rate  $\gamma_n^{\text{eff}}$  (Section 4.6), each *over-represented* particle ( $S_i > 0$ ) is marked for death with probability

$$\mathbb{P}(i \text{ dies}) = 1 - \exp(-\gamma_n^{\text{eff}} S_i \Delta t), \quad (10)$$

and particles with  $S_i < 0$  supply clone sources. The WCA and metastability backends implement this in the direct death-replacement form: killed slots are filled by clones sampled from the under-represented pool, with weights increasing with  $|S_i|$ . The entropic-bottleneck backend uses a vectorised but closely related fixed- $N$  survivor/clone-pool variant. This distinction matters for exact diagnostics, but not for the mechanism studied here: all variants move mass along the reaction coordinate by copying whole configurations from under-represented regions into over-represented ones.

Cloning copies the *whole* configuration  $(X^i, Y^i)$  — the score depends only on  $x$ , but the  $y$ -coordinate must travel with its  $x$ -coordinate, or Proposition 1 would be violated. The step is made deliberately gentle by three safeguards: (i) a burn-in and soft ramp on the rate; (ii) score clipping to  $[-S_{\max}, S_{\max}]$ ; and (iii) a hard cap  $f_{\max}$  on the fraction of particles replaced per event (`max_event_fraction`). The numerical values of  $S_{\max}$ ,  $f_{\max}$  and the FR stride `fr_every` are case-specific and are listed in each case’s design table.

#### 4.4 What marginal selection cannot do

Proposition 1 says that an  $x$ -only *bias* leaves the conditional law  $Y \mid X = x$  alone. The birth–death step is not a bias, so that proposition does not cover it, and the natural worry is that resampling perturbs the conditional. The following says something sharper, and in the opposite direction: in the mean-field limit the selection step does not perturb the conditional at all, and therefore cannot repair it either.

**Proposition 2** (The selection flow leaves every conditional invariant). *Let  $p_t(x, y)$  evolve under the mean-field birth–death flow generated by a score  $S$  that depends on the reaction coordinate alone,*

$$\partial_t p_t(x, y) = -\gamma [S(x) - \langle S \rangle_{p_t}] p_t(x, y), \quad \langle S \rangle_{p_t} = \iint S(x) p_t(x, y) dx dy, \quad (11)$$

*and write  $p_t(x) = \int p_t(x, y) dy$  for the marginal and  $p_t(y \mid x) = p_t(x, y)/p_t(x)$  for the conditional, at any  $x$  with  $p_t(x) > 0$ . Then*

$$\partial_t p_t(x) = -\gamma [S(x) - \langle S \rangle_{p_t}] p_t(x), \quad \partial_t p_t(y \mid x) = 0. \quad (12)$$

*Proof.* Integrating (11) in  $y$  and using that  $S$  does not depend on  $y$  gives the first identity. For the second, differentiate the quotient and substitute both:

$$\partial_t \frac{p_t(x, y)}{p_t(x)} = \frac{\partial_t p_t(x, y)}{p_t(x)} - \frac{p_t(x, y) \partial_t p_t(x)}{p_t(x)^2} = \frac{-\gamma [S(x) - \langle S \rangle] p_t(x, y)}{p_t(x)} + \frac{\gamma [S(x) - \langle S \rangle] p_t(x, y)}{p_t(x)} = 0. \quad \square$$

The factor  $S(x) - \langle S \rangle$  multiplies  $p_t(x, y)$  and  $p_t(x)$  by the *same* number, so it cancels out of their ratio. Selection rescales the marginal and does nothing else.

Three consequences organise the rest of the report.

*First, the mechanism is bounded a priori.* Whatever a birth–death step buys, it is bought by reweighting the marginal. It cannot be bought by improving the conditional average that (6) says the mean-force estimate is built from. This is a stronger and more useful statement than “resampling is benign”: benign would permit a conditional improvement, and Proposition 2 forbids one.

*Second, it is not a property of the Fisher–Rao score.* The proof uses only that  $S$  is a function of  $x$ . It therefore applies verbatim to the historical selection rules of Section 7.16 — the Chapter 6 rule  $S \propto \partial_x^2 p/p$  and the count-balancing rule  $S \propto 1 - p/\bar{p}$  — which sit in exactly the same class. No score built on the reaction coordinate alone can escape (12), so no such rule can be distinguished from any other by its effect on conditionals. They can differ only in how they redistribute the marginal, and Section 7.16 measures how much that distinction is worth.

*Third, it predicts a regime in which the method must fail.* If the binding difficulty is that  $p_t(y \mid x)$  has not equilibrated — the transverse relaxation time  $\tau_\perp$  is long compared with the run — then no marginal selection rule can help, because none of them moves  $p_t(y \mid x)$  at all. That is not a tuning failure to be fixed with a better rate or a better target; it is structural. Section 17 makes it the fourth regime of the map, and it is the reason the map now leads with a validity gate on the ABF baseline rather than with a discovery test.

*Remark 2* (What the finite- $N$  ensemble does instead). Proposition 2 is a mean-field statement, and the gap at finite  $N$  runs the unfavourable way. Cloning copies a whole configuration  $(x, y)$ , so a clone and its parent share a  $y$ -coordinate: the conditional is preserved *in expectation* while the number of *independent*  $y$ -samples behind each  $x$ -bin falls. The empirical conditional therefore does not stay put — it acquires resampling noise — and the two safeguards that matter are the ones

that limit how fast lineages coalesce, namely the event-fraction cap  $f_{\max}$  and the requirement that clones decorrelate in time  $\tau_{\perp}$  before the next event. This is why the report tracks ancestor diversity alongside the free-energy error, and why pushing the rate up eventually converts a marginal gain into a conditional loss.

#### 4.5 Notation across cases

The same FR principle is applied to all three systems, but the implementation and two parameter names differ slightly across backends. The birth–death *rate* is written  $\gamma$  in the analytic two-dimensional cases (Sections 5 and 6) and `fr_rate` in WCA (Section 7); they play the same role in (10). The *onset* of FR is parameterised as a burn-in *fraction*  $b$  in the analytic cases ((13),  $n_{\text{burnin}} = \lfloor b n_{\text{steps}} \rfloor$ ) and as an absolute step count `fr_start_steps` in WCA; the two are related by  $b \approx \text{fr\_start\_steps}/n_{\text{steps}}$ . Where a case uses a collective variable  $z(q)$  rather than  $x$ , read  $x \mapsto z$  throughout this section (Remark 1).

#### 4.6 Rate schedule

The FR rate is zero during an initial burn-in of  $n_{\text{burnin}} = \lfloor b n_{\text{steps}} \rfloor$  steps (burn-in fraction  $b$ ), then ramps softly to its maximum  $\gamma$ ,

$$\gamma_n^{\text{eff}} = \gamma \left( 1 - e^{-(n - n_{\text{burnin}})/n_{\text{ramp}}} \right), \quad n > n_{\text{burnin}}, \quad (13)$$

and the birth–death step is applied every `fr_every` steps. The burn-in lets ABF build a usable mean-force estimate (hence a usable target) before any particles are moved; the ramp avoids a sudden shock to the ensemble.

#### 4.7 Target marginals: estimated, uniform, oracle

The three targets differ only in which density  $q_n$  the score (9) is taken against. Recall  $B_t = \hat{F}_n$  is the current ABF bias.

**Estimated target (the practical method).** A target free energy  $\hat{F}_n^{\text{target}}$  is maintained independently of the bias by an exponential moving average of the ABF estimate,  $\hat{F}_{n+1}^{\text{target}} = (1 - r) \hat{F}_n^{\text{target}} + r \hat{F}_{n+1}$ , and the target marginal is the *biased* marginal that this target free energy would induce under the current bias,

$$q_n^{\text{est}}(x) \propto \exp \left[ -\beta (\hat{F}_n^{\text{target}}(x) - B_t(x)) \right]. \quad (14)$$

This uses only quantities available at run time and is the method whose practical value we assess.

**Uniform target (a naive ablation).**

$$q_n^{\text{unif}}(x) = \text{Unif}(x) \quad \text{on the } x\text{-grid}. \quad (15)$$

This tests whether simply *flattening* the reaction-coordinate marginal is enough. It is a target ablation, not the intended method; it ignores the fact that, under a bias  $B_t \neq F$ , the canonical biased marginal is not uniform.

**Oracle target (a diagnostic control only).**

$$q_n^{\text{oracle}}(x) \propto \exp\left[-\beta(F_{\text{ref}}(x) - B_t(x))\right]. \quad (16)$$

*Remark 3* (The oracle target is not an algorithm). Equation (16) replaces the estimated target free energy by the *reference* free energy  $F_{\text{ref}}$  — exactly the unknown quantity the whole procedure is trying to compute. The oracle target therefore *cannot* be evaluated in any real free-energy calculation. We include it solely as an ideal-information diagnostic: it asks what the FR birth–death step would do if the target marginal were anchored to the reference. This is not a universal upper bound. If oracle FR improves on estimated FR in a given regime, the gap can be read as target-shape headroom; if it ties or loses, the diagnostic instead says that the limiting mechanism is not the online target estimate, and may even be a mismatch between a fixed reference-shaped target and the transient ABF bias. The practical comparison is always

$$\text{ABF only} \quad \text{versus} \quad \text{ABF} + \text{estimated-target FR}.$$

## 4.8 The coupled ABF–FR step

One outer iteration, given the state  $\{(X_n^i, Y_n^i)\}$  and accumulators, is:

1. update the ABF accumulators from the current particles and recompute  $\hat{F}'_n, \hat{F}_n = B_t$  (Section 3.5);
2. propose the biased Langevin move (5) for all replicas;
3. if FR is active at this step (past burn-in, on the `fr_every` schedule): form  $\hat{p}_n$  (8) and the target  $q_n$  — one of (14), (15), (16) — compute the scores (9), and apply the birth–death step (10);
4. record diagnostics at the evaluation cadence.

The ABF estimator sees the post-resampling ensemble at the next step, which is the route by which improved  $x$ -coverage feeds back into a better mean-force estimate. The remainder of the report quantifies whether, and how, this feedback is beneficial.

## 5 Case I: a two-dimensional metastability model

Our first system is a two-dimensional double well: the regime in which ABF is already a strong baseline, so it sets the lower end of the difficulty spectrum.

### 5.1 Model potential and reference

We use a two-dimensional model potential with reaction coordinate  $\xi(x, y) = x$ ,

$$\begin{aligned} V(x, y) = & 3e^{-x^2} \left( e^{-(y-1/3)^2} - e^{-(y-5/3)^2} \right) - 5e^{-y^2} \left( e^{-(x-1)^2} + e^{-(x+1)^2} \right) \\ & + 0.2x^4 + 0.2(y - 1/3)^4. \end{aligned} \quad (17)$$

As shown in Figure 2,  $V$  has two metastable basins near  $x = \pm 1$  separated by a barrier near  $x = 0$  (with an upper sub-basin near  $(0, 5/3)$ ), so unbiased dynamics is metastable in  $x$ . We fix  $\beta = 4$ . The reference  $F_{\text{ref}}, F'_{\text{ref}}, p_{\text{ref}}$  are obtained by the  $y$ -quadrature of (2)–(4) on a fine grid and are treated as ground truth.




Figure 2: Problem geometry. (a) The potential  $V(x, y)$  (17) with two basins near  $x = \pm 1$ . (b) The canonical density  $\pi \propto e^{-\beta V}$  at  $\beta = 4$ , concentrated in the two basins. (c) The reference free energy  $F_{\text{ref}}(x)$  (a double well in  $x$ ) and (d) the reference mean force  $F'_{\text{ref}}(x)$ , both by  $y$ -quadrature.

## 5.2 Methods, metrics, and study design

Four methods are compared, identical except for the FR target: **ABF only** (no birth–death; baseline); **ABF+FR estimated** (14) (the practical method); **ABF+FR uniform** (15) (naive marginal-flattening ablation); and **ABF+FR oracle** (16) (diagnostic control, Remark 3). All four share the same potential, integrator, ABF estimator, seeds and initial conditions, so differences are attributable to the FR step and its target alone (matched-seed comparison).

### 5.2.1 Metrics

All profile errors are interior, domain-RMS  $L^2$  distances on the window  $x \in [-2.5, 2.5]$  (the reflecting walls and the  $x^4$  confinement make the edges uninformative), with free-energy profiles centred as in Section 3.

**Primary (practical) metrics.**

$$\left\| \hat{F}_T - F_{\text{ref}} \right\|_{L^2}, \quad \left\| \hat{F}'_T - F'_{\text{ref}} \right\|_{L^2}, \quad \int_0^T \left\| \hat{F}_t - F_{\text{ref}} \right\|_{L^2} dt. \quad (18)$$

Table 1: Production study design (read from the result CSVs and configs/two\_dim\_xi\_x\_production\_gpu.yaml).

| Quantity                        | Value                                    |
|---------------------------------|------------------------------------------|
| Inverse temperature $\beta$     | 4.0                                      |
| Time step $\Delta t$            | 0.002                                    |
| Steps per run                   | 100,000                                  |
| Particles per run $N$           | 1000                                     |
| Seeds                           | 5 (0, 1, 2, 3, 4)                        |
| Methods                         | 4 (ABF; FR estimated / uniform / oracle) |
| FR rate $\gamma$                | 0.01, 0.02, 0.05, 0.1                    |
| FR bandwidth $\eta$             | 0.075, 0.1, 0.15                         |
| FR burn-in fraction             | 0, 0.2, 0.4                              |
| FR stride <code>fr_every</code> | 5                                        |
| Total configurations            | 109                                      |
| Total runs                      | 545                                      |

The first two are *final-budget accuracy* at the end of the run; the *integrated* error rewards *fast* convergence (anytime performance), since a method that reaches a good profile earlier accumulates less error over  $[0, T]$ . These two notions need not be optimised by the same hyperparameters, a point we treat carefully in Section 5.3.

**Diagnostic metrics.** Reaction-coordinate marginal flatness ( $\|\hat{p} - \text{Unif}\|$  and  $\|\hat{p} - q\|$ ); cumulative barrier crossings at  $x = 0$ ; region fractions (left basin / barrier / right basin); the FR event fraction (particles replaced per FR application) and the FR score dispersion  $\text{std}_i S_i$ ; conditional  $Y \mid X = x$  errors (Section 5.2.1); and ancestor/diversity counts. We stress that *marginal flatness is a diagnostic, not the objective*: the goal is free-energy/mean-force accuracy (18). A method can flatten  $\hat{p}$  and yet fail to improve  $\hat{F}$  if it does so by corrupting the conditional samples.

**Conditional  $Y \mid X$  diagnostics.** At a set of  $x$ -bin centres we compare the empirical conditional  $\hat{p}(y \mid x \in I_j)$  (the  $y$ -histogram of particles whose  $x$  falls in a narrow window) to the reference conditional  $p_{\text{ref}}(y \mid x) \propto e^{-\beta V(x,y)}$ , via an  $L^2$  distance and a KL divergence. By Proposition 1 a benign FR step should leave these close to the ABF-only values.

### 5.2.2 GPU implementation and study size

The production runs use a batched PyTorch backend that simulates many runs at once and reuses the reference metrics/diagnostics code verbatim. For speed it uses the binned-smoothed ABF/FR estimator of Section 3.5 rather than the exact kernel sums (7), and a fully vectorised fixed- $N$  birth-death; the two backends agree to within the stated tolerance ( $< 1\%$  in practice) on the validation suite. We flag in Section 18 that this binned estimator is a deliberate, but not mathematically identical, surrogate for the kernel notebook implementation. The FR methods sweep the rate  $\gamma$ , the marginal bandwidth  $\eta$ , and the burn-in fraction  $b$  at a fixed FR stride; the resulting study (Table 1) comprises 545 runs over 109 configurations and 5 seeds, at 1000 particles and 100,000 steps each, with *no* NaNs in the primary metrics. The selected configurations are listed in Table 2.

Table 2: Hyperparameters of the selected configurations (integrated- $L^2(F)$  rule).  $n_{\text{FR}}$  is the Fisher–Rao stride `fr_every`.

| Method       | target    | $\gamma$ | $\eta$ | burn-in | $n_{\text{FR}}$ | seeds |
|--------------|-----------|----------|--------|---------|-----------------|-------|
| ABF only     | none      | 0        | 0.075  | 0       | 1               | 5     |
| FR estimated | estimated | 0.1      | 0.075  | 0       | 5               | 5     |
| FR uniform   | uniform   | 0.05     | 0.15   | 0.2     | 5               | 5     |
| FR oracle*   | oracle    | 0.1      | 0.075  | 0       | 5               | 5     |

Table 3: Main comparison, configurations selected by smallest median *integrated*  $L^2(F)$  (convergence speed). Lower is better except  $P(\text{beats ABF})$ .

| Method       | $\int_0^T \ \hat{F}_t - F_{\text{ref}}\  dt$ | $\ \hat{F}_T - F_{\text{ref}}\ $ | $\ \hat{F}'_T - F'_{\text{ref}}\ $ | $P(\text{beats ABF})$ | FR event frac.       |
|--------------|----------------------------------------------|----------------------------------|------------------------------------|-----------------------|----------------------|
| ABF only     | 15.91                                        | 0.0137                           | 0.0718                             | –                     | 0                    |
| FR estimated | 14.08                                        | 0.0126                           | 0.0771                             | 0.60                  | $1.6 \times 10^{-5}$ |
| FR uniform   | 14.46                                        | 0.0125                           | 0.0769                             | 0.80                  | $5.3 \times 10^{-6}$ |
| FR oracle*   | 13.27                                        | 0.0113                           | 0.0602                             | 1.00                  | $1.7 \times 10^{-5}$ |

\* The oracle target uses the reference free energy  $F_{\text{ref}}$  (the unknown being computed); it is a diagnostic positive control, *not* a usable algorithm. Per-config medians over 5 seeds; one configuration per method, selected by the smallest median integrated  $L^2(F)$ .

### 5.3 Results

Table 3 reports the four methods, each at its best configuration under the *integrated- $L^2(F)$*  rule (convergence speed; see Section 5.3.2). Reading the practical comparison (ABF only versus ABF+FR estimated): the estimated-target FR reduces the median integrated free-energy error from 15.91 to 14.08, a 11.5% reduction, while the *final* free-energy error improves only slightly (from 0.0137 to 0.0126). The birth–death is extremely gentle: the median per-event replacement fraction is  $1.6 \times 10^{-5}$ , of order one particle in  $10^5$  per FR application. Figure 3 shows the same story over time: the FR curves detach below the ABF-only curve in the early and middle phase of the  $L^2(F)$  panel.

#### 5.3.1 Final profiles and target-type ablation

Figure 4 overlays the final mean-force, free-energy and reaction-coordinate-marginal profiles. All practical methods recover  $F'_{\text{ref}}$  and  $F_{\text{ref}}$  well; the visible differences between ABF only and ABF+FR estimated at the final budget are small, consistent with the modest final-error improvements. The  $x$ -marginal (panel c) is somewhat flatter under FR, as intended, but marginal flatness is a diagnostic, not the goal.

Figure 5 shows the distribution, over the 36 configurations of each target type, of the seed-median errors. For the *integrated* error all three FR target types have medians below the ABF-only baseline, but the spread is wide and some configurations are *worse* than ABF, so FR does not help unconditionally. For the *final*  $L^2(F)$  the estimated and uniform medians sit essentially on the ABF baseline, and on the final *mean-force* error they lie slightly *above* it — only the oracle control is consistently below. Flattening the marginal with an imperfect target buys early speed but not final mean-force accuracy.



Convergence of the selected (best integrated  $L^2(F)$ ) configuration: median and IQR over 5 seeds




Figure 3: Convergence of the selected (integrated- $L^2(F)$  rule) configuration for each method: median (line) and inter-quartile range (band) over 5 seeds. The free-energy and mean-force errors (top, log scale) fall fastest for the oracle diagnostic and faster for the estimated/uniform FR than for ABF only in the early-to-middle phase; barrier crossings (bottom left) are essentially indistinguishable, and the FR event fraction (bottom right) is of order  $10^{-5}$  throughout. The oracle curve is a diagnostic control (Remark 3).

### 5.3.2 Two selection rules give different “best” hyperparameters

Because integrated and final errors measure different things, they need not be optimised by the same configuration. Selecting by *integrated*  $L^2(F)$  (Table 3) favours fast convergence and selects an aggressive configuration for the estimated target ( $\gamma = 0.1$ , burn-in 0); selecting by *final*  $L^2(F)$  (Table 4) favours final-budget accuracy and selects a gentle, delayed configuration ( $\gamma = 0.01$ , burn-in 0.4), which improves the final free-energy error by 12.5% and beats ABF on 80.0% of matched seeds. This dependence on the selection rule is itself a finding.

The *uniform* target is competitive here (integrated reduction 9.1%): under a near-converged bias  $B_t \approx F$  the canonical biased marginal is already close to uniform, so “flatten  $\hat{p}$ ” and “match the biased marginal” nearly coincide. The *oracle* diagnostic (Remark 3) achieves the lowest errors of all (integrated reduction 16.6%), but it requires  $F_{\text{ref}}$  and is unusable; the correct reading is comparative: the estimated target captures 11.5% of the reduction against the oracle’s 16.6%, so a substantial part of the mechanism’s potential is left unrealised by the online target estimate. In this easy regime, then, the oracle–estimated gap localises the practical bottleneck in the *target estimation* — a reading we revisit, and qualify, in Section 17.




Figure 4: Final profiles (mean over seeds) for the selected configurations. (a) Mean force  $\hat{F}'(x)$  and (b) centred free energy  $\hat{F}(x)$  versus the reference (dashed); (c) reaction-coordinate marginal  $\hat{p}(x)$ , with the unbiased  $p_{\text{ref}}$  (dotted). Practical methods recover the reference well; FR mildly flattens  $\hat{p}$ .




Figure 5: Target-type ablation: each box is the distribution over the 36 configurations of a target type (each point a seed-median); the dashed line is the ABF-only median. FR lowers the *integrated* error (left) but the practical targets barely change the *final*  $L^2(F)$  (middle) and slightly raise the final mean-force error (right). Only the oracle diagnostic improves all three.

## 5.4 Mechanism

The mechanism is a feedback loop between coverage and the mean-force estimate. ABF needs samples *across* the reaction coordinate: by (4),  $F'_{\text{ref}}(x)$  is a conditional average that can only be estimated where particles visit  $x$ . Early in a run the bias is poor and the ensemble is unbalanced — Figure 6(b) shows the initial occupancy split ( $\sim 0.5$  in one basin) relaxing only gradually toward the balanced  $1/3$ – $1/3$ – $1/3$  split. FR compares  $\hat{p}_t$  with the target  $q_t$  and preferentially kills over-represented and clones under-represented particles (10), raising the useful coverage of  $x$  earlier than diffusion alone would. Provided the birth–death does not corrupt  $Y \mid X = x$  (Proposition 1), better coverage feeds back into a better conditional mean-force estimate.

A notable feature of Figure 6(c) is that FR does *not* produce many more barrier crossings than ABF; the benefit comes from *reallocating* existing replicas toward starved values of  $x$  rather than from manufacturing rare transitions. Consistent with this, the event fractions are of order  $10^{-5}$ : in the successful regimes FR is a *weak* correction.

Table 4: Supplement: configurations selected by smallest median  $\text{final } L^2(F)$  (final-budget accuracy).

| Method       | $\int_0^T \ \hat{F}_t - F_{\text{ref}}\  dt$ | $\ \hat{F}_T - F_{\text{ref}}\ $ | $\ \hat{F}'_T - F'_{\text{ref}}\ $ | $P(\text{beats ABF})$ | FR event frac.       |
|--------------|----------------------------------------------|----------------------------------|------------------------------------|-----------------------|----------------------|
| ABF only     | 15.91                                        | 0.0137                           | 0.0718                             | –                     | 0                    |
| FR estimated | 15.25                                        | 0.0120                           | 0.0673                             | 0.80                  | $1.1 \times 10^{-6}$ |
| FR uniform   | 14.97                                        | 0.0118                           | 0.0680                             | 0.60                  | $7.3 \times 10^{-6}$ |
| FR oracle*   | 13.27                                        | 0.0113                           | 0.0602                             | 1.00                  | $1.7 \times 10^{-5}$ |

\* Oracle is a diagnostic control (uses  $F_{\text{ref}}$ ), not a deployable method. Configurations here are selected by smallest median  $\text{final } L^2(F)$ ; cf. Table 3, which need not pick the same hyperparameters.

#### 5.4.1 FR does not visibly corrupt $Y \mid X = x$

Figure 7 compares the conditional fidelity of all methods at the final time. The FR variants are *not* worse than ABF only; near the barrier ( $x \in \{-0.5, 0, 0.5\}$ ), where ABF-only under-samples, they are in fact slightly *better* on both the  $L^2$  and KL conditional metrics — direct evidence that the FR mechanism improves  $x$ -coverage without corrupting the orthogonal sampling, the benign behaviour Proposition 1 requires.

*Remark 4* (Conditional diagnostics are not configuration-resolved). The metastability conditional-diagnostics output records only (method, target type, seed) per row, not the FR hyperparameters. It supports the *aggregate* statement that successful FR does not visibly corrupt  $Y \mid X = x$ , but cannot isolate a specific “good” versus “bad” configuration’s effect. The entropic-bottleneck case (Section 6) closes this gap with an *analytic* conditional law.

#### 5.4.2 What determines a good FR?

Figure 8 (estimated target) shows the integrated error decreasing as  $\gamma$  grows — a stronger correction converges faster, which is why the integrated rule selects the aggressive  $\gamma = 0.1$  — but the same setting does *not* minimise the final error (Section 5.3.2): there is an explicit speed-versus-accuracy trade-off in  $\gamma$ . The marginal bandwidth  $\eta$  is robust at intermediate values, and an early burn-in helps the integrated error here because the largest imbalance is at early times. Across the successful configurations the event fractions are tiny ( $\sim 10^{-5}$ ) and the cap  $f_{\text{max}}$  never binds, so particle diversity is preserved.

### 5.5 Comparison to OPES

As a correctness gate for the OPES implementation, we run well-tempered OPES [9] on this metastability toy alongside ABF and mFR, using the identical dynamics, mean-force reconstruction, and seeds. On this *easy* case — where plain ABF already converges and the three FR variants agree within seed scatter (Section 5.6) — all methods sit near the same floor: neither mFR nor OPES has headroom to exploit, and the comparison simply confirms that the OPES implementation recovers the correct free-energy shape in the well-sampled basin-and-barrier region  $|x| \lesssim 2$  containing both basins and the saddle. No quantitative ranking is claimed here; the decisive comparisons are in Sections 6.7 and 7.13, where ABF is starved and all methods diverge.

### 5.6 Case finding

On this easy metastability example the answer to Q1 is *yes, but modestly, and chiefly as a convergence accelerator*. Practical estimated-target FR improves the integrated free-energy error by




Figure 6: Mechanistic / early-time behaviour, estimated FR (best integrated-rule config) versus ABF only. (a) Early-time free-energy error (log scale): FR converges faster. (b) Region occupancy under FR relaxes from an unbalanced start toward the balanced 1/3 split (dotted). (c) Cumulative barrier crossings are nearly identical — the benefit here is redistribution, not a large increase in raw crossings. (d) FR score dispersion decays as  $\hat{p}_t \rightarrow q_t$ .

11.5% with a smaller final-budget gain; the improvement is real but not a rescue, and it is not unconditional. Because plain ABF already converges well here, the room for FR to help is small. The oracle control does best, which — read *within this regime* — points to target estimation as the practical bottleneck; we will see in Sections 6 and 7 that this reading does not survive into the starved regimes, where the oracle’s advantage vanishes or reverses.

## 6 Case II: an entropic bottleneck

The second system tightens a transverse channel at the barrier top and, by lowering the temperature, drives ABF into a genuinely sample-starved regime. It also has a rare luxury: the conditional law  $Y \mid X = x$  is *analytic*, so we can test Proposition 1 directly rather than against a histogram.

Conditional sampling fidelity at  $T$  (aggregated over configs and seeds; not config-resolved)




Figure 7: Conditional  $Y | X = x$  fidelity at  $T$ , by  $x$ -bin: (a)  $L^2$  distance and (b) KL divergence of  $\hat{p}(y | x)$  to  $p_{\text{ref}}(y | x)$ . FR does not inflate the conditional error; near the barrier it slightly reduces it. *Limitation:* the conditional CSV identifies runs only by (method, target, seed) and is aggregated over configurations (Remark 4).

## 6.1 Model potential and reference

The potential couples a symmetric double well in  $x$  to a transverse harmonic channel whose stiffness varies with  $x$ ,

$$V(x, y) = H(x^2 - 1)^2 + \frac{1}{2}\omega(x)^2 y^2, \quad \omega(x) = \omega_{\text{out}} + (\omega_{\text{in}} - \omega_{\text{out}})e^{-x^2/2s^2}, \quad (19)$$

with reaction coordinate  $\xi(x, y) = x$ . The stiffness  $\omega(x)$  is large near  $x = 0$  (a *narrow channel at the barrier top*) and relaxes to  $\omega_{\text{out}}$  in the wells. Because the  $y$ -integral is Gaussian, the reference is analytic up to a constant:

$$F_{\text{ref}}(x) = H(x^2 - 1)^2 + \beta^{-1} \log \omega(x) + C, \quad F'_{\text{ref}}(x) = 4Hx(x^2 - 1) + \beta^{-1} \omega'(x)/\omega(x), \quad (20)$$

and, crucially, the conditional law is the exact Gaussian

$$Y | X = x \sim \mathcal{N}(0, 1/(\beta\omega(x)^2)). \quad (21)$$

The base regime is  $\beta = 8$ ,  $H = 2.5$ ,  $\omega_{\text{out}} = 1$ ,  $\omega_{\text{in}} = 25$ ,  $s = 0.25$ , with  $N = 256$  walkers,  $dt = 10^{-3}$ , and 40 000 steps. The free energy is gauged at  $x = 0$  and all  $L^2$  errors are RMS over the interior window  $x \in [-1.5, 1.5]$ .

**Why this is an entropic bottleneck — and an honest caveat.** Near  $x = 0$  the accessible  $y$ -width shrinks like  $1/\omega(x)$ : at the base setting the transverse channel is  $\omega_{\text{in}}/\omega_{\text{out}} = 25$ , i.e. its standard deviation is  $25\times$  narrower than in the wells, contributing an *entropic* term  $\beta^{-1} \log \omega(x)$  to the free energy. We are explicit, however, that in this parameterisation the entropic contribution  $\beta^{-1} \log(\omega_{\text{in}}/\omega_{\text{out}}) \approx 0.40$  is smaller than the energetic barrier height  $U(0) - U(\pm 1) = H = 2.5$ . The quantity  $4H = 10$  is the quartic force prefactor in  $F'(x)$ , not the free-energy barrier. The bottleneck therefore modulates the transverse width but is *not* the dominant contributor to reaction-coordinate starvation; that role belongs to the energetic double well. This matters for reading the  $\omega_{\text{in}}$  sweep below.




Figure 8: Estimated-target FR hyperparameter sensitivity: median integrated  $L^2(F)$  over the grid (lower = faster convergence). Left:  $\gamma$  versus  $\eta$ ; right:  $\gamma$  versus burn-in fraction. The integrated error improves with larger  $\gamma$  and earlier burn-in; intermediate  $\eta$  is robust. These trends should not be extrapolated beyond this regime.

Table 5: Entropic-bottleneck base regime and Fisher–Rao hyperparameters.

| Quantity                                                    | Value     |
|-------------------------------------------------------------|-----------|
| Inverse temperature $\beta$                                 | 8         |
| Barrier height $H$                                          | 2.5       |
| Channel stiffness $\omega_{\text{out}}, \omega_{\text{in}}$ | 1, 25     |
| Channel width $s$                                           | 0.25      |
| Walkers $N$                                                 | 256       |
| Time step $\Delta t$                                        | $10^{-3}$ |
| Steps per run                                               | 40,000    |
| FR rate $\gamma$                                            | 15        |
| FR stride <code>fr_every</code>                             | 10        |
| Score clip $S_{\text{max}}$                                 | 3         |
| Event cap $f_{\text{max}}$                                  | 0.08      |
| Target EMA rate                                             | 0.005     |

## 6.2 Methods and design

The practical comparison is again **ABF only** versus **ABF+FR estimated** (14), on matched seeds (identical initial conditions and Langevin noise per seed). The diagnostics **FR uniform** (15) and **FR oracle** (16) (built from the analytic  $F_{\text{ref}}$ , hence not deployable, Remark 3) are also run. The FR hyperparameters are listed in Table 5; an assertion guards every batched call so the reference never reaches the estimated target. For the  $\omega_{\text{in}}$  sweep only, the stiffest channel is stabilised with a smaller time step and more steps at the same physical time horizon ( $dt = 4 \times 10^{-4}$ , 100,000 steps).

## 6.3 Results

At the base regime the practical method reduces the median final free-energy error from 0.201 (ABF) to 0.098 (ABF+FR estimated), a 51.3% reduction, winning 20 of 20 matched seeds (Figure 9). Two diagnostics are immediately revealing. The *uniform* target — which carries no information about the free-energy shape at all — recovers almost the entire gain (50.5%), while the *oracle* target (the

Entropic bottleneck: convergence (median and IQR over seeds)




Figure 9: Entropic bottleneck, convergence (median and inter-quartile range over seeds). ABF+FR estimated separates from ABF early and stays below it for the entire trajectory on both the free-energy (left) and mean-force (right) errors.

*analytic*  $F_{\text{ref}}$ ) is *worse than plain ABF* ( $-28.1\%$ ). This is the opposite of the metastability ranking (Section 5.6): here a target that over-commits to a fixed shape fights the ABF bias instead of merely equalising coverage. The signature points to *coverage equalisation along the reaction coordinate*, not free-energy-shape steering, as the source of the gain — which is the only kind of mechanism Proposition 2 leaves available, since selection rescales the marginal and provably does nothing else.

Earlier drafts of this report called that mechanism “balanced resampling / variance reduction”. The first half is right and the second is not, and the distinction matters enough to state once here and use throughout. What reallocation buys in a starved bin is not primarily a smaller error bar on an estimate that was already centred: it is samples in a bin whose ABF accumulator had barely been updated, and whose bias was therefore held near zero by the trust ramp. The deficit being repaired is systematic, not stochastic. The audit discussed in Section 18 makes the same point from the data side on WCA, where estimator variance is about 1% of the error budget and *rises* under directed balancing — a mechanism cannot be variance reduction if variance is neither the bulk of the error nor decreasing.

### 6.3.1 Temperature governs the gain, not the bottleneck width

Varying the inverse temperature  $\beta$  at fixed  $\omega_{\text{in}} = 25$  (Table 6) is the cleanest mechanistic signal in the study. FR *helps when sampling is hard and hurts when it is easy*: at high temperature ( $\beta = 2, 4$ ) crossings are frequent, ABF already samples the whole coordinate (its error is  $\sim 0.05$ ), and FR’s birth–death merely injects resampling noise — it is mildly harmful and loses most seeds ( $-26.1\%$  at  $\beta = 4$ ). At low temperature ( $\beta = 8, 12$ ) crossings are rare, ABF is starved on the far side of the barrier (error  $0.2\text{--}0.4$ ), and FR’s coverage equalisation substantially repairs the deficit, winning every seed. The cross-over sits between  $\beta = 4$  and  $\beta = 8$ .

By contrast, sweeping the bottleneck width  $\omega_{\text{in}}$  at fixed  $\beta = 8$  (Figure 10) leaves the gain essentially *flat* at  $\sim 50\%$ . This is the honest consequence of the caveat above: the ABF baseline error is nearly independent of  $\omega_{\text{in}}$  because the  $\omega$ -independent energetic barrier dominates the sampling difficulty. The gain is large and reliable across  $\omega_{\text{in}}$ , but it is governed by metastability (temperature), not by the bottleneck width. We do not claim a bottleneck-strength scaling the data do not show.

Table 6: Temperature sweep at  $\omega_{\text{in}} = 25$ : the FR gain flips sign as ABF crosses from well-sampled (hot) to starved (cold).

| $\beta$ | ABF $\ F - F_{\text{ref}}\ $ | FR $\ F - F_{\text{ref}}\ $ | gain (%) | win rate |
|---------|------------------------------|-----------------------------|----------|----------|
| 2       | 0.055                        | 0.060                       | -8.8     | 0.30     |
| 4       | 0.049                        | 0.062                       | -26.1    | 0.00     |
| 8       | 0.201                        | 0.098                       | +51.3    | 1.00     |
| 12      | 0.416                        | 0.281                       | +32.3    | 1.00     |

### 6.3.2 A wide safe regime for the FR rate

Sweeping the birth–death rate  $\gamma$  at the base regime (Figure 11), the gain increases *monotonically* up to  $\gamma = 50$  (+72.6%); we did not reach a failure boundary within the swept range. The windowed ancestor effective sample size falls steadily from 101 to 9 as lineage diversity collapses, yet the free-energy error keeps improving. This differs from the WCA dimer (Section 7), where aggressive birth–death eventually degrades the estimator; the likely reason is the very short conditional-relaxation time of the analytic  $y$ -channel here, so a cloned walker re-equilibrates its orthogonal coordinate almost immediately. FR is unusually robust in this model.

### 6.4 Mechanism: conditional fidelity from an analytic law

Because  $Y \mid X = x \sim \mathcal{N}(0, 1/(\beta\omega(x)^2))$  is *analytic* (21), we can test Proposition 1 directly rather than against a noisy histogram — the gap that Remark 4 left open in the metastability case. Figure 12 compares the empirical  $\widehat{\text{Var}}(Y \mid X \in I_j)$  to  $1/(\beta\omega(x_j)^2)$  at five  $x$ -bins. The result is reassuring but not uniformly dominant: FR tracks the analytic variance well across nearly three orders of magnitude (from  $\sim 2 \times 10^{-4}$  at the bottleneck  $x = 0$  to  $\sim 0.12$  in the wells), and is most competitive in the bottleneck and middle bins, but it is slightly worse than ABF in the well bins for this diagnostic. Thus the free-energy improvement is not coming from a visible collapse of the conditional law; it is achieved while the orthogonal variance remains close to the analytic conditional, with residual errors dominated by finite bin counts rather than a systematic FR bias.

### 6.5 An entropy-dominant stress test at fixed total barrier

The caveat of the previous subsections leaves one question open: the gain here tracks *temperature*, but because the energetic barrier dominates we never isolated whether an *entropic* barrier of the same height behaves any differently. We close that gap with a controlled stress test that promotes the transverse coordinate to  $m$  dimensions,

$$V(x, \mathbf{y}) = H(x^2 - 1)^2 + \frac{1}{2}\omega(x)^2 \|\mathbf{y}\|^2, \quad \mathbf{y} \in \mathbb{R}^m, \quad F_{\text{ref}}(x) = H(x^2 - 1)^2 + \frac{m}{\beta} \log \omega(x) + C, \quad (22)$$

so the marginal free energy carries an entropic term  $\frac{m}{\beta} \log \omega(x)$  that we can enlarge at will (the conditional law generalises to  $\mathbf{Y} \mid X = x \sim \mathcal{N}(0, I_m/(\beta\omega(x)^2))$ ). Writing the total thermal barrier as  $B_0 = \beta H + m \log(\omega_{\text{in}}/\omega_{\text{out}})$  and the *entropic share*  $\phi = m \log(\omega_{\text{in}}/\omega_{\text{out}})/B_0$ , we set  $H = (1 - \phi)B_0/\beta$  and  $\omega_{\text{in}} = \omega_{\text{out}} e^{\phi B_0/m}$ , which holds the total barrier *fixed* while sliding it from purely energetic ( $\phi = 0$ ) to strongly entropic ( $\phi = 0.9$ ). We sweep  $\phi \in \{0, 0.25, 0.5, 0.75, 0.9\}$  at  $\beta = 4$ ,  $m = 2$ ,  $B_0 = 8$  (20 matched seeds, time horizon  $T = 40$ ); Table 7 lists the design and results.

The outcome is sharper than a simple yes/no, and is the most useful thing this case returns (Figure 13). The *deployable* estimated-target gain does *not* grow with the entropic share: it is U-shaped in  $\phi$  (slope  $-8.3\%$  per unit  $\phi$ , correlation  $r = -0.28$ ) and is mildly *harmful* in the middle of






Figure 10: Bottleneck-width sweep  $\omega_{in}$  at  $\beta = 8$ . The FR gain is large ( $\sim 50\%$ ) and 100% reliable at every width, but it does *not* grow with the bottleneck strength: the ABF baseline error is flat because the energetic barrier height  $H = 2.5$  dominates the smaller entropic contribution  $\beta^{-1} \log(\omega_{in}/\omega_{out})$ .

the sweep ( $-9.7\%$  at  $\phi = 0.5$ ), where the smooth entropic free-energy bump is hardest to estimate online from the high-variance instantaneous force  $\omega(x)\omega'(x) \|\mathbf{y}\|^2$ . The uniform target behaves almost identically, reaffirming that the practical mechanism is coverage equalisation rather than free-energy-shape steering. The *oracle* target tells the opposite story: handed the analytic  $F_{ref}$ , its gain rises steeply with the entropic share (slope 64% per unit  $\phi$ ,  $r = 0.94$ ), from harmful at  $\phi = 0$  to  $+43.7\%$  at  $\phi = 0.75$ . There is therefore a genuinely *entropic-specific* opportunity for birth-death — but it is *target-limited*: realising it needs a free-energy-shaped target accurate enough to steer against, which the online ABF-bias estimate is not in this regime. What the deployable method does deliver across the whole sweep is faster convergence: its integrated-error gain is positive at every  $\phi$  (median  $+20.5\%$ ), and all FR variants preserve the  $m$ -dimensional conditional law as in the scalar case.

## 6.6 Why the deployable gap is an obstruction, not a tuning failure

The stress test leaves a natural hope: if the online ABF-bias estimate is merely *too noisy* to steer against, a better online free-energy estimator inside the target should convert the oracle’s entropic headroom into a deployable gain. We tested this hope directly and report a negative result that we believe is the more useful finding, because it is structural. The mean force along the reaction coordinate is

$$\langle f_x | x \rangle = 4Hx(x^2 - 1) + \omega(x)\omega'(x) \langle \|\mathbf{y}\|^2 | x \rangle, \quad \langle \|\mathbf{y}\|^2 | x \rangle = m/(\beta\omega(x)^2), \quad (23)$$




Figure 11: FR-rate sweep  $\gamma \in \{1, \dots, 50\}$  at the base regime. Left: gain over ABF increases monotonically (no failure boundary in range). Right: windowed ancestor ESS collapses ( $101 \rightarrow 9$ ) and the replacement fraction rises, yet accuracy keeps improving.

Table 7: Entropy-dominant stress test ( $\beta = 4$ ,  $m = 2$ ,  $B_0 = 8$ , 20 seeds,  $T = 40$ ): at fixed total barrier, the deployable estimated/uniform gains are small and non-monotone in  $\phi$ , while the oracle gain rises sharply — the entropic-specific headroom is real but target-limited.

| $\phi$ | $H$  | $\omega_{\text{in}}$ | ABF $\ F - F_{\text{ref}}\ $ | est. (%) | est. wins | uni. (%) | oracle (%) |
|--------|------|----------------------|------------------------------|----------|-----------|----------|------------|
| 0      | 2.00 | 1.0                  | 0.0464                       | +13.0    | 20/20     | +13.0    | -6.5       |
| 0.25   | 1.50 | 2.7                  | 0.0360                       | +4.7     | 15/20     | +6.0     | -5.7       |
| 0.5    | 1.00 | 7.4                  | 0.0286                       | -9.7     | 4/20      | -14.3    | +9.2       |
| 0.75   | 0.50 | 20.1                 | 0.0248                       | -8.1     | 7/20      | -0.5     | +43.7      |
| 0.9    | 0.20 | 36.6                 | 0.0312                       | +12.7    | 15/20     | +12.9    | +42.4      |

Total barrier  $B_0 = \beta H + m \log(\omega_{\text{in}}/\omega_{\text{out}})$  ( $= 8$ ) is held fixed while  $\phi$  slides it from energetic to entropic ( $H = (1 - \phi)B_0/\beta$ ,  $\omega_{\text{in}} = \omega_{\text{out}}e^{\phi B_0/m}$ ). Gains are matched-seed medians in final  $\|F - F_{\text{ref}}\|$  vs ABF; the oracle target uses the analytic  $F_{\text{ref}}$  and is a diagnostic control, *not* deployable.

and the ABF bias  $B(x) = \int^x \langle f_x | x' \rangle dx'$  is the sufficient nonparametric estimator of  $F$  from the binned per-visit force samples when no external model of the transverse equilibrium is supplied. This is the key reason the entropy-dominant results are not better: the deployable target is not missing a hyperparameter, it is missing information. We constructed several deployable variants that never read  $F_{\text{ref}}$  and asked whether any beats  $B$  as either a readout or an FR target.

(i) *Moment fusion.* Estimate the entropic part from the transverse log-moment,  $F_{\text{ent}}(x) = -\frac{m}{2\beta} \log(\|\mathbf{y}\|^2 | x)$ , and the energetic part from the  $x$ -channel. Because the entropic mean force obeys  $dF_{\text{ent}}/dx = (m/\beta)\omega'/\omega = \omega\omega'\langle\|\mathbf{y}\|^2 | x\rangle$ , the reconstruction telescopes to  $B$  *identically* ( $\max_x |F_{\text{moment}} - B| = 6 \times 10^{-4}$  on the grid; median moment/ABF readout ratio 0.994 over 20 seeds). The  $y$ -channel moment is redundant with the  $x$ -force once integrated.

This redundancy also explains why the resulting FR target is weak. The deployable birth–death target has the form

$$q_{\hat{F}}(x) \propto \exp\{-\beta(\hat{F}(x) - B(x))\}. \quad (24)$$

If  $\hat{F} = B + \varepsilon$ , then  $q_{\hat{F}}(x) \propto \exp\{-\beta\varepsilon(x)\}$  after normalisation. Thus the true free-energy shape cancels out of the practical target: when  $\varepsilon \simeq 0$  the target is essentially uniform, and when  $\varepsilon$  is nonzero the birth–death score steers against estimator noise rather than against the oracle shape.




Figure 12: Conditional fidelity against the *analytic* law  $\text{Var}(Y | X = x) = 1/(\beta\omega(x)^2)$ . Empirical FR and ABF variances both track the analytic curve across three orders of magnitude. FR is most competitive near the bottleneck and middle bins, while ABF is slightly closer in the well bins; the diagnostic shows no conditional-law collapse.

This is the mathematical reason `fr_moment` and `fr_uniform` cluster together.

(ii) *Regression intercept.* Forcing a non- $B$  signal by regressing  $f_x$  on  $\|\mathbf{y}\|^2$  per bin (so the slope absorbs  $\omega\omega'$ ) is *worse* than  $B$  at every controlled per-bin sample count (ABF/regression RMSE ratio 0.64–0.73, i.e. the regression readout has the larger error), because estimating the slope re-injects variance; it also divides by  $\text{Var}(\|\mathbf{y}\|^2) \rightarrow 0$  in the tight channel and destabilises the feedback loop.

The variance penalty is not accidental. With  $u = \|\mathbf{Y}\|^2 | X = x$ ,

$$\mathbb{E}[u | x] = \frac{m}{\beta\omega(x)^2}, \quad \text{Var}(u | x) = \frac{2m}{\beta^2\omega(x)^4}. \quad (25)$$

The instantaneous entropic force contribution  $\omega(x)\omega'(x)u$  therefore has conditional variance

$$\text{Var}(\omega\omega'u | x) = \frac{2m}{\beta^2} \left( \frac{\omega'(x)}{\omega(x)} \right)^2, \quad (26)$$

and its standard error is reduced only by the number of visits in the bin. Those visits are precisely scarce in the barrier region. Estimators that try to separate the energetic and entropic pieces must estimate additional local slopes or moments from the same scarce samples, so they trade the ABF sample mean for a higher-variance object.

(iii) *Smoothing and confidence gating.* Wide-smoothing the energetic residual  $B - F_{\text{ent}}$  can break the identity  $F_{\text{moment}} \equiv B$ , but only by introducing a deterministic bias in the target. Because Equation (24) is exponential in this error, a biased  $\hat{F}$  creates a biased cloning/deletion field rather

Entropy-dominant stress test: the achievable (oracle) gain is entropic-specific, the deployable gain is not




Figure 13: Entropy-dominant stress test at fixed total barrier  $B_0$ . (a) The barrier decomposition: as  $\phi$  grows the energetic part  $\beta H$  is traded for the entropic part  $m \log(\omega_{\text{in}}/\omega_{\text{out}})$  with the total held constant. (b) Matched-seed gain over ABF: the deployable estimated (and uniform) targets are small and non-monotone in  $\phi$ , whereas the oracle gain rises sharply with the entropic share — the achievable benefit is entropic-specific, but the practical online estimate does not capture it.

than a harmless post-processing error. Per-bin confidence gating has the opposite pathology: it shrinks the very score that a successful target would need near the bottleneck. In the confirmation sweep the gated target is worse than the ungated moment target, which indicates that the removed component was not reliable shape information but noisy steering.

(iv) *Two-dimensional control variate.* Stratifying samples by the transverse coordinate  $\|\mathbf{y}\|^2$  and recombining per-stratum mean forces — the textbook “resolve the orthogonal coordinate” device — is  $2.6\times$  more accurate than plain ABF *when the stratum weights are the analytic  $\chi^2$  law*, but reduces to plain ABF *identically* once the weights are estimated from the same data:

$$\sum_k \frac{n_k}{n} \bar{f}_k = \sum_k \frac{n_k}{n} \frac{1}{n_k} \sum_{i \in k} f_i = \frac{1}{n} \sum_i f_i. \quad (27)$$

The measured difference is  $3 \times 10^{-16}$  (machine precision). The entire gain is the injected conditional law, not the stratification.

The same variance logic applies to density- or CZAR-style free-energy targets. If a bin of width  $h$  has  $n_x$  effective visits, a density estimate has relative noise of order  $n_x^{-1/2}$ , so the log-density free energy has  $\text{Var}(\hat{F}(x)) = O((\beta^2 n_x)^{-1})$ . A steering target, however, uses the *spatial shape* of this estimate. Differencing a noisy log-density to obtain a score or mean-force-like signal adds powers of  $h^{-1}$ , while an extended-system restraint contributes its own force variance. Thus the estimator can be asymptotically sensible but still be a bad online control signal at the finite counts available in the bottleneck.

All variants collapse to one statement: any deployable rearrangement of the run’s own observables either returns  $B$  exactly or departs from it only by added bias or estimation variance, so the steering signal  $F_{\text{target}} - B$  is noise rather than shape. The oracle’s advantage — and the control variate’s analytic-weight advantage — is *information that is not in the trajectory* (the true  $F_{\text{ref}}$ , equivalently the exact transverse equilibrium). There is also a chicken-and-egg barrier specific to steering: learning that the bottleneck is a high- $F$  region requires *sampling* it, the same starvation that limits  $B$ ;  $\langle \|\mathbf{y}\|^2 \mid x \rangle$  carries no earlier information than the integrated force. A 20-seed confirmation bears this out in dynamics: the moment target (`fr_moment`) and a per-bin confidence-gated

variant (`fr_moment_gated`) both sit in the ABF/uniform band ( $-5.6\%$  and  $-19.8\%$  at  $\phi = 0.75$ ;  $15.2\%$  and  $8.4\%$  at  $\phi = 0.9$ ), far from the oracle’s  $+43.7\%$ , and gating *removes* rather than improves the signal — confirming it was noise. The practical consequence is that the entropic-specific headroom is reachable only by *adding information* (enhanced sampling of a genuinely slow orthogonal coordinate — not the case here, where  $\mathbf{y}$  equilibrates fast — multiple-replica or tempered data, or a parametric prior on  $\omega(x)$  fit from the well interiors), not by a cleverer estimator on the same trajectory.

## 6.7 Comparison to OPES

As in the WCA case (Section 7.13) we check the starved-regime claim against a modern adaptive-bias method: well-tempered OPES [9] on the identical bottleneck dynamics ( $\beta=8$ ,  $\omega_{\text{in}}=25$ ), reconstructed by the same mean-force integration and evaluated on the same interior window, with the deposition barrier tuned on held-out seeds and fixed to 4 kT. OPES reaches a median  $L^2(F)=0.3252$  over five seeds, against ABF’s 0.2098 and mFR’s 0.0954: it is *worse than plain ABF* here ( $-55.0\%$ ), and roughly  $3.4\times$  less accurate than mFR ( $-240.9\%$ ). The reweighting effective sample fraction is low ( $n_{\text{eff}}=0.0871$ ), consistent with a bias that flattens the  $x$ -marginal aggressively but leaves the *conditional*  $\mathbf{y}$ -average — the entropic part of the mean force — as noisy as under ABF. This is the same lesson as Section 6.6: the entropic obstruction is about *which* average enters the force, and flattening the sampled coordinate does not address it, whereas mFR’s ensemble reallocation does. OPES does not close the gap that mFR closes.

## 6.8 Case finding

The entropic bottleneck is a controlled bridge between the easy metastability case and the hard WCA dimer. Estimated-target FR robustly and substantially reduces the free-energy error in the cold, sample-starved regime ( $\sim 51.3\%$ , 20/20 seeds), but the gain is governed by *temperature* (FR helps when ABF is starved, mildly hurts when it is not), not by the bottleneck width. The mechanism is coverage equalisation, not shape steering and not variance reduction: the uniform target nearly matches the estimated one and the oracle target is *worse* than ABF — the reverse of the metastability ranking. The safe FR regime is wide (no failure boundary up to  $\gamma = 50$ ), and the analytic conditional diagnostic shows that the gain is obtained without a systematic collapse of  $Y \mid X = x$ , even though FR is not uniformly better than ABF on every conditional-variance bin. Finally, the dedicated entropy-dominant stress test (Section 6.5) settles the bottleneck-versus-temperature question that the  $\omega_{\text{in}}$  sweep could not: holding the *total* barrier fixed and growing only the entropic share, the deployable gain still does not increase (it is U-shaped in  $\phi$ ), whereas the oracle gain rises sharply to  $+43.7\%$  at  $\phi = 0.75$ . The entropic-specific benefit is therefore real but *target-limited*, and Section 6.6 shows the limit is an *obstruction rather than a tuning failure*: every deployable estimator built from the trajectory’s own observables (moment fusion, regression/control variates, smoothing, and confidence gating) reduces to the ABF bias or to the ABF bias plus noise, so the oracle is a provable *information* upper bound that no online estimator on the same data can reach.

## 7 Case III: a condensed-phase WCA dimer

The third system is many-body and is where ABF is most starved: a dimer that must push a dense solvent aside to change its bond length. It tests whether the two-dimensional findings survive in a setting with thousands of degrees of freedom and a nonlinear collective variable.

Table 8: WCA dimer base system and tuned Fisher–Rao hyperparameters.

| Quantity                             | Value   |
|--------------------------------------|---------|
| Inverse temperature $\beta$          | 1       |
| Time step $\Delta t$                 | 0.002   |
| Steps per run                        | 250,000 |
| Replicas $N$                         | 1024    |
| Lattice spacing $a$                  | 1.5     |
| Dimer barrier $h$ , width $w$        | 2, 2    |
| FR rate <code>fr_rate</code>         | 0.10    |
| FR onset <code>fr_start_steps</code> | 20,000  |
| FR stride <code>fr_every</code>      | 5       |
| Score clip $S_{\max}$                | 2.0     |
| Event cap $f_{\max}$                 | 0.02    |
| Target EMA rate                      | 0.005   |

## 7.1 Model and reference

A dimer of two tagged particles sits in a periodic two-dimensional bath of  $n_{\text{dim}}^2 - 2$  Weeks–Chandler–Andersen (WCA) solvent particles ( $n_{\text{dim}} = 10$ , lattice spacing  $a$ , box  $L = n_{\text{dim}} a$ ). All particles interact through the purely repulsive WCA potential [14] (Lennard–Jones cut and shifted at its minimum  $r_0 = 2^{1/6}\sigma$ ). The two dimer particles additionally feel a symmetric double-well bond  $V_{\text{dim}}(r) = h(1 - ((r - r_0 - w)/w)^2)^2$  with barrier height  $h = 2$  and width  $w = 2$ . The reaction coordinate is the scaled bond length

$$z(q) = \frac{|q_1 - q_2| - r_0}{2w}, \quad (28)$$

a nonlinear collective variable of the full configuration  $q \in \mathbb{R}^{2n}$  (Remark 1):  $z \approx 0$  is the *compact* state,  $z \approx 1$  the *stretched* state, and  $z \in [0.25, 0.75]$  the *transition* region. The free-energy barrier near  $z = 0.5$  is partly entropic / crowding-driven, since the dimer must displace solvent to cross it. Dynamics are overdamped Langevin at  $\beta = 1$ ,  $dt = 0.002$ , with the minimum-image convention; the default system is  $a = 1.5$ ,  $\sigma = \epsilon = 1$ . The local mean force along  $z$  is the constrained projection of the physical forces, including the geometric contribution with the sign convention of the estimator (Remark 1). The reference  $F(z)$  is computed once per system  $a$  by constrained-MD thermodynamic integration [3, 7]; it is a high-accuracy numerical reference used for evaluation only, not an analytic exact solution. Errors are taken on the interior window  $z \in [-0.1, 1.1]$  with the additive constant aligned.

## 7.2 Methods and design

As before the practical comparison is **ABF only** versus **ABF+FR estimated** (14), with **FR uniform** and **FR oracle** diagnostics; here the oracle target is built from the TI reference and is explicitly *not* deployable (Remark 3). A no-leakage assertion ensures only the oracle method ever receives the reference, and a clone copies the *entire* WCA configuration — not just  $z$  — so cloned solvent arrangements travel with the bond length. Replica lineage is tracked to report the ancestor effective sample size  $\text{ESS} = 1/\sum_a w_a^2$ . The tuned FR configuration is given in Table 8; the rate is named `fr_rate` and the onset `fr_start_steps` (Section 4.5).

Table 9: WCA main comparison (250k steps,  $N = 1024$ ,  $a = 1.5$ ). The tuned row is the practical headline; the strong row shows the accuracy–diversity tradeoff. The mean-force gain is modest.

| Method                    | $\ F - F_{\text{ref}}\ $ | $\ F' - F'_{\text{ref}}\ $ | gain (%) | wins  | anc. ESS |
|---------------------------|--------------------------|----------------------------|----------|-------|----------|
| ABF only                  | 0.0835                   | 0.3220                     | —        | —     | —        |
| FR estimated (tuned)      | 0.0361                   | 0.2513                     | +55.0    | 10/10 | 86       |
| FR uniform                | 0.0372                   | 0.2525                     | +54.6    | 10/10 | 90       |
| FR oracle*                | 0.0387                   | 0.2475                     | +52.7    | 10/10 | 100      |
| FR estimated (strong)     | 0.0329                   | 0.2590                     | +58.9    | 10/10 | 45       |
| FR estimated (aggressive) | 0.0935                   | 0.3883                     | -13.1    | 1/10  | 22       |

\* The oracle target uses the thermodynamic-integration reference  $F_{\text{ref}}$ ; it is a diagnostic control, *not* a deployable method. Medians over 10 seeds.




Figure 14: WCA dimer, convergence (median and inter-quartile range over seeds). ABF and FR are identical until FR activates; from then on FR is consistently below ABF on  $L^2(F)$  for the rest of the run. The cumulative replacement count (right) confirms the birth–death is active but bounded.

### 7.3 Results

The main comparison (Table 9; 10 seeds,  $N = 1024$ , 250 000 steps) shows the practical method **reliably beating ABF**: the tuned FR cuts the median final  $L^2(F)$  from 0.0835 to 0.0361 (+55.0%), winning all 10/10 matched seeds. As in the entropic-bottleneck case, the uniform target (+54.6%) and the oracle target (+52.7%) are *essentially as good as* the estimated target — the online estimated target already extracts nearly all the available benefit, and there is no headroom to the oracle. The stronger estimated-target row in the same table attains the lowest final free-energy error, but with a much smaller ancestor ESS; we therefore use the tuned row as the practical headline and the strong row as the accuracy–diversity contrast. We are explicit that the headline gain is in the *free energy*: the mean-force error  $L^2(F')$  improves only modestly ( $\approx 22\%$ ), so we do not claim a large mean-force improvement.

### 7.4 Mechanism: support reallocation with a diversity cost

The diagnostics support the same broad mechanism as the entropic bottleneck rather than a stationary-coverage reshaping one. Final reaction-coordinate region fractions and the kernel-weighted ABF denominator in the transition window are nearly identical across ABF, FR estimated,




Figure 15: Final profiles (median over seeds) versus the TI reference (dashed): (a) free energy  $\hat{F}(z)$ , (b) mean force  $\hat{F}'(z)$ , (c) reaction-coordinate marginal. FR recovers the reference free energy more accurately than ABF, chiefly in and around the transition region.




Figure 16: Mechanism. (a) The kernel-weighted ABF denominator  $N_{\text{eff}}(z)$  is raised by FR in the transition window; (b) per-bin mean-force error is reduced there; (c) birth (clone) and death locations concentrate around the over/under-sampled regions; FR reshapes the *sampling support*, not the stationary marginal.

FR uniform and FR oracle: FR does *not* converge to a different stationary  $z$ -distribution. Instead the gain appears immediately after FR activates and is sustained for the rest of the run. Figure 16 shows the interpretation: birth–death acts as a *support-reallocation resampler for the ABF estimator*. Replicas stuck in over-represented configurations are replaced by whole-configuration clones of under-represented ones, improving weighted force-sample support in important  $z$ -bins even when the marginal looks unchanged. A clone is not an independent solvent arrangement at birth; the point is that useful whole configurations are moved into starved reaction-coordinate regions and then decorrelated by subsequent Langevin dynamics. This explains why target shape matters little here, consistent with uniform  $\approx$  estimated  $\approx$  oracle.

## 7.5 Failure boundary

Unlike the entropic bottleneck, the WCA dimer exhibits a clear failure boundary in the FR rate (Figure 17). The rate ladder (a shorter 150k-step, 4-seed stage) maps a clean safe regime: too gentle ( $\text{fr\_rate} \leq 0.05$ ) gives a small gain (+19–29%); the sweet spot ( $\text{fr\_rate} \approx 0.1$ –0.2) gives the maximal +58–64% with healthy ancestor ESS; the gain then *erodes* sharply ( $\text{fr\_rate} = 0.5$  falls to +12% with the error already climbing back toward the ABF baseline) and finally *reverses* for






Figure 17: FR-rate ladder (150k steps, 4 seeds). Final  $L^2(F)$  (left) is minimised in the 0.1–0.2 sweet spot, climbs back toward the ABF baseline (dashed) by  $\text{fr\_rate} = 0.5$ , and rises above it for  $\text{fr\_rate} \geq 1$ ; ancestor ESS (right) collapses as the rate grows, pinpointing diversity loss as the failure mechanism.

$\text{fr\_rate} \geq 1$ , where the median final  $L^2(F)$  is worse than ABF (−38% at rate 1, −68% at rate 2) and no seed wins, as the ancestor ESS collapses from  $\sim 337$  to  $\sim 20$ . The reversal arrives even earlier under a longer budget: the production aggressive configuration ( $\text{fr\_rate} = 0.5$  at 250k steps, 10 seeds) is already net-harmful (−13.1%, winning only 1/10 seeds, ancestor ESS  $\sim 22$ ), because the extra steps give diversity loss more time to dominate. Over-aggressive birth–death destroys the replica diversity the mean-force estimator depends on, so the estimator degrades faster than coverage improves. This is the predicted failure mode; the analytic  $y$ -channel of the entropic bottleneck masks it (Section 6) because cloned walkers re-equilibrate instantly, whereas a dense solvent does not.

## 7.6 Difficulty dependence

Varying one difficulty axis at a time (Figure 18) shows the gain is robust and, importantly, *grows with difficulty*. With *budget*, FR wins at every  $n_{\text{steps}}$  and the relative advantage grows rather than washes out (+35.5% at 50k to +56.9% at 400k): ABF’s error plateaus around 0.08 after  $\sim 150$ k while FR keeps descending, so FR is not merely an early-time accelerator that ABF later catches. With *replica count*, the gain is roughly flat (+51–59% from  $N = 256$  to 2048) — FR helps even at  $N = 256$ . With *solvent crowding*, the densest bath ( $a = 1.35$ ) makes the entropic barrier severe and ABF  $\sim 4\times$  worse in absolute error; FR still wins 5/5 and delivers by far the largest absolute error reduction ( $0.365 \rightarrow 0.225$ ). In the units a practitioner cares about — absolute free-energy error removed — the crowded regime is where FR helps most.

## 7.7 WCA phase diagram: when does marginal Fisher–Rao help?

A single WCA setting is not enough to tell whether the gain above is a generic property of the dimer or a response to sample starvation. We therefore added an additive, command-line phase-diagram study that sweeps physical parameters while keeping the existing production result intact. In the production phase diagram reported here we vary the inverse temperature  $\beta \in \{1, 2, 4\}$  and dimer barrier height  $h \in \{1, 2, 4, 6\}$  at fixed width  $w = 2$ , lattice spacing  $a = 1.5$ , and  $M = n_{\text{dim}}^2 = 100$  physical particles; we also add a targeted physical-size slice  $M \in \{49, 100, 196\}$  at  $(\beta, h) = (2, 4)$ .




Figure 18: Difficulty dependence (matched-seed gain over ABF). (a) Budget: the relative gain grows with  $n_{\text{steps}}$  as ABF plateaus and FR keeps improving. (b) Replica count: gain is roughly flat. (c) Solvent crowding: smaller  $a$  (denser) yields the largest absolute error reduction. FR wins 5/5 seeds in every cell.

Table 10: WCA phase diagram over  $(\beta, h)$  at  $M = 100$  physical particles ( $N = 1024$  replicas, 120,000 steps).  $R = L^2(F)_{\text{ABF}}/L^2(F)_{\text{mFR est}}$ ; gain is the matched-seed median percentage reduction in final  $L^2(F)$ . The oracle column is a diagnostic that uses the TI reference and is *not* deployable.

| $\beta$ | $h$ | $\beta h$ | ABF    | mFR est | mFR uni | mFR ora* | $R$  | gain (%) | wins | ESS |
|---------|-----|-----------|--------|---------|---------|----------|------|----------|------|-----|
| 1       | 1   | 1         | 0.1006 | 0.0503  | 0.0518  | 0.0521   | 2.00 | +49.9    | 4/4  | 171 |
| 1       | 2   | 2         | 0.0918 | 0.0474  | 0.0409  | 0.0480   | 1.94 | +48.0    | 4/4  | 166 |
| 1       | 4   | 4         | 0.0954 | 0.0465  | 0.0463  | 0.0530   | 2.05 | +51.3    | 4/4  | 191 |
| 1       | 6   | 6         | 0.0960 | 0.0510  | 0.0523  | 0.0601   | 1.88 | +46.8    | 4/4  | 238 |
| 2       | 1   | 2         | 0.0130 | 0.0132  | 0.0161  | 0.0109   | 0.98 | -6.5     | 1/4  | 224 |
| 2       | 2   | 4         | 0.0127 | 0.0161  | 0.0160  | 0.0076   | 0.79 | -29.0    | 0/4  | 218 |
| 2       | 4   | 8         | 0.0150 | 0.0243  | 0.0236  | 0.0103   | 0.62 | -61.5    | 0/4  | 209 |
| 2       | 6   | 12        | 0.0246 | 0.0196  | 0.0195  | 0.0167   | 1.26 | +20.4    | 4/4  | 275 |
| 4       | 1   | 4         | 0.0045 | 0.0253  | 0.0226  | 0.0079   | 0.18 | -488.8   | 0/4  | 212 |
| 4       | 2   | 8         | 0.0082 | 0.0282  | 0.0279  | 0.0107   | 0.29 | -229.8   | 0/4  | 212 |
| 4       | 4   | 16        | 0.0156 | 0.0397  | 0.0388  | 0.0171   | 0.39 | -156.0   | 0/4  | 165 |
| 4       | 6   | 24        | 0.0249 | 0.0414  | 0.0414  | 0.0231   | 0.60 | -69.1    | 0/4  | 153 |

\* Diagnostic oracle target (uses  $F_{\text{ref}}$ ); not deployable.

This  $M$  is the number of physical particles in one WCA system; it is distinct from the  $N = 1024$  ABF/FR replicas used to estimate the mean force. Each cell uses 4 matched seeds and 120000 steps. The deployable comparison is ABF versus mFR–ABF with the online estimated target; the uniform target is an ablation, and the oracle target is included as a diagnostic positive control in the table but is not a usable method because it reads the TI reference.

The result is sharper than a monotone “larger  $\beta h$  is harder” story. The measured ABF error, not the nominal product  $\beta h$ , is the useful starvation diagnostic. In the  $\beta = 1$  row ABF remains at  $L^2(F) \approx 0.09$ – $0.10$ , and mFR improves all matched seeds with +47–51% median gain; the anchor cell  $\beta = 1, h = 2$  improves from 0.0918 to 0.0474 ( $R = 1.94$ ). At  $\beta = 2$  and  $\beta = 4$ , ABF is often much more accurate (below  $L^2(F) = 0.005$  in the easiest accurate cell), and the same birth–death configuration over-corrects: only 5/12 cells have positive gain, and the worst cell  $\beta=4, h=1$  loses strongly. Thus the phase diagram supports the conditional claim rather than an unconditional one: mFR helps when the ABF free-energy estimator is sample-starved, but it can hurt when ABF has




Figure 19: Production WCA phase diagram at  $M = 100$ : ABF final  $L^2(F)$  (left), mFR-ABF with the estimated target (middle), and the matched-seed improvement ratio  $R = L^2(F)_{\text{ABF}}/L^2(F)_{\text{mFR}}$  (right). Values  $R > 1$  mean mFR improves over ABF. The gain localises in the high-error  $\beta = 1$  row, including the anchor cell; when ABF is already accurate, mFR is neutral to harmful.

already solved the marginal problem.

The mechanism diagnostics separate exploration from target-shape and genealogy effects. The uniform target is close to the estimated target in this production run (mean absolute difference in matched gain 10.77 percentage points). The oracle target is different: it improves most cells, which shows that birth–death can help beyond the  $\beta = 1$  row if the target is accurate, but that target is diagnostic only. FR also increases barrier crossings only modestly; the important pattern is that it reallocates whole configurations in the cells where the ABF estimator is noisy. Genealogy does not show catastrophic collapse in the production run: the minimum final ancestor ESS for mFR-ABF is 149 out of 1024, the mean final ESS is 213, and the largest single ancestor fraction observed is only 0.0273. The harmful cells therefore look like an unnecessary resampling/variance tradeoff, not the severe genealogical collapse seen in the aggressive-rate failure ladder.

This production phase diagram therefore refines the WCA conclusion. There is an “easy/accurate” region where ABF already works well and mFR should be reduced or disabled, a sample-starved region (the  $\beta = 1$  row, including the original WCA anchor) where gentle mFR gives a large free-energy gain, and intermediate cells where the answer depends on the online target. The targeted mass slice at  $(\beta, h) = (2, 4)$  is non-monotone: mFR helps at  $M = 49$  (+24.4%), but hurts at  $M = 100$  and  $M = 196$ , while the oracle target helps the two larger systems. Thus physical size alone is not the difficulty variable; the relevant diagnostic is whether the practical online target is correcting a genuinely starved ABF estimator without adding unnecessary birth–death noise. This is distinct from the separate over-aggressive failure mode identified by the rate ladder, where diversity collapse dominates.

## 7.8 Starvation diagnostics: measured ABF error, not $\beta h$

To make the conditional claim quantitative and reviewer-checkable we computed a per-run starvation diagnostic for every cell of the production phase diagram (`scripts/analyze_wca_starvation.py`; a pure re-analysis of the existing runs, no new dynamics). For each cell we record the measured ABF baseline error, the kernel-weighted ABF support in the transition window, the per-bin mean-force error, the reaction-coordinate marginal distances, the birth–death event fractions, and the ancestor genealogy, and we pair them seed-by-seed with the matched mFR gain.

The headline diagnostic plot (Figure 23) puts the measured ABF final  $L^2(F)$  on the  $x$ -axis and




Figure 20: Barrier-crossing counts for ABF and mFR-ABF (estimated target) in the same production cells. mFR slightly increases crossing counts in most cells, but the accuracy gain does not follow crossing count alone: unnecessary resampling can raise crossings while worsening  $L^2(F)$  when ABF is already accurate.

the matched-seed median mFR gain on the  $y$ -axis. The relationship is monotone and strong: across all 14 production cells the rank correlation between ABF error and mFR gain is  $\rho = 0.80$ , whereas the correlation with the nominal difficulty product  $\beta h$  is only  $\rho = -0.57$  and of the wrong sign for a “harder-is-better” reading. Thresholding the measured error at  $L^2(F) \geq 0.05$  (starved) and  $\leq 0.02$  (easy) cleanly separates the regimes: the 4 starved cells (the entire  $\beta = 1$  row) gain a median +49% (+47 to +51%), while the 8 easy cells lose a median −73%, with the worst,  $\beta=4, h=1$ , at −489%. This is the empirical backing for using the *measured* estimator error, which a practitioner can monitor online through convergence diagnostics, rather than the physical product  $\beta h$ , as the decision variable.

The supporting scatters (Figure 24) refine the mechanism. The final kernel-weighted ABF support in the transition window is *not* a strong discriminator (it is within a few percent of the median-bin support in every regime, because ABF eventually flattens the  $z$ -marginal), and the barrier-crossing ratio between mFR and ABF does not predict the accuracy gain: unnecessary resampling can raise crossings while worsening  $L^2(F)$ . The ancestor ESS, by contrast, stays high for the deployable mFR in the helpful cells. Together these confirm that the gain is an estimator-variance effect localised to sample-starved cells, not a stationary-coverage reshaping, and that the safe operating regime is bounded by ancestor diversity.

## 7.9 Adaptive diversity-aware Fisher–Rao

The phase diagram motivates a deployable schedule that turns the birth–death on where ABF is starved and self-limits where it is not, without ever reading the reference. We add a method `fr_estimated_adaptive` (additive to, not a replacement for, the fixed-rate `fr_estimated`) whose effective rate at each FR event is a product of three online gates,

$$\gamma_{\text{eff}} = \gamma_0 g_{\text{support}} g_{\text{diversity}} g_{\text{event}}, \quad \gamma_{\text{eff}} \in [\gamma_{\min}, \gamma_{\max}].$$




Figure 21: Genealogical diagnostics for representative production regimes. The deployable mFR run keeps many ancestors alive in the helpful  $\beta = 1$  row, and the maximum ancestor fraction remains small; the production run does not exhibit the collapse associated with over-aggressive birth–death.




Figure 22: Error-versus-time curves for representative cells. mFR separates from ABF after the FR onset in the sample-starved anchor regime, while in accurate high- $\beta$  regimes the same resampling is not beneficial.

The support gate  $g_{\text{support}}$  ramps on the  $L^2$  distance between a low-variance cumulative estimate of the particle marginal  $\hat{p}$  and the *uniform* marginal, i.e. the ABF-driven non-flatness of the reaction-coordinate marginal (broadly smoothed, accumulated from a burn-in, and held inactive for a short warm-up so the estimate settles before the gate acts). We measure the distance to uniform rather than to the FR target  $q$  deliberately: the latter is shrunk by the very birth–death it gates (a feedback loop), whereas the marginal non-flatness is largely set by ABF. The diversity gate  $g_{\text{diversity}}$  tapers linearly in the ancestor ESS fraction (full above 0.25, off below 0.10), so the schedule halts before genealogical collapse — precisely the failure mode of the over-aggressive rate ladder (Figure 17). The event gate  $g_{\text{event}}$  backs the rate off when the realised event fraction approaches the `max_event_fraction` cap. Every gate and its inputs are logged at each FR event (`adaptive_fr_event_log.csv`).

A limitation, which we state plainly because it sharpens the paper’s thesis, is that *no marginal-shape online signal cleanly separates the harmful cells from the helpful ones*. The cold (high- $\beta$ ) accurate cells have genuinely sharp, bimodal  $z$ -marginals whose non-flatness is as large as that of the warm, starved cells, so a marginal-based gate cannot tell “cold and accurate” from “warm and




Figure 23: Starvation diagnostic. ABF final  $L^2(F)$  (measured baseline error, log scale) versus the matched-seed median gain of the deployable mFR over ABF, annotated by  $(\beta, h)$  and coloured by regime. The gain is large and positive exactly where ABF is inaccurate (the starved  $\beta = 1$  cluster) and turns negative as ABF becomes accurate; the relationship with measured error ( $\rho = 0.80$ ) is far cleaner than with  $\beta h$  ( $\rho = -0.57$ ).

starved”. This is the deployable-side counterpart of the starvation result of Section 7.8: the only quantity that cleanly predicts the sign of the gain is the measured ABF error ( $\rho = 0.80$ ), which is exactly what is unavailable online. The adaptive schedule is therefore a conservative first version: it reliably preserves the large starved-regime gains and reduces the harm in accurate cells relative to the fixed rate, but it does not eliminate that harm. We do not claim otherwise.

Table 11 compares the fixed and adaptive schedules on the representative cells (matched seeds, 10 seeds where complete). In the starved anchor  $\beta = 1, h = 2$  the adaptive schedule keeps essentially the full fixed gain (+49.9996% fixed versus +50.0491% adaptive), winning all seeds; the same holds for  $\beta = 1, h = 4$ . In the intermediate  $\beta = 2, h = 6$  cell the gain is also preserved ( $\approx +27\%$ ). The diagnostic easy cell  $\beta = 2, h = 4$  is the clearest success: the fixed rate is net-harmful there ( $-50\%$ , 0/10 wins) while the adaptive schedule is essentially neutral (within a few percent of ABF, 4/10 wins), i.e. it removes almost all of the harm. The hardest cases are the extreme easy cells  $\beta = 4, h = 1$  and  $\beta = 4, h = 2$ : the gate cuts the effective rate roughly five-fold (median  $\approx 0.02$  versus the fixed 0.10) and keeps the ancestor ESS high ( $\sim 320$ – $340/1024$ ), which roughly halves the harm ( $-365\% \rightarrow -164\%$  at  $\beta = 4, h = 1$ ), but does not remove it: when the ABF baseline error is already  $\approx 0.005$  there is essentially nothing to correct and any birth–death only adds variance. Two readings follow. First, the diversity and support gates do their job — the schedule never collapses genealogy and strongly throttles itself in accurate cells. Second, the irreducible residual harm in the most accurate cell confirms the conditional recommendation of Section 7.8: deploy mFR (fixed or adaptive) only when an independent starvation diagnostic, not an online marginal statistic, indicates the ABF estimator is genuinely under-sampled.




Figure 24: Transition support (kernel-weighted ABF denominator, normalised by the median-bin support) versus mFR gain. End-of-run support is balanced across regimes, so it is a weak online starvation signal on its own; the measured free-energy error remains the cleanest predictor.

Table 11: Fixed versus adaptive mFR on the representative cells (matched seeds, 10 seeds where complete). Columns: matched-seed median gain over ABF for the fixed and adaptive FR schedules, and the adaptive schedule’s median effective sample-size fraction and birth–death event fraction. The adaptive schedule preserves the starved-regime gains and reduces (without eliminating) the harm in the accurate cells.

| cell           | fixed mFR gain% | adaptive gain% | adaptive ESS frac | adaptive event frac |
|----------------|-----------------|----------------|-------------------|---------------------|
| $\beta=1, h=2$ | 50.0            | 50.0           | 0.160             | 0.0001              |
| $\beta=1, h=4$ | 49.1            | 50.4           | 0.186             | 0.0001              |
| $\beta=2, h=4$ | -50.1           | -2.1           | 0.315             | 0.0000              |
| $\beta=2, h=6$ | 28.0            | 24.7           | 0.413             | 0.0000              |
| $\beta=4, h=1$ | -365.2          | -164.2         | 0.335             | 0.0000              |
| $\beta=4, h=2$ | -249.8          | -123.9         | 0.313             | 0.0000              |

## 7.10 Equal-compute comparison

A natural objection is that mFR wins only because birth–death is extra compute. We control for this by comparing ABF and the deployable mFR at a fixed force-evaluation budget  $B = N \cdot n_{\text{steps}}$  (per physical  $M$ ), running ABF at three shapes of equal budget —  $(N, n_{\text{steps}}) = (1024, 120k), (2048, 60k)$  and  $(512, 240k)$  — plus a doubled-budget ABF point, on a starved, an easy and an intermediate cell (`scripts/run_wca_followup.py`, config `configs/wca_equal_compute.yaml`). Table 12 and Figure 27 report final and time-integrated  $L^2(F)$  against budget. The question this answers is whether, in the starved regime, simply spending ABF’s budget differently (more replicas or more steps) closes the gap to mFR at the base budget.

It does not. In the starved cell  $\beta = 1, h = 2$  the deployable mFR at the base budget ( $L^2(F) = 0.0412$ ) is far below ABF at *every* equal-budget shape (0.0832 at  $512 \times 240k$ , 0.0867 at  $1024 \times 120k$ , 0.1061 at  $2048 \times 60k$ ); indeed the *worst* mFR shape still beats the *best* ABF shape. ABF’s error has plateaued: even *doubling* ABF’s budget ( $2048 \times 120k$ ) leaves it near 0.09, more than twice the base-budget mFR error, because the starved-regime deficit is a systematic estimator-variance effect that






Figure 25: Fixed versus adaptive mFR on the representative cells (median and IQR over seeds). Left: per-cell final  $L^2(F)$  for ABF, fixed mFR and adaptive mFR. Right: matched-seed median gain over ABF; the adaptive schedule tracks the fixed gain in the starved cells and reduces (but does not fully remove) the harm in the accurate cells, where no online marginal signal separates the cold-sharp marginal from a starved one.

more raw sampling does not remove. (For ABF, longer trajectories help more than more replicas at fixed budget, but neither closes the gap.) The intermediate cell behaves the same way (mFR below ABF at all shapes); only in the already-accurate easy cell does ABF win at equal budget, where birth–death merely adds variance. The starved-regime advantage is therefore attributable to the birth–death correction, not to raw compute.

### 7.11 Serial one-walker ABF at equal force-evaluation budget

The equal-compute comparison above holds the budget fixed but keeps the method *parallel* ( $N$  replicas sharing one ABF estimate). A sharper fairness objection, raised in review, is that an  $N$ -replica ABF/mFR run at time  $T$  spends  $N \times$  the force evaluations of ordinary single-walker ABF at time  $T$ : perhaps a lone ABF walker would catch up if simply run  $N$  times longer. We test this directly with a *serial one-walker* control: an  $N_{\text{serial}} = 1$  ABF run, with its own accumulators and no birth–death, advanced along the force-evaluation-budget axis  $B = N \cdot n_{\text{steps}}$  (rather than physical simulation time, on which the serial walker runs far longer) toward the exact equal budget of the base parallel run,  $N_{\text{parallel}} \cdot n_{\text{steps}}^{\text{base}} = 1024 \times 120\text{k} = 122,880,000$ . Because a single walker is inherently serial, this is a Python-loop-bound multi-day computation at the exact budget; as a budget-conscious choice we run all three cells to the 1/4-budget cutoff (30,720,000 steps, `scripts/run_wca_serial_abf.py`, config `configs/wca_serial_abf_equal_budget.yaml`), which is enough to read the shape of the error-vs-budget curve; the exact  $N \cdot T$  endpoint is a resumable extension. Table 13 and Figure 28 report  $L^2(F)$  against budget, with the serial walker drawn as a curve over its accumulated budget so one can see whether it is plateauing at the ABF error floor or still catching up.

The parallel and serial runs have comparable force-evaluation budgets but are *not* dynamically equivalent: 1024 replicas explore from 1024 independent starting points in parallel, whereas one walker must diffuse through the whole reaction coordinate serially. The result actually favours the serial walker as an *ABF* estimator: in the starved cell, at the 1/4-budget cutoff ( $B = 3.07e + 07$ , run complete), serial one-walker ABF reaches  $L^2(F) = 0.0730$  — *below* the parallel-ABF floor (0.0867 at the full base budget, a +15.7% improvement, 9/10 seeds beating it), because a single






Figure 26: Adaptive FR schedule trajectories for a starved and an easy cell. The effective rate stays near its base value while the (exogenous, marginal-flatness) support gate is high in the starved cell, and is throttled by the support gate and warm-up in the already-accurate cell, with the diversity gate providing the genealogy safeguard.

long, well-mixed trajectory decorrelates better than many short replicas. Yet it remains a factor  $1.77\times$  *above* base-budget mFR (0.0412), and the error-vs-budget curve (Figure 28) has *flattened* at that ABF floor over the last decade of budget, with no sign of closing the gap. The intermediate cell corroborates (serial ABF  $\approx$  parallel ABF, both plateauing above mFR), while in the easy cell — where mFR is *harmful* — the serial walker plateaus at the low ABF floor (matching parallel ABF) and sits nearly an order of magnitude *below* mFR, exactly as the thesis predicts. In all three cells the one-walker curve has plateaued by the 1/4 cutoff, so the plateau, not a still-descending transient, is what the comparison reads. So even granting ordinary ABF the most favourable single-long-trajectory form of the same force-evaluation budget, it does not recover the starved-regime birth-death reallocation: the mFR advantage there is the correction, not raw compute. We do not overclaim the exact  $N\cdot T$  endpoint (the cells are run to the 1/4 cutoff, a resumable extension), but the  $\sim 1.8\times$  gap and its flat trend give no indication that a further  $4\times$  budget would close it. Either way the central WCA thesis stays conditional: mFR helps when ABF is sample-starved and is unnecessary or harmful when ABF is already accurate.

## 7.12 Frozen-bias validation

Finally we separate the quality of the *learned bias* from any online adaptive transient. In a two-stage protocol, stage 1 learns a bias  $B(z) = \hat{F}(z)$  online (for ABF, fixed mFR, and adaptive mFR); stage 2 freezes that bias — no ABF updates, no birth-death — and runs fresh dynamics with independent seeds, collecting the biased marginal  $p_B(z)$  and reconstructing

$$F_{\text{recon}}(z) = B(z) - \beta^{-1} \log p_B(z) + C, \quad (29)$$

with  $C$  fixed by centring on the evaluation window. No reference enters the dynamics; the TI profile is used only for the post-hoc  $L^2$  of  $F_{\text{recon}}$ . Table 14 reports, per learned bias, the online  $L^2(F)$ , the frozen reconstructed  $L^2(F)$ , and the learned-bias  $L^2(F)$ ; Figure 29 overlays the reconstructed profiles on the TI reference. The validation gives two results. First, the saved biases are faithful and the online ranking is real, not a transient artifact: the recomputed learned-bias  $L^2(F)$  matches the online value method-by-method, and the mFR-learned bias is genuinely about twice

Table 12: Equal-compute comparison at fixed force-evaluation budget  $B = N \cdot n_{\text{steps}}$ . For a starved ( $\beta=1, h=2$ ), an intermediate ( $\beta=2, h=6$ ) and an easy ( $\beta=4, h=1$ ) cell, ABF is run at three equal-budget shapes (plus a doubled-budget point) and compared with deployable mFR at the base budget by final  $L^2(F)$  (median over seeds). In the starved and intermediate cells the worst mFR shape still beats the best ABF shape; only in the easy cell does ABF win.

| cell           | method       | $N$  | $n_{\text{steps}}$ | budget | $L^2(F)$ |
|----------------|--------------|------|--------------------|--------|----------|
| $\beta=1, h=2$ | abf          | 512  | 240000             | 1e+08  | 0.0832   |
|                | abf          | 1024 | 120000             | 1e+08  | 0.0867   |
|                | abf          | 2048 | 60000              | 1e+08  | 0.1061   |
|                | abf          | 2048 | 120000             | 2e+08  | 0.0919   |
|                | fr-estimated | 512  | 240000             | 1e+08  | 0.0364   |
|                | fr-estimated | 1024 | 120000             | 1e+08  | 0.0412   |
|                | fr-estimated | 2048 | 60000              | 1e+08  | 0.0586   |
|                | fr-estimated | 2048 | 120000             | 2e+08  | 0.0586   |
| $\beta=2, h=6$ | abf          | 512  | 240000             | 1e+08  | 0.0249   |
|                | abf          | 1024 | 120000             | 1e+08  | 0.0279   |
|                | abf          | 2048 | 60000              | 1e+08  | 0.0269   |
|                | abf          | 2048 | 120000             | 2e+08  | 0.0258   |
|                | fr-estimated | 512  | 240000             | 1e+08  | 0.0175   |
|                | fr-estimated | 1024 | 120000             | 1e+08  | 0.0190   |
|                | fr-estimated | 2048 | 60000              | 1e+08  | 0.0212   |
|                | fr-estimated | 2048 | 120000             | 2e+08  | 0.0212   |
| $\beta=4, h=1$ | abf          | 512  | 240000             | 1e+08  | 0.0041   |
|                | abf          | 1024 | 120000             | 1e+08  | 0.0046   |
|                | abf          | 2048 | 60000              | 1e+08  | 0.0112   |
|                | abf          | 2048 | 120000             | 2e+08  | 0.0048   |
|                | fr-estimated | 512  | 240000             | 1e+08  | 0.0172   |
|                | fr-estimated | 1024 | 120000             | 1e+08  | 0.0232   |
|                | fr-estimated | 2048 | 60000              | 1e+08  | 0.0359   |
|                | fr-estimated | 2048 | 120000             | 2e+08  | 0.0359   |

as accurate as the ABF-learned bias ( $\approx 0.044$  versus  $0.088$  at  $\beta = 1, h = 2$ ). Second, the single-window reconstruction itself is mixing-limited at this budget: the reconstructed  $L^2(F)$  ( $\approx 0.2122$  at  $\beta = 1, h = 2$ ) is dominated by the residual difficulty of equilibrating the  $z$ -marginal under an imperfect frozen bias rather than by the bias quality, so it does not finely separate the methods (the error grows from the transition outward into the basins, not from any low-density blow-up). The deployable conclusion — that the mFR-learned bias is genuinely better — rests on the first result; the reconstruction is reported as an honest, and limited, independent check.

### 7.13 Comparison to a modern enhanced-sampling baseline (OPES)

A natural objection is that the comparison so far pits marginal Fisher–Rao (mFR) only against plain ABF, and that a modern adaptive-bias method might already close the starved-regime gap on its own. We therefore add *on-the-fly probability enhanced sampling* (OPES, the well-tempered METAD variant of Invernizzi and Parrinello [9]) as a third baseline, run on the identical WCA dynamics, grid and seeds. To keep the comparison about the *biasing strategy* rather than the estimator, OPES free energies are reconstructed by the *same* mean-force integration used for ABF and mFR; the native OPES reweighting estimate is reported alongside it in the supplement and is *not* uniformly worse (it is comparable in the starved anchor and better in  $\beta=1, h=4$ ), so we report both rather than selecting the more favourable one. The deposition barrier is tuned on held-out seeds {20, 21} and fixed to 4 kT for all production seeds 0–9. *Preliminary:* in this pilot the kernel




Figure 27: Equal-compute comparison. Final (left) and time-integrated (right)  $L^2(F)$  versus force-evaluation budget  $N \cdot n_{\text{steps}}$  for ABF (varying the  $N/n_{\text{steps}}$  split at fixed budget) and the deployable mFR, on representative cells.

bandwidth  $\sigma$  and deposition pace were held fixed rather than tuned; a closure study sweeping  $\sigma \times \text{pace} \times \text{barrier}$  on held-out seeds (Section 7.13) tunes them per cell and supersedes the single- $\sigma$  numbers below (the tuned  $\sigma$  is well below this fixed value in every cell).

Across all six representative cells OPES samples aggressively — up to  $\sim 6 \times 10^5$  barrier round trips and a reweighting  $n_{\text{eff}}$  of 0.74 in the starved cell — yet its *free-energy accuracy* does not beat ABF in any cell, and does not beat mFR in five of six (Table 15). In the starved anchor  $\beta=1, h=2$  OPES reaches  $L^2(F)=0.1333$  versus ABF’s 0.0878 and mFR’s 0.0431; mFR is roughly a factor of three more accurate than the OPES baseline at matched compute. Only in the easy cell  $\beta=4, h=1$ , where every method is already near the floor, does OPES edge mFR (+9.7%). The flat-target ablation ( $\gamma \rightarrow \infty$ , no well-tempering) is uniformly worse (e.g.  $L^2(F)=0.201$  vs 0.0856’s well-tempered counterpart in  $\beta=4, h=2$ ), confirming that the shipped variant is the favourable one along the *well-tempering* axis. We caution that this ablation does not by itself rule out a bandwidth-tuning artefact:  $\sigma$  was fixed across all cells, and the closure study (Section 7.13) tunes it directly before drawing a final conclusion about the size of the mFR advantage over OPES.

The reading is that OPES’s strength — rapid, well-tempered flattening of the bias — buys sampling throughput, not directed reallocation along the mean force. Where ABF is starved, the residual error is dominated by how the local mean force is *estimated* from a finite ensemble, and it is exactly there that mFR’s birth–death reallocation helps and OPES does not. The mFR advantage documented in this section therefore holds against a modern enhanced-sampling method, not merely against textbook ABF.

## Closure study

The pilot above fixed the OPES kernel bandwidth  $\sigma$  and deposition pace and tuned only the barrier, so its single- $\sigma$  numbers in Table 15 are a *preliminary* lower bound on OPES accuracy rather than the method’s tuned best. A closure study resolves this: it sweeps  $\sigma \times \text{pace} \times \text{barrier}$  by successive halving on the disjoint tuning seeds  $\{20, 21\}$  across all three systems (WCA, metastability, entropic bottleneck), freezes the winning configuration *per cell*, and reruns at production budget (10 seeds, 120000 steps for WCA). The headline is that  $\sigma$ —the axis the pilot never swept—is first-order everywhere, and the tuned optimum sits well below the pilot’s fixed  $\sigma$  in every system: WCA

Table 13: Serial one-walker ABF at equal force-evaluation budget versus parallel ABF/mFR. “gain” is the median  $L^2(F)$  improvement over the base parallel ABF ( $N=1024$ , 120k); the serial win count is the number of serial seeds beating that base-ABF median. Budget =  $N \cdot n_{\text{steps}}$ .

| cell                        | method               | $N$  | $n_{\text{steps}}$ | budget   | $L^2(F)$ | $L^2(F')$ | gain%  | wins |
|-----------------------------|----------------------|------|--------------------|----------|----------|-----------|--------|------|
| starved $\beta=1, h=2$      | serial ABF ( $N=1$ ) | 1    | 30,720,000         | 3.07e+07 | 0.0730   | 0.2966    | 15.7   | 9/10 |
|                             | parallel ABF         | 2048 | 60,000             | 1.23e+08 | 0.1061   | 0.3722    | -22.4  | –    |
|                             | parallel ABF         | 1024 | 120,000            | 1.23e+08 | 0.0867   | 0.3254    | 0.0    | –    |
|                             | parallel ABF         | 2048 | 120,000            | 2.46e+08 | 0.0919   | 0.3404    | -6.0   | –    |
|                             | parallel ABF         | 512  | 240,000            | 1.23e+08 | 0.0832   | 0.3154    | 4.0    | –    |
|                             | parallel mFR         | 2048 | 60,000             | 1.23e+08 | 0.0586   | 0.2857    | 32.4   | –    |
|                             | parallel mFR         | 1024 | 120,000            | 1.23e+08 | 0.0412   | 0.2522    | 52.5   | –    |
|                             | parallel mFR         | 512  | 240,000            | 1.23e+08 | 0.0364   | 0.2527    | 58.0   | –    |
| intermediate $\beta=2, h=6$ | serial ABF ( $N=1$ ) | 1    | 30,720,000         | 3.07e+07 | 0.0234   | 0.0995    | 16.1   | 5/5  |
|                             | parallel ABF         | 2048 | 60,000             | 1.23e+08 | 0.0269   | 0.1861    | 3.5    | –    |
|                             | parallel ABF         | 1024 | 120,000            | 1.23e+08 | 0.0279   | 0.1660    | 0.0    | –    |
|                             | parallel ABF         | 2048 | 120,000            | 2.46e+08 | 0.0258   | 0.1625    | 7.3    | –    |
|                             | parallel ABF         | 512  | 240,000            | 1.23e+08 | 0.0249   | 0.1424    | 10.7   | –    |
|                             | parallel mFR         | 2048 | 60,000             | 1.23e+08 | 0.0212   | 0.1793    | 23.9   | –    |
|                             | parallel mFR         | 1024 | 120,000            | 1.23e+08 | 0.0190   | 0.1497    | 31.7   | –    |
|                             | parallel mFR         | 512  | 240,000            | 1.23e+08 | 0.0175   | 0.1296    | 37.0   | –    |
| easy $\beta=4, h=1$         | serial ABF ( $N=1$ ) | 1    | 30,720,000         | 3.07e+07 | 0.0041   | 0.0296    | 10.7   | 3/5  |
|                             | parallel ABF         | 2048 | 60,000             | 1.23e+08 | 0.0112   | 0.0378    | -145.5 | –    |
|                             | parallel ABF         | 1024 | 120,000            | 1.23e+08 | 0.0046   | 0.0263    | 0.0    | –    |
|                             | parallel ABF         | 2048 | 120,000            | 2.46e+08 | 0.0048   | 0.0249    | -6.1   | –    |
|                             | parallel ABF         | 512  | 240,000            | 1.23e+08 | 0.0041   | 0.0255    | 10.7   | –    |
|                             | parallel mFR         | 2048 | 60,000             | 1.23e+08 | 0.0359   | 0.0994    | -687.8 | –    |
|                             | parallel mFR         | 1024 | 120,000            | 1.23e+08 | 0.0232   | 0.0650    | -409.6 | –    |
|                             | parallel mFR         | 512  | 240,000            | 1.23e+08 | 0.0172   | 0.0493    | -276.2 | –    |

per-cell  $\sigma \in \{0.035, 0.02, 0.05\}$ , entropic bottleneck  $\sigma \approx 0.02$ , metastability  $\sigma \approx 0.045$  (all vs a pilot  $\sigma$  of 0.05–0.12).

On the WCA anchor cells the closure also separates the well-tempering target estimator into three labelled variants—adaptive- $\gamma$ , fixed- $\gamma$ , and flat (no tempering)—rather than collapsing them. Well-tempering is first-order: the flat target is worst in every cell (e.g. intermediate  $L^2(F) = 0.1521$  vs adaptive 0.0892). Adaptive- and fixed- $\gamma$  are a near-tie (starved  $0.1276 \pm 0.0029$  vs  $0.1295 \pm 0.0032$ ), so the adaptive schedule neither helps nor hurts at these operating points. On the toys the effect of tuning  $\sigma$  alone is large and cleanly isolated: for metastability, holding barrier/pace/budget/seeds fixed and changing only  $\sigma$  from the pilot 0.12 to the tuned 0.045 improves  $L^2(F)$  from 0.1169 to 0.0571 ( $2.05\times$ , 51%); the entropic bottleneck tuned point reaches  $L^2(F) = 0.1839$ . The qualitative pilot ordering (mFR ahead in the starved regime) is unchanged; the closure sharpens the OPES numbers and confirms the single- $\sigma$  pilot understated OPES accuracy.

## 7.14 Case finding

At the original condensed-phase WCA anchor, ABF+FR with an online estimated target reliably and substantially improves free-energy accuracy: +55.0% in median final  $L^2(F)$  at the tuned operating point, winning 10/10 seeds, with the gain concentrated in the free energy (not the mean force). The mechanism is directed coverage equalisation after warm-up — marginal reallocation,




Figure 28: Serial one-walker ABF (green curve, median and IQR over seeds) versus parallel ABF (three  $N/n_{\text{steps}}$  shapes) and the deployable mFR at equal force-evaluation budget  $B = N \cdot n_{\text{steps}}$ , for the starved, intermediate and easy cells. The dotted line marks the base equal budget  $1024 \times 120\text{k}$ . Serial ABF is drawn along its accumulated budget so its plateau (or catch-up) is visible; the parallel methods are points at their budgets.

Table 14: Frozen-bias reconstruction (29). For each learned bias (ABF, mFR, adaptive mFR) on a starved ( $\beta=1, h=2$ ) and an easy ( $\beta=4, h=1$ ) cell: the online  $L^2(F)$ , the frozen reconstructed  $L^2(F)$  from fresh unbiased dynamics, and the recomputed learned-bias  $L^2(F)$  (median over seeds). The recomputed learned-bias error matches the online value method-by-method, confirming the online ranking; the single-window reconstruction is mixing-limited and does not finely separate the methods.

| cell           | learned bias          | online $L^2(F)$ | frozen recon $L^2(F)$ | learned-bias $L^2(F)$ |
|----------------|-----------------------|-----------------|-----------------------|-----------------------|
| $\beta=1, h=2$ | abf                   | 0.0878          | 0.2190                | 0.0878                |
|                | fr-estimated          | 0.0431          | 0.2122                | 0.0442                |
|                | fr-estimated-adaptive | 0.0430          | 0.2168                | 0.0432                |
| $\beta=4, h=1$ | abf                   | 0.0053          | 0.0177                | 0.0052                |
|                | fr-estimated          | 0.0242          | 0.0188                | 0.0244                |
|                | fr-estimated-adaptive | 0.0136          | 0.0169                | 0.0134                |

which is all Proposition 2 permits it to be; at the anchor, target shape matters little (uniform  $\approx$  estimated  $\approx$  oracle, and the estimated target already reaches the oracle’s accuracy). The method has a clear safe regime,  $\text{fr\_rate} \approx 0.1\text{--}0.2$ , and a real failure boundary: too-aggressive birth–death reverses the benefit as diversity collapses. The new phase diagram qualifies the claim: mFR is most useful where the ABF estimator is visibly sample-starved (the  $\beta = 1$  WCA row in the production phase diagram), but when ABF already has a low free-energy error the same birth–death correction can be unnecessary or harmful. Two budget-matched controls guard against the objection that mFR simply buys more sampling: at equal force-evaluation budget mFR still beats every parallel-ABF shape in the starved cell (Section 7.10), and a serial one-walker ABF control run at the exact  $N \cdot T$  force-evaluation budget (Section 7.11) tests whether a lone ABF walker catches up; on the evidence available the starved-regime advantage tracks the directed reallocation rather than raw compute. Thus the practical recommendation is conditional, not automatic: use gentle, diversity-preserving mFR in starved WCA regimes, and monitor ABF error, transition support, FR event fractions, and genealogy before increasing the birth–death strength.




Figure 29: Frozen-bias reconstruction (29) from the learned ABF, fixed-mFR and adaptive-mFR biases (median and IQR over independent frozen seeds) against the TI reference, for a starved and an easy cell.

Table 15: WCA dimer, median final  $L^2(F)$  over production seeds 0–9: well-tempered OPES versus ABF and marginal Fisher–Rao (mFR), all reconstructed by the same mean-force integration at matched compute. OPES is less accurate than ABF in every cell and less accurate than mFR in five of six; the sole exception is the easy cell, where all methods are near the floor. Percentages are the OPES error change relative to each baseline (negative = worse).

| Cell                          | OPES   | ABF    | mFR    | OPES vs ABF | OPES vs mFR |
|-------------------------------|--------|--------|--------|-------------|-------------|
| $\beta=1, h=2$ (starved)      | 0.1333 | 0.0878 | 0.0431 | -51.8%      | -209.7%     |
| $\beta=2, h=6$ (intermediate) | 0.1099 | 0.0265 | 0.0184 | -315.2%     | -498.8%     |
| $\beta=4, h=1$ (easy)         | 0.0218 | 0.0053 | 0.0242 | -311.4%     | +9.7%       |

### 7.15 Does the directional mechanism survive a matched sham? A many-body test

Everything above compares mFR against ABF and against other *score-driven* arms (`fr_uniform`, `fr_oracle`), each of which derives its replacement direction from the same Fisher–Rao score and differs only in how the target is built. So the WCA positive, however reproducible, was a positive against ABF and never against *turnover of the same intensity*. Section 15 supplies that control in a two-dimensional model; this subsection supplies it here, in the many-body system.

Nothing is retuned. The cell is the accepted `b1_h2` configuration ( $\beta = 1$ ,  $h = 2$ ,  $w = 2$ ,  $M = 100$ ,  $a = 1.5$ ), the budget is 120 000 steps at  $N = 1024$ , and every Fisher–Rao knob is read back from the production YAML and asserted, so an edit to that file cannot move this experiment. Only the seeds are new (16 of them, verified unused anywhere in the study). Because the WCA sampler runs one method per process, a sham cannot observe its partner online: it replays the per-event replacement count sequence its partner recorded *on the same seed*, choosing which replicas die and which are copied uniformly at random. Across all seeds and both partners the realised totals matched with 0 discrepancies.

Table 16 gives the result, and both preregistered criteria pass. Practical mFR improves the integrated free-energy error by -17.97%  $[-22.73, -17.16]$  on 16/16 matched seeds. The attribution

Table 16: Matched-sham control at the accepted WCA-positive cell (b1\_h2,  $M = 100$  particles, 120 000 steps,  $N = 1024$ , FR rate 0.10) on 16 fresh seeds, with no hyperparameter reopened. Percentages are paired differences in the integrated free-energy error against ABF on the same seed; **negative is better**. Each sham replays its own partner’s per-event replacement counts on that seed and differs only in randomising which replicas die and which are copied (0 count mismatches across all seeds and both partners). The *direct* rows are the preregistered attribution test. **Quote the corrected-reference column:** the cached reference the v1 column was scored against is wrong near  $z \approx 0.26$  (Remark 7). The v1 column is retained because the run happened and its criteria were assessed against it, not because it is still current.

| comparison                                                          | corrected reference (phase_hp_v3) |                    |           | cached ref.       |
|---------------------------------------------------------------------|-----------------------------------|--------------------|-----------|-------------------|
|                                                                     | $\Delta I_F$                      | 95% CI             | seeds won | $\Delta I_F$ (v1) |
| <i>against ABF</i>                                                  |                                   |                    |           |                   |
| practical mFR                                                       | <b>-17.97%</b>                    | $[-22.73, -17.16]$ | 16/16     | -22.83%           |
| sham (practical)                                                    | -3.19%                            | $[-6.20, +2.89]$   | —         | +2.60%            |
| oracle mFR                                                          | -19.54%                           | $[-25.45, -15.98]$ | —         | -17.46%           |
| <i>direct: each arm against its own sham — the attribution test</i> |                                   |                    |           |                   |
| practical mFR vs its sham                                           | <b>-20.96%</b>                    | $[-25.04, -11.54]$ | 15/16     | -26.38%           |
| oracle mFR vs its sham                                              | —                                 | —                  | —         | -21.08%           |

Median integrated  $L^2(F)$  against the corrected reference: ABF 42.21, practical mFR 33.05, oracle mFR 32.93, sham 41.43. Health at the practical arm: min ancestor ESS/ $N = 0.170$  and max ancestor fraction 0.019, against preregistered floors of 0.10 and 0.05. The looser ESS floor than the gateway’s 0.30 was declared in advance, because the accepted WCA configuration already runs at ESS/ $N \approx 0.17$  and the gateway’s floor would have failed the configuration under test rather than tested it. The oracle-versus-its-own-sham contrast was not recomputed against the corrected reference: the practical arm carries the preregistered attribution claim, and the oracle is not a deployable method.

criterion — the *direct* contrast of each arm against its own sham, which holds the event schedule and the replacement count fixed so that only the selection direction differs — gives -20.96%  $[-25.04, -11.54]$  on 15/16 seeds, with the interval excluding zero. Verdict: **TRANSFERS** (Figure 30).

These are the corrected-reference figures. The experiment was first scored against a cached thermodynamic-integration reference and read -22.83% and -26.38%; that reference was later found wrong inside the transition region, and Remark 7 gives the correction and what it did and did not change.

*Remark 5* (The sham is not shown equivalent to ABF, and the side it misses on changed). The equivalence test was written expecting the sham to land *on* ABF, and against the corrected reference it still does not, though only just: matched random turnover reads -3.19%  $[-6.20, +2.89]$  on 9/16 seeds, so it is not resolvably different from ABF, but its TOST interval  $[-5.60, +0.80]$  runs past the  $\pm 5\%$  margin and it is therefore not shown equivalent.

The side it misses on is worth stating plainly, because it reversed under the reference correction and the two sides support opposite readings. Against the cached reference the sham read +2.60% — resolvably *harmful* — and that failure strengthened attribution, since a control that makes the estimate worse cannot explain an arm that makes it better. Against the corrected reference the miss is on the *beneficial* side, which is the case the preregistration did anticipate and the less comfortable one: it leaves open that a small part of the gain is generic turnover rather than direction. What closes that gap is not this test but the direct one. The preregistered attribution statistic contrasts




Figure 30: The matched-sham control on the WCA dimer, 16 fresh seeds at the accepted cell. **(a)** Paired per-seed change in the integrated free-energy error against ABF on the same seed; bars are medians. Both mFR arms sit below zero, both shams above it. **(b)** The attribution test: each arm against *its own* sham, which fired the same number of clone and delete events at the same times and differs only in which replicas it acted on. **(c)** Round trips per run differ by about 1% across arms while the free energy differs by more than seventeen — the mechanism redistributes population rather than creating barrier crossings. Panel values are the cached-reference scoring the figure was drawn from; the corrected-reference numbers are Table 16’s and move the accuracy bars without changing any sign.

each arm against its *own* sham at an identical event schedule, and it gives  $-20.96\%$   $[-25.04, -11.54]$  with the interval excluding zero — larger in magnitude than the arm-versus-ABF gain itself, so the direction of the mass movement, not the fact of it, is what carries the effect.

Two cautions apply to reading this remark at all. The equivalence criterion is failed in both scorings, so neither polarity is a predeclared outcome; and the sham’s shift from  $+2.60\%$  to  $-3.19\%$  is itself evidence that a control this small is sensitive to reference error at the scale the audit found.

*Remark 6* (mFR is not buying barrier crossings). Round trips per run are nearly identical across the arms: 780858 for ABF against 790296 for practical mFR. Paired on the same seeds, practical mFR changes barrier crossings by  $+1.09\%$   $[+0.77, +1.50]$  while changing the integrated free-energy error by  $-17.97\%$  — a separation of roughly sixteen-fold between the transport change and the accuracy change. Both figures are measured on the same runs against the same corrected reference, which matters here: pairing a transport number from one run tree with an accuracy number from another would make the ratio an artifact of the pairing. The crossing change is small but resolvably nonzero, so the claim is that transport is nearly unchanged, not that it is untouched. So the mechanism is not acting by increasing transport across the barrier; it is redistributing population among replicas that were already crossing. That is exactly what the two-state reduction of Section 15.7 predicts — the reallocation term changes the relaxation rate toward the target without touching the physical crossing rate  $k_{01}$  — and it is an independent confirmation of that model on a system it was not built from.

*Remark 7* (The reference this subsection was first scored against was wrong). This subsection was first scored against a cached thermodynamic-integration reference. A parallel audit later found that reference wrong near  $z \approx 0.26$  — by up to  $67\sigma$  against a high-precision consensus built from four independently prepared solvent pools — which is inside the transition region where the arms differ most. The cached acquisition grid ( $\Delta z = 0.028$ ) is too coarse to represent a feature with structure at 0.010, and the smoothing kernel took its width in *grid cells* rather than in the physical



coordinate, so the two errors compounded: the feature was first unresolved and then smoothed away.

Correcting it required **fresh dynamics, not a rescore**. The original Case IX runs retained the scored scalars and the final profile but never the time-resolved  $\hat{F}_t(z)$ , so the integrated end-point could not be recomputed from what was stored. The cell was therefore re-run in full at the frozen configuration with nothing retuned and 16 paired seeds, scored online against the corrected reference (`cache/phase_hp_v3`: unsmoothed, four preparations, acquisition grid equal to the evaluation grid so interpolation is the identity). That moves the headline from  $-22.83\%$  to  $-17.97\%$   $[-22.73, -17.16]$ , the direct arm-versus-sham contrast from  $-26.38\%$  to  $-20.96\%$   $[-25.04, -11.54]$ , and the sham from  $+2.60\%$  to  $-3.19\%$   $[-6.20, +2.89]$ . Because the seeds are new, the comparison between the two columns of Table 16 carries seed noise as well as the reference change, and the split between them is not identified.

Three things are worth separating. The verdict is unchanged: every preregistered criterion that passed against the cached reference passes against the corrected one, at 16/16 seeds with the attribution interval still excluding zero. The magnitude moved by about five percentage points, which is comparable to effects this report elsewhere treats as meaningful, so the corrected value is the one quoted everywhere and the cached one is retained only as a record. And the sham’s move is not the unambiguous improvement it first appears: it removes the harmful-side reading of Remark 5, but replaces it with a beneficial-side one, which is the direction that would dilute rather than reinforce attribution. The general rule in Section 18.2 — that no arm-level effect is trustworthy unless the reference uncertainty is much smaller than the effect — was adopted because of this episode rather than in anticipation of it. The concrete repair is that every run now stores  $\hat{F}_t(z)$  alongside the reference it was scored against, so a future reference change costs an analysis pass rather than a GPU campaign.

## 7.16 Is the Fisher–Rao direction better than the directed selection ABF already had?

Every control so far contrasts mFR with ABF or with *undirected* turnover. Neither asks the question a reader of the birth–death literature would ask first: selection along a reaction coordinate is not new, so is the Fisher–Rao *direction* better than the directions already in use? Proposition 2 sharpens what is at stake. All three rules below are functions of the collective variable alone, so all three leave every conditional exactly invariant and can differ *only* in how they redistribute the marginal. The question is whether that difference is worth anything.

Five arms are run at the accepted cell on 16 fresh paired seeds: `abf`; `fr_estimated`, frozen exactly as Section 7.15 ran it; its matched sham; the Chapter 6 rule  $S = c \partial_z^2 p / p$  (`book_laplacian`); and the count-balancing rule  $S = c(1 - \hat{p}/\bar{p})$  (`count_balancing`), the simple alternative that multiple-walker implementations actually ship.

**Turnover is matched, and matching it was the hard part.** An arm must not win by selecting harder, so each prior-art intensity  $c$  was chosen on *held-out* seeds to bring its replacement count closest to `fr_estimated`’s on the same cell; the mFR configuration was not retuned. This is not a formality. `book_laplacian`’s score is pinned at the clip  $S_{\max}$  over the whole intensity range one would naturally search: a  $30\times$  change in  $c$  moved its turnover by  $1.44\times$ , so the obvious ladder bottomed out at  $5.5\times$  the target and would have handed that arm a five-fold turnover advantage. Matching it requires  $c = 0.012$ , two decades below where the search started. The final intensities sit at  $0.96\times$  and  $1.16\times$  of mFR’s own turnover.

Table 17: Q1 — the Fisher–Rao direction against the directed selection rules already in use, at matched turnover on 16 paired seeds. Endpoint is the time-integrated  $L^2$  free-energy error, scored against the corrected reference; negative favours the first arm. The two rows below the rule state why the novelty claim fails: count balancing recovers the whole gain on its own, and the Chapter 6 rule recovers none of it.

| Arm                                                          | Comparator                   | Median (%) | 95% CI           | Wins  |
|--------------------------------------------------------------|------------------------------|------------|------------------|-------|
| practical mFR                                                | ABF                          | −17.62     | [−23.60, −11.61] | 15/16 |
| practical mFR                                                | its own matched sham         | −17.89     | [−21.22, −12.13] | 16/16 |
| practical mFR                                                | <code>book_laplacian</code>  | −21.66     | [−25.38, −15.14] | 15/16 |
| practical mFR                                                | <code>count_balancing</code> | +0.66      | [−5.18, +4.36]   | 8/16  |
| <i>Context: what each prior-art rule is worth on its own</i> |                              |            |                  |       |
| <code>book_laplacian</code>                                  | ABF                          | +0.73      | [−1.85, +8.22]   | 7/16  |
| <code>count_balancing</code>                                 | ABF                          | −18.23     | [−21.94, −12.90] | 15/16 |

Equivalence against `count_balancing` is established by TOST: the 90% interval [−3.96, +2.55] lies strictly inside the  $\pm 5\%$  margin. Selection intensities were matched on held-out seeds to  $0.96\times$  and  $1.16\times$  of mFR’s own turnover; the mFR configuration was not retuned.

**The primary result replicates.** Against the corrected reference and with the arms freshly run rather than rescored, mFR improves the integrated error by −17.62% [−23.60, −11.61] on 15/16 seeds and beats its own sham by −17.89% [−21.22, −12.13] on 16/16, while the sham is not better than ABF (−1.59%). This agrees with the rescored Section 7.15 value of −17.97% to within a fraction of a point, from independent trajectories.

**The novelty claim does not survive.** Against `book_laplacian` mFR is clearly better, −21.66% [−25.38, −15.14] on 15/16 seeds. But the context row is what decides the question: `book_laplacian` does not beat plain ABF either (+0.73%, 7/16), so mFR’s margin over it is not evidence of beating *directed* selection — that rule simply does not function on this system, and beating an inert arm is not a novelty result.

Against `count_balancing` mFR is **tied**: +0.66% with a TOST interval of [−3.96, +2.55], strictly inside the  $\pm 5\%$  equivalence margin, on 8/16 seeds. Equivalence here is *demonstrated*, not inferred from an interval that happens to contain zero. And count balancing recovers the whole effect on its own: −18.23% [−21.94, −12.90] against ABF on 15/16 seeds, statistically indistinguishable from mFR’s −17.62%.

The honest reading is that the establishment-limited gain is real, replicated and attributable — it beats its own sham 16/16 — but that the Fisher–Rao geometry is not what delivers it. A one-line count-balancing rule gets the same gain at matched turnover. What this report can claim is a principled selection rule consistent with the establishment mechanism and a diagnosis of when selection helps at all; it cannot claim a performance win over prior art, and Section 18 states the consequence for how the method should be positioned.

This outcome is what Proposition 2 would lead one to expect. The three rules are in one equivalence class as far as conditionals are concerned, and they can be separated only by the marginal they steer toward. Two of them — mFR’s bias-aware target and count balancing’s flat one — evidently steer toward marginals that are close enough to make no measurable difference at this budget, while the Laplacian rule does not steer usefully at all.

*Remark 8* (The genealogy gate was being read off the wrong statistic). The preregistered health gate  $\text{ESS}_{\text{anc}}/N \geq 0.30$  had never been evaluated on this system, and evaluating it exposed a

measurement error worth recording. Read off the ancestry the sampler already tracked, mFR scores 0.178 — an apparent failure. But the matched-turnover *random* sham scores 0.192, no better, and a sham is non-pathological by construction. A gate its own null control fails is not measuring what it claims.

The cause is definitional. That ancestry is traced from  $t = 0$  and never reset, so its ESS decays monotonically for *any* birth–death process, healthy or not, and a fixed floor on it is a cap on run length rather than a health criterion. The gateway confirmatory that set the 0.30 floor measured ESS over a rolling 4000-step window and read 0.37–0.73. On the windowed statistic mFR reads 0.807, and the takeover gate  $w_{\max} \leq 0.05$  passes at 0.0176. Both statistics are now recorded so they cannot be conflated again. The general lesson — check what the null control scores before believing a gate failure — is the same one Remark 5 arrived at from the other side.

## 8 Molecular torsion benchmarks: united-atom alkanes

The first three cases use one- or two-dimensional model potentials. To test the Fisher–Rao correction on a genuinely molecular reaction coordinate we add two united-atom alkanes whose slow coordinate is a signed dihedral: **butane** (one dihedral) and **pentane** (two coupled dihedrals). Both have essentially exact, independently constructed references, so the same “does mFR help ABF, and why” question can be asked where the collective variable is a nonlinear function of a three-dimensional molecular configuration.

### 8.1 United-atom model

Each  $\text{CH}_2/\text{CH}_3$  group is a single site in  $\mathbb{R}^3$ ; a dihedral is degenerate in 2D, so the molecules are three dimensional. Butane is  $\text{CH}_3\text{--CH}_2\text{--CH}_2\text{--CH}_3$  with sites  $q_1, \dots, q_4$  and one signed dihedral  $\varphi_1 = \varphi(q_1, q_2, q_3, q_4)$ ; pentane is  $\text{CH}_3\text{--CH}_2\text{--CH}_2\text{--CH}_2\text{--CH}_3$  with sites  $q_1, \dots, q_5$  and two dihedrals  $\varphi_1 = \varphi(q_1, q_2, q_3, q_4)$  and  $\varphi_2 = \varphi(q_2, q_3, q_4, q_5)$ . The potential, in reduced units following [2, 11], is

$$V = \underbrace{\frac{k_0}{2} \sum_{\text{bonds}} (r - d_0)^2}_{V_{\text{bond}}} + \underbrace{\frac{k_\theta}{2} \sum_{\text{angles}} (\theta - \theta_0)^2}_{V_{\text{angle}}} + \underbrace{\sum_{\text{dih.}} V_4(\varphi)}_{V_{\text{tors}}} + \underbrace{\sum_{\text{nb pairs}} 4\epsilon[(\sigma/r)^{12} - (\sigma/r)^6]}_{V_{\text{nb}}}, \quad (30)$$

with  $d_0 = 1$ ,  $k_0 = 1000$ ,  $\theta_0 = 1.187$  (the bend between successive bond vectors; the interior C–C–C angle is  $\pi - \theta_0 = 112^\circ$ ),  $k_\theta = 208$ . The Ryckaert–Bellemans torsion [13]

$$V_4(\varphi) = c_1(1 - \cos \varphi) + 2c_2(1 - \cos^2 \varphi) + c_3(1 + 3 \cos \varphi - 4 \cos^3 \varphi), \quad (31)$$

with  $c_1 = 1.18$ ,  $c_2 = -0.23$ ,  $c_3 = 2.64$ , has (in the polymer/trans=0 convention we use) its global minimum at the *trans* state  $\varphi = 0$  ( $V_4 = 0$ ), local minima at the *gauche* states  $\varphi = \pm 116.6^\circ$  ( $V_4 = 1.38$ ), a trans–gauche barrier of  $5.53 k_B T$  at  $\pm 61.6^\circ$  and a cis barrier of  $7.64 k_B T$  at  $\pm 180^\circ$  (all at  $\beta = 1$ ). Thus  $\varphi$  is genuinely metastable at  $\beta = 1$ .

**Nonbonded convention (documented and tested).** We use the Ryckaert–Bellemans exclusion rule: every pair separated by  $\leq 3$  bonds (1–2, 1–3, 1–4) is excluded, and Lennard–Jones acts only between sites  $\geq 4$  bonds apart. Consequently **butane has no LJ pair**, so its dihedral free energy is *exactly*  $F(\varphi_1) = V_4(\varphi_1) + C$  (the bond/angle internal-coordinate Jacobian is  $\varphi$ -independent); butane is the clean easy control. **Pentane has a single 1–5 pair**; this “pentane effect” couples  $\varphi_1$  and  $\varphi_2$  and makes  $\varphi_2$  a hidden slow coordinate when only  $\varphi_1$  is biased. We resolve the LJ




Figure 31: Molecular models and the torsional potential. Trans is  $\varphi = 0$  ( $V_4 = 0$ ), gauche  $\pm 116.6^\circ$  ( $1.38 k_B T$ ), trans-gauche barrier  $5.53 k_B T$ , cis barrier  $7.64 k_B T$ .

range from the physical united-atom ratio  $\sigma/d_0 = \sigma_{\text{CH}_2}/d_{\text{CC}} \approx 2.5$  and set  $\sigma = 2.3$ ,  $\epsilon = 1$ , which disfavors the  $g^+g^-/g^-g^+$  basins by  $\sim 8 k_B T$  while barely shifting the  $\varphi_1$  marginal (a  $\sigma$  sensitivity scan is reported in the supplement). The molecules are isolated (no solvent, no periodic box); overdamped Langevin dynamics at time step  $dt = 5 \times 10^{-4}$  remove centre-of-mass drift each step, and rotational invariance is left free ( $V$  and  $\varphi$  are rotation invariant).

## 8.2 Signed dihedral and the correct generalized mean force

The dihedral is computed by an `atan2` construction (not `arccos`), wrapped to  $[-\pi, \pi)$ , and is validated for range, periodic continuity, translation/rotation invariance, the expected reflection sign flip, known trans/gauche conformers, and finite-difference gradient agreement. All histograms, KDEs, interpolations, differentiations and integrations on  $\varphi$  are periodic (wrapped-Gaussian kernels; the one-dimensional mean force is integrated after enforcing  $\int_{-\pi}^{\pi} F'(\varphi) d\varphi = 0$ ).

Because  $\varphi$  is a nonlinear collective variable, the local mean force is *not*  $\partial V/\partial \varphi$ . Using any vector field  $v$  with  $v \cdot \nabla \varphi = 1$  [4, 7], the local mean force whose conditional average is  $dF/d\varphi$  is

$$f_{\text{loc}}(q) = \nabla V \cdot v - \beta^{-1} \nabla \cdot v, \quad v = \frac{\nabla \varphi}{|\nabla \varphi|^2}, \quad (32)$$

with the geometric term  $\nabla \cdot v = \Delta \varphi / |\nabla \varphi|^2 - 2(\nabla \varphi^\top H \varphi \nabla \varphi) / |\nabla \varphi|^4$  ( $H \varphi$  the Hessian). This term does *not* vanish for a dihedral (unlike the linear toy CV  $\xi = x$ ); it is evaluated by exact automatic differentiation, validated against finite differences to  $10^{-11}$  and against the analytic decoupled reduction. The ABF bias applies  $+A'(\varphi)\nabla \varphi$ , a conservative force that flattens the  $\varphi$  marginal when  $A = F$ . The reference free energies are constructed independently of the tested estimators by exact internal-coordinate free-energy perturbation (bonds sampled from  $r^2 e^{-\beta k_0(r-d_0)^2/2}$ , angles from  $\sin \theta e^{-\beta k_\theta(\theta-\theta_0)^2/2}$ , Jacobian-exact), with a convergence-ladder cross-check; only post-hoc metrics and the labelled oracle diagnostic ever read them. Decoupled (LJ off) butane and pentane recover  $V_4(+V_4) + C$  to machine precision (the B0/P0 exact gates).

Table 18: Butane ( $\beta=1$ , localized): median over production seeds. mFR and OPES compared to ABF on the same matched seeds.

| Method            | final $L_2(F)$ | int. $L_2(F)$ | final $L_2(F')$ | transitions | anc. ESS |
|-------------------|----------------|---------------|-----------------|-------------|----------|
| ABF               | 0.275          | 7.732         | 0.287           | 470387      | –        |
| ABF+mFR (est.)    | 0.275          | 7.735         | 0.288           | 469736      | 351      |
| ABF+mFR (uniform) | 0.275          | 7.751         | 0.287           | 469984      | 350      |
| ABF+mFR (oracle)  | 0.274          | 7.734         | 0.287           | 470090      | 361      |
| OPES              | 0.353          | 9.674         | 0.364           | 165928      | –        |

## 9 Case IV: butane — an easy molecular control

Butane isolates a single torsional slow coordinate. Because its central dihedral  $\varphi_1$  resolves *all* of the important conformational metastability (there is no other slow degree of freedom and no nonbonded pair, Section 8), a well-tuned ABF along  $\varphi_1$  should already sample the free-energy landscape well. The scientific question is whether the marginal Fisher–Rao (mFR) correction adds anything, or merely adds resampling variance.

### 9.1 Design and reference

The reaction coordinate is  $\xi = \varphi_1$ ; the reference is the analytic  $F(\varphi_1) = V_4(\varphi_1) + C$  (exact because butane has no LJ pair), confirmed by the internal-coordinate free-energy-perturbation pipeline to machine precision. We compare, at matched seeds, initial conditions,  $N$ , time step, estimator and force-evaluation budget: shared-bias  $N$ -replica **ABF**; the deployable **ABF+mFR** with the online estimated target; the **uniform** and **oracle** mFR diagnostics; and tuned periodic **OPES**. The mFR rate was tuned on held-out seeds by median integrated  $L_2(F)$ ; the ladder is reported in Section 10.3. Production uses  $N = 384$  replicas over disjoint seeds (localized trans initialization, with a dispersed three-basin initialization as a secondary diagnostic).

**Practical-equivalence protocol (pre-registered).** Because a non-significant difference is not evidence of equivalence, we pre-declared a practical-equivalence margin of **10% relative change** in final and time-integrated  $L_2(F)$ , and test it with matched-seed bootstrap 95% confidence intervals: mFR is declared *practically equivalent* to ABF if the CI of the relative change lies within  $\pm 10\%$ , *mildly harmful* if it lies entirely above  $+10\%$ , and a *genuine improvement* if entirely below  $-10\%$ .

### 9.2 Results

Under a well-tuned ABF, butane’s dihedral mixes rapidly: replicas cross the trans $\leftrightarrow$ gauche barrier many times and the  $\varphi_1$  marginal is flat long before the estimator has converged. Consequently the estimated mFR target matches the empirical marginal almost immediately, the log-density-ratio score is near zero, and the birth–death event rate is essentially zero ( $\approx 0.000$ ): **mFR turns itself off**. Its ancestor ESS stays at the full population and its free-energy error is byte-for-byte that of ABF. The equivalence test (Table 19) returns **equivalent** for the deployable mFR at  $\beta = 1$ : the median final  $L_2(F)$  is 0.275 versus 0.275 for ABF, a relative change of  $\sim 0.1\%$  with a 95% CI comfortably inside  $\pm 10\%$  and a win rate near one-half — the signature of two methods that are, for practical purposes, the same run. The uniform and oracle mFR diagnostics are equally equivalent, confirming that at this event rate the FR *target* is irrelevant because the mechanism barely acts. Tuned OPES, by contrast, reaches only 0.353 and is classified **harmful** by the same test (a 17–29%

Table 19: Butane practical-equivalence test (matched-seed bootstrap 95% CI of the relative change in final  $L_2(F)$  vs ABF;  $\pm 10\%$  margin).

| Cell                       | Method        | rel. med. | 95% CI           | win rate | verdict           |
|----------------------------|---------------|-----------|------------------|----------|-------------------|
| butane $\beta=0.5$ , trans | mFR (est.)    | 0.001     | [-0.001, 0.002]  | 0.42     | <b>equivalent</b> |
| butane $\beta=0.5$ , trans | mFR (oracle)  | 0.000     | [-0.001, 0.003]  | 0.50     | <b>equivalent</b> |
| butane $\beta=0.5$ , trans | mFR (uniform) | 0.000     | [-0.001, 0.002]  | 0.50     | <b>equivalent</b> |
| butane $\beta=0.5$ , trans | OPES          | 0.189     | [0.181, 0.193]   | 0.00     | <b>harmful</b>    |
| butane $\beta=1$ , disp.   | mFR (est.)    | -0.001    | [-0.002, 0.002]  | 0.56     | <b>equivalent</b> |
| butane $\beta=1$ , disp.   | mFR (oracle)  | -0.002    | [-0.002, 0.000]  | 0.69     | <b>equivalent</b> |
| butane $\beta=1$ , disp.   | mFR (uniform) | -0.001    | [-0.002, 0.002]  | 0.56     | <b>equivalent</b> |
| butane $\beta=1$ , disp.   | OPES          | 0.174     | [0.167, 0.183]   | 0.00     | <b>harmful</b>    |
| butane $\beta=1$ , trans   | mFR (est.)    | 0.001     | [0.001, 0.004]   | 0.38     | <b>equivalent</b> |
| butane $\beta=1$ , trans   | mFR (oracle)  | -0.001    | [-0.001, 0.001]  | 0.69     | <b>equivalent</b> |
| butane $\beta=1$ , trans   | mFR (uniform) | 0.002     | [0.001, 0.003]   | 0.06     | <b>equivalent</b> |
| butane $\beta=1$ , trans   | OPES          | 0.292     | [0.280, 0.300]   | 0.00     | <b>harmful</b>    |
| butane $\beta=2$ , trans   | mFR (est.)    | 0.002     | [0.001, 0.005]   | 0.33     | <b>equivalent</b> |
| butane $\beta=2$ , trans   | mFR (oracle)  | -0.000    | [-0.001, 0.001]  | 0.58     | <b>equivalent</b> |
| butane $\beta=2$ , trans   | mFR (uniform) | 0.002     | [0.001, 0.005]   | 0.17     | <b>equivalent</b> |
| butane $\beta=2$ , trans   | OPES          | 0.765     | [0.748, 0.790]   | 0.00     | <b>harmful</b>    |
| pentane $\beta=1$ , disp.  | mFR active    | 0.051     | [0.040, 0.068]   | 0.06     | <b>equivalent</b> |
| pentane $\beta=1$ , disp.  | mFR (est.)    | 0.007     | [-0.003, 0.014]  | 0.38     | <b>equivalent</b> |
| pentane $\beta=1$ , disp.  | mFR (oracle)  | 0.002     | [-0.007, 0.006]  | 0.44     | <b>equivalent</b> |
| pentane $\beta=1$ , disp.  | mFR (uniform) | 0.008     | [-0.005, 0.012]  | 0.38     | <b>equivalent</b> |
| pentane $\beta=1$ , disp.  | OPES          | 0.236     | [0.221, 0.276]   | 0.00     | <b>harmful</b>    |
| pentane $\beta=1$ , trans  | mFR active    | 0.059     | [0.043, 0.070]   | 0.00     | <b>equivalent</b> |
| pentane $\beta=1$ , trans  | mFR (est.)    | 0.004     | [-0.002, 0.013]  | 0.38     | <b>equivalent</b> |
| pentane $\beta=1$ , trans  | mFR (oracle)  | -0.003    | [-0.007, 0.005]  | 0.50     | <b>equivalent</b> |
| pentane $\beta=1$ , trans  | mFR (uniform) | 0.004     | [-0.002, 0.012]  | 0.38     | <b>equivalent</b> |
| pentane $\beta=1$ , trans  | OPES          | 0.373     | [0.357, 0.391]   | 0.00     | <b>harmful</b>    |
| pentane $\beta=2$ , trans  | mFR active    | 0.021     | [0.014, 0.024]   | 0.00     | <b>equivalent</b> |
| pentane $\beta=2$ , trans  | mFR (est.)    | 0.003     | [0.002, 0.007]   | 0.12     | <b>equivalent</b> |
| pentane $\beta=2$ , trans  | mFR (oracle)  | -0.003    | [-0.004, -0.000] | 0.69     | <b>equivalent</b> |
| pentane $\beta=2$ , trans  | mFR (uniform) | 0.003     | [0.001, 0.006]   | 0.25     | <b>equivalent</b> |
| pentane $\beta=2$ , trans  | OPES          | 0.585     | [0.573, 0.614]   | 0.00     | <b>harmful</b>    |

relative increase in  $L_2(F)$ , CI entirely above the margin, winning 0/16 seeds): its smooth deposited bias converges more slowly than ABF’s direct mean-force estimate on this easy, smooth torsional landscape at matched force-evaluation budget. The easy-regime hierarchy at  $\beta = 1$  is therefore  $\text{ABF} \approx \text{mFR} \gg \text{OPES}$ .

*Remark 9* (Shared estimator floor). The absolute  $L_2(F) \approx 0.27$  is a converged *estimator-resolution* floor, not under-sampling: the median error plateaus by step 5000 and is identical for ABF and every mFR variant to within  $5 \times 10^{-4}$ . It arises from the finite kernel bandwidth and grid discretization of the mean-force estimator on this sharp torsion (the FEP reference itself matches  $V_4$  to  $10^{-9}$ , and the recovered barrier is 89% of the analytic value with the correct shape and sign). Because every method is compared at *identical* estimator settings, this common floor cancels in the matched-seed differences, which is exactly what the equivalence and harm/benefit conclusions rest on.




Figure 32: Butane: ABF, mFR variants and OPES converge to the same  $F(\varphi_1)$ ; the mFR relative change vs ABF sits inside the  $\pm 10\%$  practical-equivalence band.

### 9.3 Temperature slice

To probe whether a colder butane creates any early-time headroom for mFR — where a higher barrier slows the ABF transient and leaves the marginal non-flat for longer — we sweep  $\beta \in \{0.5, 1, 2\}$  ( $\beta=1$  is the cell above). Figure 32b places every temperature on the equivalence axis. At the tuned rate the deployable mFR stays inside the  $\pm 10\%$  band at all three temperatures: even at the colder  $\beta=2$ , where the barrier is  $\sim 11 k_B T$  and ABF takes longer to flatten  $\varphi_1$ , the tuned mFR does not convert the slightly larger transient into a measurable convergence gain, and the rate ladder (Section 10.3) shows that pushing the rate to obtain events only degrades accuracy and diversity. The headline easy-regime conclusion is the honest, defensible statement “*no meaningful gain within the predefined margin*”, which is stronger than claiming mFR “does not work”: the mechanism is simply not needed when ABF already resolves the only slow coordinate, and forcing it to act adds resampling variance for no return.

## 10 Case V: pentane — coupled torsions and an incomplete reaction coordinate

Pentane has two dihedrals,  $\varphi_1$  and  $\varphi_2$ , coupled by the 1–5 Lennard–Jones interaction (the “pentane effect”; Section 8). We bias only  $\varphi_1$ , so  $\varphi_2$  is a *hidden* slow coordinate: sampling  $F(\varphi_1)$  correctly requires the conditional law  $p(\varphi_2 | \varphi_1)$  to equilibrate, which ABF does not directly drive. This is the demanding test of whether mFR reallocation improves reaction-coordinate support *without* corrupting the conditional law.

### 10.1 Design and reference

The joint reference  $F(\varphi_1, \varphi_2)$  is built by internal-coordinate free-energy perturbation (Figure 33a); the one-dimensional reference  $F(\varphi_1)$  is its marginal (Figure 33b) and is barely shifted from  $V_4(\varphi_1)$  — the coupling lives almost entirely in the *conditional*, not the marginal. This is exactly what makes a “false improvement” possible: a method can flatten the  $\varphi_1$  population while mis-sampling  $\varphi_2 | \varphi_1$ . We therefore report, against the joint reference, per- $\varphi_1$ -bin conditional total variation and KL, conditional basin probabilities over  $\{T, G^+, G^-\}$ , a reference-weighted aggregate conditional error, and visits to all nine torsional basins. Methods, seeds, budget and tuning mirror Section 9;




Figure 33: Pentane reference. The 1–5 coupling forbids  $g^+g^-$  (left) but barely shifts the  $\varphi_1$  marginal (right): the coupling lives in the *conditional* law, enabling “false improvement”.

Table 20: Pentane ( $\beta=1$ , localized): marginal accuracy *and* hidden-conditional fidelity. **cond.** TV is the reference-weighted total variation of  $p(\varphi_2|\varphi_1)$ .

| Method            | final $L_2(F)$ | int. $L_2(F)$ | final $L_2(F')$ | transitions | anc. ESS | cond. TV | cond. basin | basins |
|-------------------|----------------|---------------|-----------------|-------------|----------|----------|-------------|--------|
| ABF               | 0.284          | 9.236         | 0.297           | 551668      | –        | 0.028    | 0.022       | 9.000  |
| ABF+mFR (est.)    | 0.287          | 9.265         | 0.298           | 550917      | 346      | 0.027    | 0.020       | 9.000  |
| ABF+mFR (uniform) | 0.287          | 9.265         | 0.298           | 550950      | 345      | 0.027    | 0.021       | 9.000  |
| ABF+mFR (oracle)  | 0.285          | 9.226         | 0.296           | 551134      | 361      | 0.027    | 0.020       | 9.000  |
| OPES              | 0.393          | 12.082        | 0.442           | 195576      | –        | 0.027    | 0.020       | 9.000  |

we run  $\beta \in \{1, 2\}$  with localized trans–trans initialization (dispersed nine-basin as a secondary diagnostic).

## 10.2 Results

The central distinction is *healthy* improvement (better  $F(\varphi_1)$  with better-or-unchanged conditional fidelity) versus *false* improvement (flatter  $\varphi_1$  but worse  $p(\varphi_2 | \varphi_1)$ ). The data give a clear answer, and it is neither: at  $\beta = 1$  ABF attains a median final  $L_2(F(\varphi_1))$  of 0.284 with a reference-weighted conditional TV of 0.028; the deployable (tuned, gentle) mFR attains 0.287 and 0.027 — **practically equivalent on both the marginal and the hidden conditional**, because at the tuned rate the birth–death barely fires (ancestor ESS stays at  $\sim 90\%$  of the population). The  $\varphi_1$  marginal is simply not starved: ABF flattens it as easily as in butane, so there is no support deficit for mFR to correct, and the coupling-induced difficulty lives in  $\varphi_2$ , which mFR does not bias. The labelled **fr\_active** probe (rate 0.20) makes the consequence of forcing the mechanism explicit: it *worsens* the marginal (0.301 vs 0.284), collapses ancestor ESS to  $\sim 180$ , and at this primary trans–trans setting also *worsens* the conditional TV (0.038 vs 0.028) — harm on every axis, not a healthy gain. Figures 34a and 34b show the marginal accuracy, the hidden-conditional fidelity, and the basin occupancies for every method. Tuned OPES is markedly worse on the marginal (0.393 vs 0.284) at matched force-evaluation budget — its deposited bias under-explores  $\varphi_1$  (and at  $\beta = 2$  misses two of the nine basins) — while its conditional is comparable; both its native-reweight and common mean-force estimates are reported.






Figure 34: Pentane hidden-coordinate diagnostics: whether a method’s  $\varphi_1$  accuracy is matched by faithful  $p(\varphi_2|\varphi_1)$  (left) and how it populates the nine torsional basins (right).




Figure 35: mFR rate ladder on the alkanes: accuracy (integrated  $L_2(F)$ ) and ancestor ESS versus FR rate, identifying the useful window (if any) and the over-aggressive boundary.

### 10.3 Rate–diversity behaviour and frozen-bias validation

Figure 35 maps the mFR rate against accuracy and ancestor ESS on both alkanes. The result is unambiguous and is the same for butane and pentane: there is *no useful window*. Integrated  $L_2(F)$  is smallest at the gentlest rate and rises monotonically as the rate increases — always at or above the ABF baseline — while ancestor ESS collapses (from  $\sim 95\%$  of the population toward  $\sim 35\%$  at the strongest rate). For pentane the reference-weighted conditional TV  $p(\varphi_2 | \varphi_1)$  *also* worsens monotonically with rate, so an aggressive mFR simultaneously degrades the marginal free energy, corrupts the hidden conditional, and destroys lineage diversity. The pre-registered selection (smallest median integrated  $L_2(F)$  subject to genealogy) therefore returns the effectively inactive rate (0.02), and the labelled `fr_active` probe at 0.20 confirms the harm at production scale with matched seeds and full conditional diagnostics. A representative frozen-bias validation freezes the learned  $B(\varphi_1)$ , runs fresh dynamics with no updates and no birth–death, and reconstructs  $F_{\text{recon}} = B - \beta^{-1} \log p_B + C$ ; agreement of the reconstructed profile with the reference certifies that the learned bias — not a resampling artifact — carries the free-energy information. Where the conditional law is the binding constraint, the analysis relates the mFR gain to the measured ABF error, the  $\varphi_1$  support, the hidden-conditional error, the transition counts, the genealogy cost and  $\beta$ , rather than to the nominal barrier alone.

The frozen-bias reconstruction (Table 21) confirms that the learned  $B(\varphi_1)$  — whether from ABF or mFR — carries the free-energy information: fresh dynamics under the frozen bias, with

Table 21: Frozen-bias validation (Part F): freeze the learned bias  $B(\varphi_1)$ , run fresh dynamics with no updates and no birth–death, and reconstruct  $F_{\text{recon}} = B - \beta^{-1} \log p_B + C$  from the biased marginal. The reconstruction error is far below the online mean-force-integral error — so the learned bias is much more accurate than the online estimator floor suggests — and is the same for ABF and mFR, confirming the free energy resides in the learned bias, not in the online transient or in resampling.

| Cell              | Method     | online $L_2(F)$ | frozen recon. $L_2(F)$ |
|-------------------|------------|-----------------|------------------------|
| butane $\beta=1$  | ABF        | 0.273           | 0.049                  |
| butane $\beta=1$  | mFR (est.) | 0.274           | 0.050                  |
| pentane $\beta=1$ | ABF        | 0.285           | 0.104                  |
| pentane $\beta=1$ | mFR (est.) | 0.290           | 0.104                  |
| pentane $\beta=2$ | ABF        | 0.198           | 0.050                  |
| pentane $\beta=2$ | mFR (est.) | 0.198           | 0.051                  |

no updates and no birth–death, reconstruct  $F(\varphi_1)$  from the biased marginal to an error *well below* the online mean-force-integral estimate. This both certifies that the equivalence of ABF and tuned mFR is a property of the learned bias itself (not of the online transient or of resampling) and localizes the shared  $\sim 0.27$  online floor (Remark 9) specifically to the finite resolution of the mean-force estimator: the underlying bias is far more accurate than that floor suggests, and it is the same for ABF and mFR.

## 11 Harder one-dimensional molecular coordinates: end-to-end distance

The  $\varphi_1$  studies (Sections 9 and 10) left one mechanism untested: in both alkanes ABF resolved the biased coordinate *easily*, so marginal Fisher–Rao (mFR) never had a support deficit to repair. To probe whether mFR helps when ABF is genuinely sample-starved *along the biased coordinate*, we replace the torsion by a harder but chemically meaningful scalar: the pentane end-to-end distance

$$\xi(q) = R_{15}(q) = |q_5 - q_1|, \quad (33)$$

which measures molecular extension/compactness and depends jointly on both torsions and on the bond/angle fluctuations. Many torsional states map to the same  $R_{15}$ , so the hidden conditional here is the full joint torsion law  $p(\varphi_1, \varphi_2 \mid R_{15})$ . As a secondary, easier control we also study the butane  $R_{14} = |q_4 - q_1|$ .

### 11.1 Distance CV, generalized mean force, and reference

With  $r = q_5 - q_1$ ,  $R = |r|$ ,  $e = r/R$  we have  $\nabla_{q_5} R = e$ ,  $\nabla_{q_1} R = -e$ ,  $|\nabla R|^2 = 2$ , and (because  $H_R \nabla R = 0$ ) the den Otter geometric term is exactly  $\nabla \cdot v = 2/R$  for  $v = \nabla R / |\nabla R|^2$ , giving the local mean force

$$f_R(q) = \frac{1}{2} e \cdot (\nabla_{q_5} V - \nabla_{q_1} V) - \frac{2}{\beta R}. \quad (34)$$

The analytic gradient/divergence match autodiff to  $< 10^{-9}/10^{-7}$ ;  $R$  never approaches 0 under the stiff bonds, and identical soft walls confine every method to a reliable interval. Because  $R_{15}$  is non-periodic we use reflected-boundary KDE and a left-anchored trapezoidal PMF. The reference is built by internal-coordinate importance sampling: bonds/angles from their exact Boltzmann+Jacobian law, dihedrals from  $e^{-\beta V_4}$ , reweighted by  $e^{-\beta V_{\text{nb}}}$  (the 1–5 LJ; butane has none, so its reference is




Figure 36: Pentane  $R_{15}$ : (a) reference free energy (minimum at the extended state, compact configurations strongly disfavoured); (b) ABF convergence of the thermal-window  $L_2(F)$  for the four screen cells — the  $\beta=2$  cells remain far above the  $\beta=1$  ones (1.43–1.58 vs 0.67–0.74 kT); (c) the ABF-sampled  $R_{15}$  marginal (solid) against the *unbiased* Boltzmann law (dashed, same colour): ABF broadens the marginal towards compact  $R$  as intended, but the residual structure at  $\beta=2$  is the support deficit the screen detects.

exact by construction). The primary (‘ $V_4$ -proposal’) estimator has effective-sample fraction 0.90; an independent uniform-proposal scheme agrees to 0.180 kT in the thermal window (range 9.8 kT), well inside the 15% success threshold. A minimal resolution gate ( $n_{\text{grid}} \in \{128, 256\}$ , bandwidth  $\{h, h/2\}$ ) leaves the reference shape and the ABF diagnosis unchanged.

## 11.2 ABF-only starvation screen: $R_{15}$ is the first genuinely starved coordinate

We ran ABF alone across  $\beta \in \{1, 2\}$  (butane  $R_{14}$ :  $\beta \in \{1, 2, 3\}$ ) with localized-trans and dispersed initializations, classifying each cell by *distinct evidence families* of the pre-registered criteria: a magnitude/plateau family (normalized residual  $L_2$  above the estimator floor), a support family (under-supported thermal  $R$  bins), a mixing family (compact $\leftrightarrow$ extended round trips per replica), and a conditional family (biased  $p(\varphi_1, \varphi_2 | R)$ ). Because the raw magnitude and plateau signals are not statistically independent, we count them as one family and require  $\geq 2$  *distinct* families for a “starved” verdict. Table 22 reports the outcome.

Unlike the  $\varphi_1$  torsion — which ABF resolves easily — the pentane  $R_{15}$  coordinate is classified **starved** at  $\beta = 2$  (thermal  $L_2 \approx 14$ –16% of the free-energy range,  $\sim 2.5\times$  the estimator floor, with  $\sim 22\%$  of thermal  $R$  bins under-supported and, in the dispersed cell, near-absent compact $\leftrightarrow$ extended mixing). It is easy at  $\beta = 1$ . The butane control  $R_{14}$  is **easy** at  $\beta \in \{1, 2\}$  and only **intermediate** at  $\beta = 3$ . This yields a clean, graded picture: a chemically meaningful *distance* CV *can* make ABF sample-starved along the biased coordinate once the torsional barriers that change the distance are rarely crossed, and the difficulty is ordered by molecule — pentane  $R_{15}$  reaches a full *starved* verdict already at  $\beta = 2$ , whereas the easier butane  $R_{14}$  never does at any temperature tested, reaching only *intermediate* (a single evidence family: high error, but no support or mixing deficit) at the colder  $\beta = 3$ . A minimal resolution gate leaves the  $R_{15}$  starvation diagnosis unchanged ( $L_2$  21–22% across  $n_{\text{grid}} \in \{128, 256\}$  and bandwidth  $\{h, h/2\}$ ). A run-length control settles the obvious alternative explanation — that the cell is merely under-converged at a fixed budget rather than starved: doubling the ABF budget from 80k to 160k steps reduces the thermal  $L_2$  only from 15.1% to 13.7% of the range (a 9% relative gain for  $2\times$  the compute), leaves the error *plateaued* (final-quarter change  $-0.04$ , i.e. no longer decreasing) still far above the  $\sim 6$ –7% estimator floor, and does *not* fill the under-supported bins (low-support fraction  $0.22 \rightarrow 0.28$ ). The deficit is

Table 22: Pentane  $R_{15}$  and butane  $R_{14}$  ABF-only starvation screen (round-trips per replica; families = distinct evidence families,  $\geq 2 \Rightarrow$  starved).

| cell               | final $L_2(F)$ | r.t./N | families | verdict      |
|--------------------|----------------|--------|----------|--------------|
| R15 $\beta=1$ disp | 0.666          | 5.6    | 0        | easy         |
| R15 $\beta=1$ TT   | 0.744          | 9.6    | 0        | easy         |
| R15 $\beta=2$ disp | 1.582          | 0.3    | 2        | starved      |
| R15 $\beta=2$ TT   | 1.428          | 9.0    | 2        | starved      |
| R14 $\beta=1$ TT   | 0.505          | 47.5   | 0        | easy         |
| R14 $\beta=2$ TT   | 0.876          | 22.0   | 0        | easy         |
| R14 $\beta=3$ TT   | 1.391          | 17.9   | 1        | intermediate |

therefore a genuine *support* starvation that more of the same sampling does not cure. Because a starved cell exists, the pre-registered gate triggers the full  $R_{15}$  mFR/OPES comparison below — the study’s one clean opportunity for mFR to help.

### 11.3 Does mFR rescue a starved ABF? No.

On the starved  $R_{15}$  ( $\beta = 2$ , localized trans) cell we ran the full ladder at matched seeds, initial conditions, replicas, steps and force budget. The result is unambiguous and, notably, *the same as when ABF is not starved*: mFR gives **no improvement**. The deployable estimated-target mFR (rate 0.02) is **equivalent** to ABF (median change in final  $L_2$ : 2.0%; win rate 0.00). At production scale (eight matched seeds) ABF reaches a median final  $L_2$  of 1.500 and the deployable mFR 1.532 (+2.1%, inside the  $\pm 10\%$  practical-equivalence margin) with the hidden-conditional fidelity unchanged (0.072 both) — and mFR fails to beat ABF on *every one* of the eight matched seeds (win rate 0.00). The labelled aggressive probe (rate 0.10) is severely *harmful*: 2.015 (+34.4%), with the ancestor ESS collapsing to 0.06  $N$  and the conditional degrading to 0.087.

Two diagnostics make the negative result decisive rather than a tuning artifact. First, the *uniform* and *oracle* targets perform identically to the estimated one (1.533, +2.2% each): even a target built from the *exact reference free energy* — the quantity being computed — buys nothing. The failure is therefore not poor target estimation but the reallocation mechanism itself. Second, the OPES baseline was swept independently on held-out seeds for this cell over barrier  $\in \{8, 16, 24\}$ , kernel width  $\sigma \in \{0.05, 0.10\}$  and deposition pace  $\in \{200, 400\}$ : its *best* configuration (barrier 8,  $\sigma = 0.05$ , pace 200) still trails ABF by +88%, larger barriers are monotonically worse (+153 to +156%), and — most tellingly — *every* configuration records **zero** compact $\leftrightarrow$ extended round trips (the production configuration also degrades the conditional to 0.192). An OPES bias grows only where walkers have already been, so on this steep cold profile no setting mobilises the barrier crossing at all. The same bottleneck therefore defeats a different marginal-biasing strategy, which is evidence about the *coordinate*, not about any one algorithm. The rate ladder is correspondingly *monotone in the wrong direction*: final  $L_2$  rises from the ABF value as the mFR rate increases (+0.4, +1.0, +4.4, +13.4% at rates 0.005, 0.02, 0.05, 0.10) while the ancestor ESS falls (0.97  $\rightarrow$  0.87  $\rightarrow$  0.62  $\rightarrow$  0.17 of  $N$ ). There is no useful window. Figure 37 shows the paired comparison, the rate ladder, and the hidden-conditional fidelity.

**Provenance of the deployed rate.** We record this precisely, because the ladder is monotone and therefore makes the selection rule degenerate. The pre-registered rule — the smallest median integrated  $L_2(F)$  on the held-out tuning seeds subject to the genealogy guard — has its argmin at the *gentlest* rung, 0.005 (median integrated  $L_2 = 83.30$  against ABF’s 83.16), not at the 0.02 we

Table 23: R15 (starved cell) production comparison (median over seeds).

| cell             | method       | $L_2(F)$ | #round-trips | cond. TV | anc. ESS/N |
|------------------|--------------|----------|--------------|----------|------------|
| R15 $\beta=2$ TT | abf          | 1.500    | 9088         | 0.072    | —          |
| R15 $\beta=2$ TT | fr active    | 2.015    | 6132         | 0.087    | 0.06       |
| R15 $\beta=2$ TT | fr estimated | 1.532    | 8998         | 0.072    | 0.79       |
| R15 $\beta=2$ TT | fr oracle    | 1.533    | 8978         | 0.072    | 0.79       |
| R15 $\beta=2$ TT | fr uniform   | 1.533    | 8979         | 0.072    | 0.79       |
| R15 $\beta=2$ TT | opes         | 3.319    | 0            | 0.192    | —          |




Figure 37:  $R_{15}$  mFR/OPES on the *starved*  $\beta = 2$  cell: (a) matched-seed relative change in final  $L_2(F)$  vs ABF (band =  $\pm 10\%$  equivalence, dashed =  $-15\%$  positive threshold); (b) the mFR rate ladder — integrated  $L_2$  rises and ancestor ESS collapses with rate, no useful window; (c) hidden-conditional  $p(\phi_1, \phi_2 | R)$  fidelity per method.

deploy (83.58); the two stiffer rungs, 0.05 and 0.10, are worse still (84.50 and 85.94), and only 0.10 fails the genealogy guard. In other words the rule, applied literally to a monotone ladder, selects the rate that does the *least*, because doing less is what minimises the error. We deploy 0.02 instead — the gentlest rung at which birth–death is meaningfully active (ancestor ESS 0.87  $N$  rather than 0.97  $N$ ) — so that the reported comparison actually exercises the mechanism under test. This choice is *conservative against* mFR, not for it: the rule’s own pick would have scored slightly better (+0.3% versus +1.3% on the tuning seeds). Every genealogy-passing rung is classified *equivalent* to ABF, so the verdict does not depend on which is quoted; what the degeneracy shows is that on this coordinate the best available mFR configuration is the one closest to not running mFR at all.

The mechanism is diagnostic (Section 13, Q3):  $R_{15}$  starvation is a *dynamical* bottleneck — extended replicas rarely cross the cold torsional barriers needed to reach compact  $R$  — not a *support* deficit that marginal reallocation can fill. The support diagnostics locate the failure precisely, and the answer is more interesting than “reallocation cannot fill the gap”. At the deployable rate mFR barely fires (ancestor ESS 0.79  $N$ ), so the support profile is essentially untouched (under-supported fraction 0.218 versus ABF’s 0.223) and the error with it. Pushed to rate 0.10 it does exactly what it is designed to do — the under-supported fraction falls by nearly two thirds,  $0.223 \rightarrow 0.082$  — and yet the free energy becomes markedly *worse* (+34%). The reason is that the newly “supported” bins are filled with *clones* rather than independent samples: ancestor ESS collapses to 0.06  $N$  and the round-trip count per replica actually *falls* ( $8.9 \rightarrow 6.0$ ), because independently evolving replicas are replaced by copies of the few that had already crossed. Reallocation thus converts a support deficit into a *diversity* deficit: nominal occupancy improves while the effective sample size — the

quantity the mean-force estimator actually consumes — degrades. Amplification is not discovery. The independently tuned OPES baseline and the diagnostic uniform/oracle targets are reported alongside in Table 23.

The frozen-bias fairness control (freeze  $B(R)$ , run fresh dynamics with no updates and no birth-death, reconstruct  $F = B - \beta^{-1} \log p_B + C$ ) delivers a second, independent confirmation that this cell is genuinely starved — and it behaves *oppositely* to the converged  $\varphi_1$  case. There, freezing the learned bias reconstructed  $F$  far more accurately than the online mean-force integral (0.05 vs 0.27), proving the online residual was a mere estimator-resolution floor. Here the reconstruction is *worse* than the online estimate for both methods (ABF 1.90 vs 1.48; mFR 1.92 vs 1.51, seed-averaged profiles): the learned  $B(R)$  is itself inaccurate because the underlying sampling is deficient, and freezing it does not help since the frozen dynamics remain trapped by the same rare-crossing bottleneck. The residual is therefore *not* a resolution artifact but real under-sampling — exactly the distinction the starvation screen is meant to detect — and mFR’s learned bias is no better than ABF’s, consistent with their practical equivalence.

## 12 Two-dimensional ABF and marginal Fisher–Rao on the torsion torus

The mandatory test promotes the hidden pentane torsion into the CV itself,

$$\xi(q) = (\varphi_1(q), \varphi_2(q)) \in \mathbb{T}^2, \quad (35)$$

with *the same* joint CV driving ABF and mFR ( $\xi_{\text{ABF}} = \xi_{\text{FR}}$ ), so mFR cannot silently reshape  $p(\varphi_2 | \varphi_1)$  behind a one-dimensional ABF. This asks directly whether two-dimensional mFR reallocation improves two-dimensional ABF.

### 12.1 Vector generalized mean force and periodic Poisson projection

For  $g_a = \nabla \varphi_a$ ,  $H_a = \nabla^2 \varphi_a$  (exact autodiff), the Gram matrix is  $G_{ab} = g_a \cdot g_b$  and the dual fields  $w_a = \sum_b (G^{-1})_{ab} g_b$  satisfy  $w_a \cdot g_c = \delta_{ac}$  exactly. The mean-force components  $f_a = \nabla V \cdot w_a - \beta^{-1} \nabla \cdot w_a$  use a divergence computed *analytically* from the per-dihedral gradients and Hessians (via  $dG^{-1} = -G^{-1} dG G^{-1}$ ), avoiding nested autodiff; it matches a finite-difference divergence of the dual field to  $< 10^{-4}$ , and on the decoupled model reduces to  $\nabla V \cdot w_a \rightarrow V'_4(\varphi_a)$  despite the two dihedrals sharing three atoms (a non-diagonal  $G$ ). The Gram matrix is well conditioned on physical geometries ( $\max \kappa = 21.3$ , 0 regularisations).

The scattered, smoothed mean-force field  $g = (g_1, g_2)$  on the torus grid is not exactly conservative; the scalar bias is its Hodge projection onto gradients,  $B = \arg \min_{\tilde{B}=0} \int_{\mathbb{T}^2} |\nabla B - g|^2$ , i.e. the periodic Poisson problem  $\Delta B = \nabla \cdot g$  solved by FFT ( $\hat{B}(k) = -i k \cdot \hat{g}(k) / |k|^2$ ,  $\hat{B}(0) = 0$ ). The projection recovers exact gradient fields and removes divergence-free noise to  $< 10^{-9}$  (Table 24). The configuration-space bias force is  $+\sum_a (\partial_{z_a} B) \nabla \varphi_a$ , applied every step, while the (Hessian) mean force is accumulated on a stride — a standard ABF subsampling that leaves the estimator unbiased.

**Exact decoupled gate.** With the 1–5 LJ switched off the exact free energy is  $V_4(\varphi_1) + V_4(\varphi_2) + C$ . The full 2-D pipeline recovers it to a thermal-window  $L_2$  of 0.463 kT on the  $48 \times 48$  grid, essentially unchanged from 0.471 kT on  $32 \times 32$  — the residual is resolution-limited in the *estimator*, not in the grid — and the internal-coordinate FEP reference reproduces  $V_4 \oplus V_4$  to  $1.8\text{e-}15$  — the residual is the shared mean-force-estimator resolution floor (Remark 9), not a systematic error.

Table 24: Two-dimensional ABF mathematical validation.

| check                                                           | result             |
|-----------------------------------------------------------------|--------------------|
| dual biorthogonality $w_a \cdot \nabla \varphi_c = \delta_{ac}$ | $< 10^{-9}$ (test) |
| analytic $\nabla \cdot w_a$ vs FD                               | $< 10^{-4}$ (test) |
| Poisson exact-gradient recovery                                 | $< 10^{-9}$ (test) |
| decoupled exact gate max abs err                                | 1.8e-15            |
| decoupled 2-D thermal $L_2$ (grid 48)                           | 0.463 kT           |
| Gram max $\kappa$ / regularisations                             | 21.3 / 0           |




Figure 38: (a) the joint reference  $F(\varphi_1, \varphi_2)$ ; (b) a representative 2-D ABF reconstruction; (c) the exact decoupled gate  $L_2$  versus grid resolution — flat, showing the residual is the mean-force estimator floor rather than a grid-discretisation error.

## 12.2 Two-dimensional density, score, and cloning

The joint marginal  $\hat{p}_t(z)$  is a separable wrapped-Gaussian KDE on the torus; the centred Fisher–Rao score  $S_t(z) = \log[(\hat{p}_t + \epsilon)/(\hat{q}_t + \epsilon)] - \int \hat{p}_t \log[\dots]$  is interpolated bilinearly to each replica, and fixed-population kill-and-clone copies the whole molecular configuration with ancestor labels and independent future noise. The deployable target is  $\hat{q}_{\text{est}} \propto e^{-\beta(F_{\text{EMA}} - B)}$ ; uniform and reference-oracle targets are diagnostic only, the latter guarded from deployable code.

**Frozen-bias check.** Freezing the learned  $B(\varphi_1, \varphi_2)$  and running fresh dynamics with no updates and no birth–death reconstructs  $F = B - \beta^{-1} \log p_B + C$  to 0.378, roughly *half* the online mean-force error (0.722), and identically so for ABF and mFR. The learned two-dimensional bias is therefore considerably more accurate than the online  $L_2$  suggests — the residual is a shared mean-force-estimator resolution floor, not a systematic error — and the ABF/mFR equivalence is a property of the learned bias itself, not of the online transient or of resampling.

## 12.3 ABF-only starvation screen: promoting the torsion removes the starvation

Before any mFR tuning we ran 2-D ABF alone on cells P2-A..D (Table 25), the primary being  $\beta = 2$  with localized trans–trans initialization, using the same distinct-evidence-family criteria as Section 11. The screen classifies the primary cell **easy**: at  $\beta \in \{1, 2\}$  and either initialization, ABF discovers all nine torsional basins (by step  $\sim 3800$  even when started cold and localized in trans–trans), sustains very strong mixing (100–360 joint round trips per replica), and converges to 5.0–7.3% of the free-energy range, i.e. at or within about 1.5 percentage points of the shared 5.8%

Table 25: Two-dimensional torsion  $(\varphi_1, \varphi_2)$  ABF-only starvation screen.

| cell              | final $L_2(F)$ | r.t./N | families | verdict      |
|-------------------|----------------|--------|----------|--------------|
| 2D $\beta=1$ disp | 0.504          | 358.4  | 0        | easy         |
| 2D $\beta=1$ TT   | 0.497          | 348.0  | 0        | easy         |
| 2D $\beta=2$ disp | 0.724          | 122.4  | 0        | easy         |
| 2D $\beta=2$ TT   | 0.732          | 131.9  | 0        | easy         |
| 2D $\beta=3$ TT   | 1.035          | 20.2   | 1        | intermediate |




Figure 39: 2-D ABF starvation screen: joint nine-basin occupancy (left) and thermal-window  $L_2(F)$  convergence (right) across the P2 cells.

estimator floor measured on the decoupled gate — and, decisively, with no support deficit: every basin is found and mixing is strong.

Because that is a uniformly negative screen, we added a colder escalation (P2-E,  $\beta = 3$ , localized) to check the verdict is not an artifact of insufficient difficulty. It is not: even there ABF finds all nine basins (by step  $\sim 8400$ ), keeps  $\sim 20$  round trips per replica, and plateaus at 10.4% — classified only *intermediate*, with the error creeping just past the conservative threshold but *no support deficit at all*. The contrast with the distance CV is the essential point: at the *colder*  $\beta = 3$  the two-dimensional CV still enjoys complete basin discovery and healthy mixing, whereas  $R_{15}$  at the *warmer*  $\beta = 2$  is genuinely starved with  $\sim 22\%$  of its thermally relevant bins under-supported. Biasing  $(\varphi_1, \varphi_2)$  drives both torsional barriers *directly*, so the support deficit that mFR exists to repair simply never forms. No cell being starved, the pre-registered gate calls for a compact neutrality control rather than a full successive-halving mFR study (Section 13).

### 13 Does promoting the hidden torsion make marginal Fisher–Rao useful?

We can now compare the three coordinates on a common axis: the one-dimensional torsion  $\varphi_1$  (Section 10), the harder one-dimensional distance  $R_{15}$  (Section 11), and the two-dimensional torsion  $(\varphi_1, \varphi_2)$  (Section 12). The organizing question is whether mFR’s benefit tracks *measured starvation of the biased coordinate*, and whether any apparent gain is a genuine support gain or a lineage-diversity loss.



### 13.1 Method comparison

On a cell the ABF screen flags as *starved* the pre-registered gate calls for the full ladder — multi-replica shared-bias ABF, ABF+mFR with estimated / uniform / oracle targets, and an independently swept OPES — which is what Section 11 reports for  $R_{15}$ . Since no two-dimensional cell is starved, the gate instead calls for the compact *neutrality control* reported here: ABF, the deployable estimated-target mFR, a labelled aggressive probe, and OPES (the uniform/oracle diagnostic targets are informative only where mFR has a support deficit to repair, and on  $R_{15}$  they proved indistinguishable from the estimated target). Both are run at matched seeds, initial conditions, replicas, steps and force-evaluation budget. We are explicit about the asymmetry in how hyperparameters were set, because it is not the same on the two arms. On  $R_{15}$  the mFR rate was chosen from a ladder run on disjoint held-out tuning seeds, and the OPES baseline was swept independently over barrier, kernel width and pace. On the two-dimensional arm neither search was run: the mFR rate (0.01) and the OPES configuration (barrier 8,  $\sigma = 0.30$ , pace 400) were fixed *a priori* in the configuration and never tuned, because the pre-registered gate does not authorise a hyperparameter study on a cell no screen classifies as starved. That is a real limitation of the two-dimensional OPES number in particular, which should be read as “OPES at a plausible default” rather than “OPES at its best”. It does not weaken the mFR verdict, which is a null: an untuned mFR that lands within 0.2% of ABF cannot have been flattered by the absence of tuning. The principal causal contrast is shared-bias ABF versus ABF+*estimated-target* mFR. Table 26 collects the two-dimensional production comparison and Figure 40 the matched-seed effect sizes and genealogy.

The deployable 2-D mFR verdict on the primary cell is **equivalent** (median change in final  $L_2(F)$  vs ABF: 0.2%): ABF reaches 0.724 and mFR 0.725, winning 0/6 matched seeds, with every genealogy, event-cap and basin guard satisfied — a genuine null, not a masked false improvement. A *positive* result requires all twelve pre-registered criteria; a lower free-energy error accompanied by degraded basin occupancy, conditional structure or genealogy beyond the stated bounds is a *false* improvement, not a success. OPES at its default configuration is again far worse (1.682, +132%); as noted above it was not swept here, so this is a weaker statement than the independently tuned  $R_{15}$  comparison.

Comparing the two coordinates exposes a second regularity, about the *cost* rather than the benefit of birth–death. The aggressive probe (rate 0.10) costs only +0.9% here, with ancestor ESS holding at  $0.35 N$ , whereas the same probe on the starved  $R_{15}$  cell cost +34% with ESS collapsing to  $0.06 N$ . The difference is how fast the system re-equilibrates after a clone: on the torus ABF sustains  $\sim 130$  joint round trips per replica, so cloned configurations decorrelate almost immediately and aggressive reallocation is merely useless; on the barrier-limited distance coordinate the clones stay correlated, so the same mechanism actively destroys the mean-force estimate. Birth–death is therefore cheap where it is unnecessary and expensive precisely where one would want to use it.

**Frozen-bias control: the same test, the opposite sign.** Freezing the learned two-dimensional bias  $B(\varphi_1, \varphi_2)$ , running fresh dynamics with no updates and no birth–death, and reconstructing  $F = B - \beta^{-1} \log p_B + C$  reproduces the reference far *better* than the online mean-force integral: 0.378 against 0.722 for ABF and 0.378 against 0.723 for mFR, i.e. the online residual overstates the error in the learned bias by about a factor of two. This completes a three-way control with a consistent reading. Where the screen calls the coordinate easy — the  $\varphi_1$  torsion (0.05 vs 0.27) and the two-dimensional torus here — freezing *improves* the reconstruction, so the online residual is a mean-force estimator-resolution floor and there is no systematic error for mFR to remove. Where the screen calls the coordinate starved —  $R_{15}$  — freezing makes it *worse* (1.92 vs 1.51; Section 11),

Table 26: Two-dimensional production comparison (median over seeds).

| cell            | method       | $L_2(F)$ | basins | cond. TV | anc. ESS/N |
|-----------------|--------------|----------|--------|----------|------------|
| 2D $\beta=2$ TT | abf          | 0.724    | 9      | 0.389    | —          |
| 2D $\beta=2$ TT | fr active    | 0.730    | 9      | 0.404    | 0.35       |
| 2D $\beta=2$ TT | fr estimated | 0.725    | 9      | 0.390    | 0.84       |
| 2D $\beta=2$ TT | opes         | 1.682    | 9      | 0.268    | —          |




Figure 40: Matched-seed relative change in final  $L_2(F)$  of each method vs ABF (left; the shaded band is the  $\pm 10\%$  practical-equivalence margin, the dashed line the  $-15\%$  positive threshold) and the final ancestor ESS of each mFR configuration (right; the dashed line is the  $0.30 N$  genealogy floor).

because the learned bias is itself built from deficient sampling. The control therefore discriminates the two situations independently of the starvation screen, and agrees with it in all three cases. That the two methods’ reconstructions here are identical to three decimals ( $0.3778$  vs  $0.3777$ ) is a further statement of the null: mFR’s learned bias is neither better nor worse than ABF’s.

### 13.2 Cross-case synthesis

Table 27 places all biased coordinates on one starvation map. The pattern across the study is that mFR’s usefulness is gated by *marginal starvation of the biased coordinate*: where ABF already resolves the coordinate (the  $\varphi_1$  marginals, and any coordinate the screen classifies easy), mFR has no support deficit to repair and is at best practically equivalent; the residual difficulty then lives in an *orthogonal* hidden coordinate that mFR — acting only on the biased marginal — cannot drive. Promoting the hidden torsion into the CV removes the orthogonal-coordinate excuse, so the two-dimensional comparison isolates whether mFR reallocation on the *actual* slow manifold helps, and distinguishes a support gain from a lineage-diversity loss (ancestor ESS) and from two-dimensional density-estimation noise (empty-cell score spikes, curl before projection).

Table 28 collects the alkane outcome in one place, on the closure sign convention. It is arranged in three blocks because the three coordinates fail for three different reasons, and flattening them into a single “mFR does not help on alkanes” would lose the only part of this that generalises. On the torsion coordinates ABF already resolves the barrier and there is nothing to reallocate. On  $R_{15}$  there is a genuine support deficit and reallocation still cannot use it, because the deficit is




Figure 41: Spatial birth (clone-source) and death maps of the 2-D mFR birth-death process on the torsion torus, showing where reallocation moves probability mass.

Table 27: Cross-case ABF starvation synthesis along the biased coordinate.

| coordinate                     | molecule       | ABF starved?    |
|--------------------------------|----------------|-----------------|
| $\varphi_1$ (1-D dihedral)     | butane/pentane | no (prior work) |
| $R_{15}$ (1-D distance)        | pentane        | starved         |
| $R_{14}$ (1-D distance)        | butane         | intermediate    |
| $(\varphi_1, \varphi_2)$ (2-D) | pentane        | easy (primary)  |

maintained by a rare barrier crossing rather than by an imbalance among reachable configurations — which the `fr_active` row makes concrete: pushed hard enough to repair the nominal support, mFR makes the free energy substantially worse. On the torsion torus the deficit is removed by construction, by putting the slow coordinate into the CV.

The three coordinates answer Q3 cleanly, and the answer is not the same in each. For the one-dimensional torsion  $\varphi_1$  and for the two-dimensional torus CV, the limiting factor is **(i) lack of marginal starvation**: ABF resolves the biased coordinate to the estimator floor (in 2-D it biases both dihedrals directly, discovering all nine basins and mixing strongly even cold and localized), so mFR has no support deficit to repair — its measured verdict on the primary 2-D cell is equivalent. The density-estimation machinery is *not* the bottleneck — the Gram matrix is well conditioned, the FFT Poisson projection removes real curl, and the decoupled gate recovers  $V_4 \oplus V_4$  to the shared floor — so we can exclude (ii). For the harder distance CV  $R_{15}$ , ABF *is* starved, which isolates the decisive distinction: the starvation is a **dynamical bottleneck**, not a support deficit reallocation can fill with useful samples. Extended replicas rarely cross the cold torsional barriers that change  $R_{15}$ , and mFR can only clone those that already have. The support diagnostics show the resulting trade explicitly: pushed hard enough to act, mFR *does* repair the nominal support (under-supported bin fraction  $0.223 \rightarrow 0.082$ ) yet makes the free energy 34% *worse*, because the newly occupied bins hold clones rather than independent samples (ancestor ESS  $\rightarrow 0.06 N$ , and round trips per replica *falling*  $8.9 \rightarrow 6.0$ ). The limiting factors are therefore **(iii) poor discovery** — amplification is not discovery — inseparably compounded by **(v) genealogical collapse**: reallocation here converts a support deficit into a diversity deficit, improving occupancy while degrading the effective sample size the mean-force estimator actually consumes. Conditional/orthogonal relaxation (iv) is a smaller

Table 28: Alkane closure. All entries are the paired relative change in `final_12_F` against ABF, so **negative is better**; the seeds column counts matched seeds on which mFR is strictly more accurate. The three blocks are three different limitations, not three instances of one. Regenerated by `scripts/build_closure_inventory.py`.

| Arm                                                                                      | $\Delta\text{err}$ (%) | 95 % CI        | seeds | Reading                                              |
|------------------------------------------------------------------------------------------|------------------------|----------------|-------|------------------------------------------------------|
| <i>Torsion CV: ABF already resolves the barrier</i>                                      |                        |                |       |                                                      |
| Butane $\varphi_1$ , $\beta=1$                                                           | +0.10                  | [0.07, 0.42]   | 6/16  | ABF-sufficient                                       |
| Pentane $\varphi_1$ , $\beta=1$                                                          | +0.39                  | [-0.23, 1.26]  | 6/16  | ABF-sufficient                                       |
| Pentane $\varphi_1$ , $\beta=2$                                                          | +0.27                  | [0.21, 0.71]   | 2/16  | ABF-sufficient                                       |
| <i>Distance CV <math>R_{15}</math>, <math>\beta=2</math>: the one starved coordinate</i> |                        |                |       |                                                      |
| practical rate                                                                           | +2.03                  | [1.62, 2.38]   | 0/8   | discovery-limited                                    |
| oracle target                                                                            | +2.01                  | [1.66, 2.66]   | 0/8   | ideal target does not help either                    |
| pushed (rate 0.20)                                                                       | +32.81                 | [28.09, 34.32] | 0/8   | <b>adverse</b> : support repaired, lineages collapse |
| <i>Torsion torus (<math>\varphi_1, \varphi_2</math>): promoting the hidden torsion</i>   |                        |                |       |                                                      |
| practical rate                                                                           | +0.22                  | [0.11, 0.29]   | 0/6   | ABF-sufficient                                       |

effect here, since the hidden torsion conditional  $p(\varphi_1, \varphi_2 \mid R)$  stays well captured. The unifying statement is that marginal starvation of the biased coordinate is *necessary but not sufficient* for mFR to help: it also requires that the starvation be a support deficit its whole-configuration reallocation can fill, rather than a barrier-crossing bottleneck it can only amplify. In every molecular collective variable we tried — easy torsion, hard distance, two-dimensional torsion — at least one of these preconditions fails, and the matched-seed verdicts reported above follow accordingly. The diagnostic oracle target and the frozen-bias control are reported alongside to separate the accuracy of the learned bias from the shared online mean-force resolution floor, so that a residual online  $L_2$  common to every method is not mistaken for a systematic error mFR could have removed.

## 14 Atomistic neutrality controls: two dipeptides

The united-atom alkanes of Sections 9 and 10 are deliberately simple, and the limitation this study carried longest was that its molecular negative rested on a single system with a single kind of coordinate. This section adds two all-atom peptides. Both were selected in advance as candidates where mFR ought to help, both cleared every screening gate that precedes the decisive one, and both then failed it in the same way: *ABF establishes the populations on its own*. They are reported because a screen that only reports the systems that passed is not a screen.

The two are not redundant. Alanine dipeptide is the standard benchmark and was chosen for familiarity; valine dipeptide was chosen for a specific reason — a side-chain rotamer barrier that ought to have produced a persistently under-populated state — and it is the stronger negative precisely because that reason turned out to be measurable, and wrong.

### 14.1 Screening gates, and the one that decides

Both peptides were run through the same pre-registered sequence. A state map locates the metastable states in the full torsional space. A *distinguishability* gate then asks whether the selected two-dimensional CV can resolve those states at all: marginal mFR sees exactly  $p(\xi)$ , so if two states project onto overlapping regions of the CV, a deficit in one is not a resolvable feature of  $p(\xi)$  and no marginal reallocation could preferentially clone into it. A coarse umbrella/MBAR

reference with the omitted coordinate left *free* supplies the state populations. Only then does the decisive gate run.

That gate asks for a state that is **discovered early and stays under-populated**. It requires, for each state  $k$ , the persistence-based hitting time  $T_{\text{hit},k}$  and the establishment time  $T_{\text{est},k}$  measured against the *bias-aware* target

$$Q_k^*(t) = \frac{\int_{C_k} e^{-\beta(F_{\text{ref}}(z) - B_t(z))} dz}{\int e^{-\beta(F_{\text{ref}}(z) - B_t(z))} dz}, \quad (36)$$

normalised over the reference’s support with the observed populations conditioned the same way. Scoring against the *unbiased* populations instead would report a state as starved exactly when ABF has correctly flattened it — manufacturing the very signal mFR is meant to remove, from a run in which nothing is wrong. Three outcomes are possible and only one licenses going further: an important state not found early is *discovery-limited*, the  $R_{15}$  regime where mFR cannot act; every state found *and* established means *ABF is already sufficient*; a discovered but persistently under-populated state is the regime mFR addresses. Both peptides returned the middle verdict.

## 14.2 Alanine dipeptide: the mechanism ran, and changed nothing

Vacuum Ace-Ala-Nme (22 atoms, AMBER ff14SB, 300 K, BAOAB at  $dt = 1$  fs with  $\gamma = 1$  ps<sup>-1</sup>, no constraints and no hydrogen-mass repartitioning) on the standard CV  $\xi = (\varphi, \psi)$ . The reference is umbrella sampling with MBAR over 576 windows and 16 independently thermalised copies each, 9.216  $\mu$ s aggregate; a copy-level block bootstrap with MBAR re-solved on every replicate gives  $\Delta G(\varphi > 0) = 3.419 \pm 0.079 k_B T$  and a median FES standard error of  $0.244 k_B T$ , with a further  $\approx 0.25 k_B T$  inter-sampler systematic carried explicitly from an independent OPES cross-check.

The reference settles the question before the sampler runs. The C7eq $\leftrightarrow$ C7ax barrier is a genuine  $15.8 k_B T$ , so the system is metastable and not trivially easy — but the internal barrier of  $F(\psi | \varphi)$  is at most  $0.75 k_B T$  at every populated  $\varphi$  (median 0.17). There is no hidden slow coordinate: two-dimensional ABF on  $(\varphi, \psi)$  biases the only slow directions that exist.

The oracle pilot then confirms the consequence with the mechanism verifiably running. Starting every walker in C7eq, plain ABF discovers C7ax at 3.1–4.1 ps in every seed and holds it at 5.6% occupancy over the 20 – –100 ps analysis window — against a reference equilibrium population of 0.0317. The rare basin is not merely reached; the transient ensemble *over*-populates it. Against that baseline, oracle mFR changes the kernel-matched integrated FES error by  $-0.010\%$   $[-0.063, -0.002]$  on 3/4 paired seeds, against a preregistered improvement bar of  $-10\%$  — a difference three orders of magnitude smaller than the threshold, and roughly four orders below the reference’s own uncertainty. Repeating from a reference-equilibrium initialisation instead of a pure C7eq start gives  $+0.031\%$ . Verdict: **equivalent**.

An FR-rate ladder spanning  $22\times$  in realised replacement events ( $3021 \rightarrow 66818$  over rates  $0.02 \rightarrow 0.45$ ; the ratio of the *rates* is  $22.5\times$  and of the per-opportunity intensities  $20.4\times$ , so the denominator is worth naming) leaves the result flat:  $-0.010\%$ ,  $-0.002\%$ ,  $+0.024\%$ . Over that same span the age-aware ancestor ESS falls from 0.966 to 0.602 — still above the 0.30  $N$  floor, with  $w_{\text{max}} \leq 0.0137$ . So the null is not inertia: reallocation is demonstrably doing work to the ensemble, the cost of that work is visible in the genealogy, and none of it reaches the free energy.

*Remark 10* (Two bugs that would each have produced this same verdict for the wrong reason). Both were caught, and both are worth recording because each would have printed the word *equivalent* while the study measured nothing.

First, the accepted reference stores  $+\infty$  in the 226 of 9409 cells no umbrella window ever visited. Centring it computes  $\infty - \infty$ , so the oracle target is NaN over the *entire* grid, no birth or death weight

is ever positive, and the oracle arm silently degenerates into plain ABF — reporting zero events and perfect equivalence with itself. It was caught only because a test asserted that birth–death actually fires.

Second, the basin partition originally ranked reference minima *by depth* and kept the deepest few. That returns six shallow sub-minima inside the  $\varphi < 0$  megabasin and *omits C7ax entirely*, which is seventh by depth and is the one basin the study is about. Prominence merging replaces it.

A third, in the reference analysis, is the most instructive: the “internal barrier” of  $F(\psi | \varphi)$  was computed as the column’s global maximum, i.e. the height of the sterically forbidden region rather than of any barrier a walker crosses. It reported a median of  $15.74 k_B T$  and the conclusion “ $\psi$  is a hidden slow coordinate” — the exact reverse of the truth, and the premise on which the entire study design rested. The corrected periodic min–max barrier gives the  $0.17 k_B T$  quoted above. The lesson these share is that a plausible summary statistic computed on the wrong domain fails silently and in the direction that confirms the design.

### 14.3 Valine dipeptide: selected for a barrier that was not there

Ac–Val–Nme has genuine coupled backbone–rotamer structure. A lattice-seeded state map over the whole  $(\varphi, \psi, \chi_1)$  torus finds seven metastable states which factor almost exactly as backbone megabasin  $\times \chi_1$  rotamer, and the selected CV  $\xi = (\varphi, \chi_1)$  recovers the full three-dimensional state label with 0.973 cross-validated balanced accuracy and a worst pairwise footprint overlap of 0.189, so a population deficit would be visible to a marginal score. The umbrella reference on  $(\varphi, \chi_1)$ , with  $\psi$  free and started from four values, resolves 8 regions whose populations span 0.42 down to 0.0014.

The decisive gate then fails outright. Over a 300 ps run with 16 seeds of 2048 walkers, ABF reaches every one of the 8 regions within 5.4 ps, holds them persistently from 18 ps, and establishes every one inside the  $\pm 50\%$  band by 52 ps — against a discovery threshold of 10% of the run. The worst relative deficit over the second half is 0.223 against a 0.50 threshold, and no region sits below half its target for more than 0.060 of the run against a 0.20 threshold. The rarest region, at an unbiased population of 0.0014, is nonetheless found by every seed and ends at 1.46 of its bias-aware target. ABF’s own free energy lands  $0.248 k_B T$  from the reference (seeds 0.237–0.257; marginal TV 0.069), so this is not a badly converged run flattering itself. Verdict: **FAIL-B (ABF already sufficient)**.

**The reason valine was chosen was measurable, and wrong.** The system was selected because a preliminary umbrella scan put the  $\chi_1$  rotamer barrier at  $11\text{--}18 k_B T$ , which would make the rotamers effectively uncrossable and, on that reading, would have made a  $\chi_1$ -hidden design discovery-limited. Direct measurement contradicts it:  $\chi_1$  rotamers interconvert at 2.70 changes per walker per ns, a rate that is flat from 30 to 300 ps and identical for walkers seeded in a well and on a barrier — an equilibrium rate, not a seeding transient. The  $11\text{--}18 k_B T$  figures are *backbone-clamped conditional* barriers, measured with  $\varphi$  and  $\psi$  restrained at  $\kappa = 500 \text{ kJ mol}^{-1} \text{ rad}^{-2}$ ; the clamp forbids the backbone relaxation that accompanies the rotation. With the backbone free the effective two-dimensional min–max path costs  $1.1\text{--}7.4 k_B T$ . The genuinely slow coordinate is  $\varphi$ , which crossed its megabasin four times in 2581 ns of aggregate unbiased sampling — and  $\varphi$  is in the CV. That is the whole explanation for the verdict, and it is the same explanation as for every other ABF-sufficient system in this study.

*Remark 11* (Conditional barriers are not the barrier a walker pays). Conditioning on the other coordinates can inflate a barrier severalfold: for valine the  $\chi_1$  barrier at fixed  $\varphi$  has median  $4.1 k_B T$




Figure 42: Ace-Val-Nme, the decisive gate. **Top left:** region occupancy (solid) against the bias-aware target  $Q_k^*(t)$  of Equation (36) (dashed), all eight regions over three decades of population. **Top right:** the same as a ratio; every region enters the establishment band and stays there. **Bottom left:** worst-case hitting and establishment times against the discovery threshold. **Bottom right:** ABF's own accuracy, with the concentrated and stratified initialisations giving the same answer — the diagnostic control doing its job, since the verdict is then not an artifact of where the walkers started.

and reaches 15.3, while the effective path with a free backbone costs 1.1–7.4. A restrained scan is the natural way to measure a torsional barrier and it systematically over-states the one that matters for sampling. This is worth stating separately because it is the kind of error that survives review: the measurement is correct, it is simply not a measurement of the quantity the design depends on.

#### 14.4 Three diagnostics that were wrong, and are worth not repeating

The valine screen produced three measurement errors. None can change its verdict — the decisive condition fails by a wide margin, and two of the three touch conditions that only ever gated a *pass* — but all three are properties of machinery meant to be reused, and each is a mistake that produces a confident wrong answer rather than an obvious one.

Table 29: Dipeptide closure. Entries are the paired relative change in the kernel-matched integrated FES error against ABF, so **negative is better**; the pre-declared threshold for calling an improvement was  $-10\%$ . Note the arm: only the *oracle* target was ever run on alanine, so this is the ideal-information control returning nothing, which is a stronger null than a practical arm would have been. Valine has no mFR row because its establishment gate never fired — ABF reaches all 8 regions by 5.4 ps and establishes all 8 by 52 ps of a 300 ps run — so no arm was justified and none was run. Regenerated by `scripts/build_closure_inventory.py`.

| Arm                                                               | $\Delta\text{err}$ (%) | 95 % CI            | seeds | Classification               |
|-------------------------------------------------------------------|------------------------|--------------------|-------|------------------------------|
| <i>Alanine dipeptide, oracle target, pilot stages</i>             |                        |                    |       |                              |
| $N=2048$ , C7eq                                                   | +0.0129                | [0.0117, 0.0296]   | 0/4   | equivalent                   |
| $N=4096$ , C7eq                                                   | $-0.0100$              | [-0.0632, -0.0024] | 3/4   | equivalent                   |
| $N=2048$ , reference-equilibrium start                            | +0.0313                | [-0.0965, 0.0407]  | 1/4   | equivalent (crossed control) |
| <i>Alanine dipeptide, FR-rate ladder (post-pilot sensitivity)</i> |                        |                    |       |                              |
| rate 0.02 (3 021 events, ESS/ $N$ 0.966)                          | $-0.0100$              | [-0.0632, -0.0024] | 3/4   | equivalent                   |
| rate 0.15 (22 830 events, ESS/ $N$ 0.813)                         | $-0.0025$              | [-0.0540, 0.0141]  | 2/4   | equivalent                   |
| rate 0.45 (66 818 events, ESS/ $N$ 0.602)                         | +0.0236                | [0.0085, 0.0352]   | 0/4   | equivalent                   |

**An unpaired comparison in paired clothing.** The omitted-coordinate check compared  $p_{\text{ABF}}(\psi \mid \text{region})$  against  $p_{\text{ref}}(\psi \mid \text{region})$  and reported a worst-region total variation of 0.280, apparently failing a 0.15 threshold. But ABF deliberately *flattens* within a region while the reference is Boltzmann-weighted inside it, so the two weight the region’s interior differently even when the  $\psi$  conditional at every fixed  $(\varphi, \chi_1)$  cell is identical. Comparing cell by cell and aggregating with *common* weights gives 0.055, and the condition passes. The signature was visible in the original numbers — the two largest, most internally varied regions disagreed while the six smaller ones did not — which is what a weighting artifact looks like and not what a genuine conditional error looks like.

**A counter that reads zero exactly when the barrier is real.** State entries were counted only between *consecutively labelled* frames. Reaching the  $\varphi > 0$  regions means crossing a corridor above the region ceiling, where frames are unlabelled, so the counter reported zero entries into  $B3, B6, B7$  — three regions with finite hitting times that were demonstrably reached. The failure mode is perverse: the higher the barrier, the more certainly the counter reads zero. Carrying the last labelled state across the corridor gives 1655/704/527 entries respectively.

**A target that concentrated where the reference knew nothing.** The first version of Equation (36) capped cells the reference never filled at  $F_{\min} + 30 k_B T$  and normalised over the whole torus. Because ABF flattens,  $B_t$  grows large exactly in those cells and  $e^{-\beta(F_{\text{capped}} - B_t)}$  explodes there: 97% of the target mass landed in the cells the reference was least able to speak about. The fix is to define the target on the reference’s labelled support and condition the observed fractions the same way. The diagnostic that caught it printed the offending number — and was noticed only because a human read the log. It is now an assertion, and the closure artifacts record  $|\sum_k Q_k^* - 1| < 5 \times 10^{-16}$  and the matching statement for the observed populations.



## 14.5 What the two neutrality controls establish

Table 29 collects every dipeptide arm that was run. Two features of it carry more weight than the numbers. First, the alanine arm is the *oracle* target, built from the reference free energy itself — the quantity the method is trying to estimate — so what the table reports is the ideal-information control returning nothing, across a  $22\times$  range of realised event rate over which the age-aware ancestor ESS falls from 0.966 to 0.602. The mechanism demonstrably worked harder and the accuracy did not move. That is a stronger null than a practical arm would have produced, because a practical arm’s null is always partly attributable to a noisy online target. Second, valine has no row at all: its establishment gate never fired, so no mFR arm was justified and none was run. Its classification therefore rests on the ABF-only screen and not on a measured mFR outcome, and we state that rather than presenting an absence as a null result.

Two independent all-atom systems, one of them selected specifically for a barrier that mFR should have been able to exploit, both return the same verdict: when ABF’s collective variable contains the slow coordinate, ABF discovers *and* establishes the populations by itself, and marginal reallocation has no deficit to repair. This is not evidence that mFR fails; it is evidence about which regime it belongs to, and it is what motivates constructing that regime deliberately in Section 15 rather than continuing to sample molecules.

It also sets a standing prohibition that the screening plans state explicitly and that we record here because the temptation is structural: both  $T_{\text{hit}} < 0.1T$  and “starved for at least  $0.2T$ ” scale with run length, so shortening the run, cutting walkers or widening the establishment band would manufacture a deficit rather than measure one. The margins here are not close — discovery is more than an order of magnitude faster than required and the worst deficit is well under half its threshold — and a negative obtained at this margin is worth more than a positive obtained by retuning.

## 15 Building the missing regime: an entropic gateway

Every system in this study so far has landed in one of two regimes, and neither of them tests marginal reallocation. The alkane torsions, the two-dimensional torus CV and the two dipeptides of Section 14 are *ABF-sufficient*: the slow coordinate is in the collective variable, ABF flattens it, and every state is populated long before the run ends, so no deficit ever forms. The  $R_{15}$  distance CV of Section 11 is *discovery-limited*: the state is never reached, and reallocation cannot clone a walker that does not exist. The regime mFR is actually supposed to serve — a state that is *found early* but *populated slowly* — has never appeared on its own.

The response is not to try more molecules. If the slow coordinate is included in the CV, ABF will flatten it again; if it is omitted, the system becomes discovery-limited or conditionally non-ergodic, where marginal reallocation provably cannot act. Either way the outcome is decided by the choice of coordinate rather than by the method, and searching for a molecule that happens to fall between the two is example hunting. We instead construct the missing regime deliberately, in a model where discovery and establishment can be moved independently and the free energy is known in closed form.

### 15.1 Model and screening design

We reuse the smooth entropic channel of Section 6,

$$V_{s,r}(x,y) = H(x^2 - 1)^2 + \frac{1}{2} \omega_{s,r}(x)^2 y^2, \quad \omega_{s,r}(x) = \omega_{\text{out}} + (\omega_{\text{in}} - \omega_{\text{out}})e^{-x^2/2s^2}, \quad \xi(x,y) = x, \quad (37)$$

with  $r = \omega_{\text{in}}/\omega_{\text{out}} > 1$  and the analytic free energy  $F_{s,r}(x) = H(x^2 - 1)^2 + \beta^{-1} \log \omega_{s,r}(x) + C$ . Because  $F$  is exact there is no reference simulation to converge and no reference error to confound the classification — the single largest advantage over any molecular alternative. The constriction at  $x = 0$  is a narrow entropic gateway: a walker can occasionally cross, repeated independent crossings stay slow, the far chamber is visible in  $\xi$ , and a discovered walker there is a valid sample that reallocation could multiply. That is precisely the mechanism mFR claims.

The screen is **ABF only** — no Fisher–Rao arm of any kind — over  $s \in \{0.10, 0.15, 0.20, 0.30\}$ ,  $r \in \{4, 8, 16, 32\}$  and  $\beta \in \{2, 4, 8, 16\}$ , at 16 seeds,  $N = 2048$  walkers and  $T = 40$ , in both a concentrated left-basin initialisation and a control in which one walker starts in the right basin. We hold  $\beta H = 8 k_B T$  fixed, so that  $\beta F(x) = \beta H(x^2 - 1)^2 + \log \omega(x)$  carries no residual  $\beta$  dependence: the *dimensionless landscape is identical in every cell*, the total barrier is  $8 + \log r \in [9.4, 11.5] k_B T$  throughout, and  $\beta$  varies only how fast walkers move through it, at an unchanged sampling budget.

*Remark 12* (The  $\beta$  axis is a time budget, not a family of landscapes). A sizing scan established that  $(s, r)$  alone does not span the regimes: at fixed  $\beta$  the gateway geometry moves the establishment gap only weakly, monotonically and in the expected direction. Reporting a single  $(s, r)$  slice would therefore have hidden the axis that does the work — but it is essential to be precise about what that axis *is*, and it is not a family of different physical problems.

Because  $\beta H$  is held fixed,  $\beta F(x)$  is identical in every cell by construction. Substituting  $\tilde{y} = \sqrt{\beta} y$  and  $\tilde{b} = \beta b$  and rescaling time by  $\tau = t/\beta$  gives

$$dx = -[G(x, \tilde{y}) - \tilde{b}(x)] d\tau + \sqrt{2} dW_\tau, \quad d\tilde{y} = -\beta \omega^2 \tilde{y} d\tau + \sqrt{2\beta} dW_\tau, \quad (38)$$

with  $G = 4\beta H x(x^2 - 1) + \omega \omega' \tilde{y}^2$ . The *longitudinal* equation — the one the collective variable lives on — is exactly  $\beta$ -free.  $\beta$  survives only as the transverse relaxation rate relative to longitudinal motion, and in the direction that makes large  $\beta$  *more* adiabatic: as  $\beta \rightarrow \infty$  the fast  $\tilde{y}$  averages,  $\langle \omega \omega' \tilde{y}^2 \rangle = (\log \omega)'$ , so  $\langle G \rangle = d(\beta F)/dx$  and the limiting dynamics has no  $\beta$  at all.

The measurement agrees. Across the  $8 \times$  ladder,  $T_{\text{est}}/T_{\text{run}}$  varies  $6.62 \times$  — which is what made it look like a regime axis — while  $\tau_{\text{est}}$  varies only  $1.21 \times$  and  $\tau_{\text{hit}}$  is identical to four digits (0.100). The residual drifts downward with  $\beta$ , the sign predicted by finite- $\beta$  non-adiabaticity.

**We therefore report this table as a finite-budget establishment map**, not as a phase diagram over physically different equilibrium problems. The collapse is not exact, and what breaks it is worth naming: the ABF estimator does *not* rescale — counts accumulate per step, so each cell receives  $\beta$  times more samples per unit  $\tau$ , and the time-averaged error still varies *nan* $\times$  — followed by finite- $\beta$  non-adiabaticity and the FR event frequency per unit  $\tau$ . Bandwidths live in  $x$  and are invariant.

None of this weakens the comparison that follows. Methods are always evaluated at finite budget, and a regime defined by a ratio of timescales can only be reached by moving that ratio. What it forbids is the claim that varying  $\beta$  produced different free-energy landscapes; it did not, and Equation (38) is why.

With  $T_{\text{hit}}$  the first time  $\hat{P}_t(B_+) > 0$  persistently and  $T_{\text{est}}$  the first time  $\frac{1}{2} Q_+^*(t) \leq \hat{P}_t(B_+) \leq \frac{3}{2} Q_+^*(t)$  persistently — both on a trailing window of 5% of the run, against the *bias-aware* target  $Q_+^*(t) = \int_{B_+} e^{-\beta(F-\hat{F}_t)} / \int e^{-\beta(F-\hat{F}_t)}$  — each cell is labelled by a rule fixed before any of it ran: *discovery-limited* if  $T_{\text{hit}} \geq 0.1 T$  in at least a quarter of seeds; otherwise *establishment-limited* if  $T_{\text{est}} - T_{\text{hit}} > 0.25 T$  or the state is below half its target for at least  $0.2 T$ ; otherwise *ABF-sufficient* if  $T_{\text{est}} - T_{\text{hit}} < 0.1 T$ ; otherwise *intermediate*. Discovery is tested first because a state that is not reliably found cannot be judged on how fast it fills.

Gateway  $\beta$  axis:  $\beta F(x)$  is identical in every cell, and the longitudinal SDE is  $\beta$ -free under  $\tau = t/\beta$ ,  $\tilde{y} = \sqrt{\beta} y$




Figure 43: The  $\beta$  axis collapses under the rescaling of Equation (38). **Left:** as originally reported, the establishment time as a fraction of the run varies  $6.62\times$  across the ladder, which reads like a regime axis. **Middle:** in rescaled time  $\tau = t/\beta$  it varies  $1.21\times$  against an  $8\times$  change in  $\beta$ , and  $\tau_{\text{hit}}$  is identical to four digits. **Right:** the same points against the available budget  $\tau_{\text{max}} = T/\beta$ , which is what the ladder actually varies. Small points are individual runs, large points the median over the 16 seeds and 16  $(s, r)$  cells at that  $\beta$ . The residual drift is downward in  $\beta$ , the sign predicted by finite- $\beta$  non-adiabaticity of the transverse channel.

## 15.2 The finite-budget establishment map

Figure 44 reports all 128 cells, including every cell where mFR could not possibly help. Of these, 32 are ABF-sufficient, 42 intermediate, 54 establishment-limited and 0 discovery-limited. The map is layered in  $\beta$ :  $\beta = 2$  is ABF-sufficient everywhere,  $\beta = 4$  intermediate everywhere,  $\beta = 8$  mixed, and  $\beta = 16$  establishment-limited everywhere. Per Remark 12 that layering is a statement about how much rescaled time a fixed step budget buys, not about different landscapes. No cell of this model is discovery-limited; that regime is represented in this study by  $R_{15}$ , not here.

## 15.3 The anchor, chosen before any Fisher–Rao arm existed

The classification was written to disk, hashed and committed *before* a single Fisher–Rao arm was run on this system, and the anchor follows from a rule stated in that same artifact: the headline initialisation; regime establishment-limited; *interior, not knife-edge*, meaning every neighbour along  $s$ ,  $r$  and  $\beta$  present in the map shares its regime; then the largest median integrated deficit among the survivors. This selects

$$\beta = 16, \quad s = 0.1, \quad r = 32.$$

Under ABF alone the right chamber there is reached at  $T_{\text{hit}}/T = 0.040$  by *every* seed, reaches its establishment band only at  $T_{\text{est}}/T = 0.465$ , sits below half its bias-aware target for 0.423 of the run, and ends at 0.828 of that target. A state found in the first 4% of the run and still short at the end is exactly the deficit mFR claims to repair. The  $\beta = 8$  cells were rejected by the same rule despite lying nearer the regime boundary: their gap of  $\approx 0.21$  falls below the 0.25 threshold and they qualify only through the below-half criterion at  $\approx 0.20$  against a 0.20 threshold — the knife-edge the rule exists to exclude.

## 15.4 Four arms, and what the sham controls

Every arm inside a seed shares initial conditions and Langevin noise, so all comparisons are paired. Alongside ABF, oracle mFR (target from the analytic  $F$ ; non-deployable) and practical mFR (target




Figure 44: ABF-only regime map of the entropic gateway, all 128 cells. Rows are the two initialisations, columns the transport ladder  $\beta$ ; each panel is the  $(s, r)$  plane and each tile is annotated with the establishment gap  $(T_{\text{est}} - T_{\text{hit}})/T$ . The dimensionless landscape is identical in every tile ( $\beta H = 8 k_B T$  fixed), so the layering in  $\beta$  is a statement about timescales, not about thermodynamics. The black box marks the preregistered mFR anchor. Within a fixed  $\beta$  the gateway geometry modulates the gap only at the margin — a result, not a nuisance, and the reason the whole map is shown rather than one slice.

from the online EMA of the bias), we run a **matched sham**: it fires at the same event times and performs the *same realised number* of clone and delete events as the oracle arm — copied from its partner row, not re-derived from its own diverged state — and differs only in drawing the identities uniformly at random. Across 128 paired comparisons the counts matched with 0 mismatches.

This control has not existed anywhere in this study before. Every previous arm (`fr_estimated`, `fr_uniform`, `fr_oracle`, `fr_active`) derives its replacement direction from the same Fisher–Rao score and differs only in how the target is built, so the proposition “the gain comes from turnover rather than from Fisher–Rao steering” had never been tested against an arm that lacks the direction.

## 15.5 A calibration run, then a confirmatory one

The first pass at this anchor selected each arm’s FR rate from a four-point ladder *and* reported that rate’s confidence interval on the same runs. That is tuning-selection bias: the interval is conditioned on having chosen the best-looking rung. We treat that pass as calibration only. It fixed  $\gamma_{\text{practical}} = 1.5$  and  $\gamma_{\text{oracle}} = 0.5$  — the gentlest rates that pass the preregistered health gates ( $\text{ESS}_{\text{anc}}/N \geq 0.30$ ,  $w_{\text{max}} \leq 0.05$ ) — and those, together with every other knob, the primary metric and the success rule, were frozen in a committed preregistration before the confirmatory run began.

The confirmatory run uses 32 fresh seeds, disjoint from the calibration seeds, with no ladder and no tuning. Its success rule, fixed in advance, required the practical arm to reach a median  $\leq -10\%$  with a 95% CI upper endpoint below  $-5\%$  on at least 24 of 32 seeds, with both health gates passing.

Table 30: **Confirmatory** run at the preregistered establishment-limited anchor ( $\beta = 16$ ,  $s = 0.1$ ,  $r = 32$ ): 32 fresh seeds, rates frozen from a separate calibration pass, no ladder and no tuning. Percentages are paired differences against ABF on the same seed with a bootstrap CI over seeds; **negative is better**. Each sham fires at its own partner’s event times with its partner’s realised clone/delete counts, and differs only in randomising which walkers they act on. The *direct* rows are the attribution statistic: each FR arm against its own sham on the same seed, which holds the schedule and the count fixed so that only the selection direction differs. Health gates  $\text{ESS}_{\text{anc}}/N \geq 0.30$  and  $w_{\text{max}} \leq 0.05$  pass for every arm shown.

| comparison                                                                           | $\gamma$ | $\Delta I_F$   | 95% CI                  | seeds won | frozen-bias $\Delta$ |
|--------------------------------------------------------------------------------------|----------|----------------|-------------------------|-----------|----------------------|
| <i>against ABF</i>                                                                   |          |                |                         |           |                      |
| practical mFR                                                                        | 1.5      | <b>-12.12%</b> | $[-14.50, -8.41]$       | 31/32     | -12.40%              |
| sham (practical)                                                                     | 1.5      | +0.35%         | $[-1.66, 2.62]^\dagger$ | —         | -4.38%               |
| <i>direct: each arm against its own sham — the attribution test</i>                  |          |                |                         |           |                      |
| practical mFR vs sham                                                                | 1.5      | <b>-11.79%</b> | $[-14.10, -8.19]$       | 31/32     | -6.42%               |
| oracle mFR vs sham                                                                   | 0.5      | -8.89%         | $[-13.72, -5.60]$       | —         | —                    |
| <i>mechanism control: one walker seeded across the gateway, so discovery is free</i> |          |                |                         |           |                      |
| practical mFR                                                                        | 1.5      | -10.75%        | $[-16.24, -8.08]$       | 29/32     | —                    |

<sup>†</sup> 90 % CI, the interval for the two-one-sided-tests equivalence check against a  $\pm 5\%$  margin; the practical sham is equivalent to ABF on this endpoint. The oracle arm’s sham is not shown equivalent (90 % CI  $[-5.63, 1.05]$ ), so only the practical arm supports an equivalence claim.

Table 30 gives the confirmatory result. Practical mFR improves the integrated free-energy error by  $-12.12\%$   $[-14.50, -8.41]$  on 31/32 matched seeds, at  $\text{ESS}_{\text{anc}}/N = 0.399$  and  $w_{\text{max}} = 0.012$ : every criterion passes. The deficit the anchor was selected for is largely repaired, and  $T_{\text{est}}$  falls by  $-76.9\%$ .

**The attribution statistic.** Comparing each arm against ABF and then eyeballing the two contrasts is weaker than comparing an arm against *its own sham* directly on the same seed, which holds the event schedule and the replacement count fixed by construction so that only the selection direction differs. That contrast gives

practical mFR vs its sham:  $-11.79\%$   $[-14.10, -8.19]$ , 31/32 seeds,      oracle vs its sham:  $-8.89\%$   $[-13.72, -5.60]$ .

Both intervals exclude zero. Separately, the practical sham is **equivalent** to ABF by two one-sided tests against a  $\pm 5\%$  margin ( $+0.35\%$ , 90% CI  $[-1.66, 2.62]$ ).

*Remark 13* (Two statistical points, and two places this study got them wrong). A *directional label needs an interval that excludes zero*. The calibration analysis labelled a sham cell an “improvement” on a median of  $-1.1\%$  with an interval of  $[-11.1, +6.4]$ , purely because the interval was wide. That reports the imprecision of the estimate rather than the size of the effect; a directional call now requires the interval to exclude zero, and anything else is equivalent or inconclusive.

“*No interval excluded zero*” is not equivalence. Failure to detect an effect is not evidence that the effect is negligible, and the calibration write-up leaned on exactly that phrasing. Equivalence is now a TOST against a predeclared margin. It matters: under the correct test the *oracle* arm’s sham is not shown equivalent (90% CI  $[-5.63, 1.05]$ , whose lower end just clears the margin), so only the practical arm’s sham supports an equivalence claim. The headline rests on the practical arm, and the oracle’s is reported as inconclusive.




Figure 45: The anchor comparison. **Top left:** right-chamber occupancy against the bias-aware target; ABF and the sham track each other while both mFR arms reach the target far earlier. **Top right:** free-energy error. **Bottom left:** the FR-rate ladder with bootstrap CIs over matched seeds, against the  $\pm 10\%$  practical-equivalence margin — the sham straddles zero throughout. **Bottom right:** the cost side. Gains keep growing with rate, but ancestor diversity collapses through the preregistered floor, which is why the headline is quoted at a health-passing rate and not at the largest number.

*Remark 14* (A second replicate, and where the preregistered rule turned out to be blunt). The confirmatory run was executed twice. The first pass carried a pairing defect found by code inspection: the two FR rates were run as two batches, the Langevin stream is keyed to the batch, and the oracle arms were therefore scored against a baseline from the other batch (matched initial conditions, unmatched noise). The primary claim and both direct contrasts lived inside a single batch and were unaffected. The fix was structural — the FR rate now rides on the method, so all five arms share one batch — which necessarily re-runs every arm, so the second pass is reported as a *replicate* rather than a correction, and the first is neither withdrawn nor overwritten.

**The primary claim replicates.** Across two independent noise realisations, practical mFR gives  $-12.12\%$   $[-14.50, -8.41]$  on 31/32 seeds and  $-12.48\%$   $[-16.30, -9.84]$  on 31/32; the estimator-independent frozen-bias endpoint gives  $-12.40\%$  and  $-10.19\%$ . The direct arm-versus-its-own-sham contrast gives  $-11.79\%$  and  $-14.95\%$   $[-16.22, -12.34]$  on 32/32 seeds, excluding zero in both.

Amendment 2 fixed in advance what to do if both passed: *quote the replicate, cite both*. Both do pass the primary rule (PASS and PASS), so the quoted gateway figure throughout this report is the replicate's  $-12.48\%$ , with the first pass's  $-12.12\%$  cited beside it. That the two agree to

within four tenths of a percentage point is the substantive point; which one is nominated as the headline is a bookkeeping rule that was written down before either number existed, and following it is cheaper than defending a choice made afterwards.

**The sham equivalence test does not replicate, and the rule is at fault.** The practical sham reads +0.35% in the first pass (90 % CI  $[-1.66, 2.62]$ , equivalent) and +1.91% in the second (90 % CI  $[-1.87, 5.40]$ , not shown equivalent — it misses the +5% bound by about four tenths of a percentage point). Both point estimates are *positive*: the sham is never better than ABF, and in the replicate it is slightly worse.

The preregistered headline rule required the primary criteria *and* sham equivalence, so by its letter the replicate does not support the headline. We report that. But the rule is blunt in a way worth naming rather than working around: it does not distinguish the direction in which equivalence fails, and a sham that makes the estimate worse cannot be the explanation for an arm that makes it better. On the frozen-bias endpoint the same sham is equivalent (+1.14%). The statistic that actually isolates direction is the direct contrast, which is decisive in both passes — and which, for this system, was added during analysis of the first pass and is therefore *post hoc* here. It was preregistered for the WCA comparison (Section 7.15), where it also passes, and that is the stronger claim of the two because nothing about it was chosen after the fact.

## 15.6 Three things the design was built to rule out

**It is establishment, not discovery.** In the mechanism control one walker starts across the gateway, so discovery is handed to the sampler for free and any remaining acceleration must be population establishment. Practical mFR still gives  $-10.75\%$   $[-16.24, -8.08]$  on 29/32 seeds with the health gates passing. The effect survives when first passage is removed.

**It is a better bias, not a better-behaved estimator.**  $I_F$  is computed from the same accumulators the FR mechanism perturbs, and birth–death makes replicas correlated descendants, so a gain measured that way could in principle be a change in the estimator’s statistics rather than a genuinely better bias. We therefore freeze each arm’s learned mean force at the same physical time, launch a *fresh* population — identical across arms, so no arm inherits an advantage from where its walkers ended up — with no adaptation and no birth–death, and reconstruct  $\hat{F} = B - \beta^{-1} \log \hat{p}_B$  from the sampled density. On that estimator-independent endpoint practical mFR gives  $-12.40\%$   $[-16.93, -7.48]$ : the same sign and essentially the same size as the online number. The bias itself is better.

This endpoint also refines the attribution rather than merely confirming it. The practical sham reads  $-4.38\%$   $[-11.93, -1.38]$  here — small, but *not* equivalent within the margin, unlike on the online endpoint. Uniform turnover does improve the learned bias a little. The direct arm-versus-sham contrast on the same endpoint is  $-6.42\%$   $[-11.49, -1.30]$ , still excluding zero. The honest summary is therefore a split rather than a clean dichotomy: generic turnover contributes a minority of the benefit, and the Fisher–Rao direction contributes the majority.

**The headline is not the largest available number.** At the rate  $\gamma = 15$  inherited from the accepted entropic-bottleneck setup, the calibration gains are considerably larger —  $-33.51\%$  for the oracle and  $-30.16\%$  for the practical arm — but *both health gates fail* ( $\text{ESS}_{\text{anc}}/N = 0.017$  and  $0.034$ ,  $w_{\text{max}} = 0.134$ ). The confirmatory rates are the gentlest that pass, not the best-performing. This is the same accuracy–diversity boundary the WCA dimer exhibits (Section 7), reached here from the other side.

## 15.7 Why the regime exists at all: a two-state reduction

The separation the gateway was built to create is not special to this potential. Coarse-grain the reaction coordinate into “not in the rare state” and “in it”, let  $p(t)$  be the population fraction of the rare state and  $q$  its target mass, and write the mean-field kinetics as

$$\dot{p} = k_{01}(1 - p) - k_{10}p, \quad p(\infty) = \frac{k_{01}}{k_{01} + k_{10}} = q. \quad (39)$$

Two timescales follow, and they are not the same one.

*Discovery* is a first-passage event among  $N$  independent walkers, so it happens on

$$T_{\text{hit}} \sim \frac{1}{Nk_{01}}, \quad (40)$$

which is fast simply because  $N$  is large: a rare crossing becomes likely once you have enough tries. *Establishment* is the relaxation of Equation (39) toward  $q$ , whose rate is the eigenvalue  $k_{01} + k_{10}$  and which therefore takes

$$T_{\text{est}} \sim \frac{1}{k_{01} + k_{10}}, \quad (41)$$

with *no factor of  $N$* . Adding walkers buys discovery and does not buy establishment. The ratio  $T_{\text{est}}/T_{\text{hit}} \sim Nk_{01}/(k_{01} + k_{10})$  grows with  $N$ , so the establishment-limited window is not a curiosity of this model — it is what a many-walker sampler generically produces once  $k_{01}$  is small, which is exactly the condition under which the coordinate was worth biasing in the first place.

Marginal Fisher–Rao reallocation adds a term whose local linear rate near  $q$  is set by the resampling intensity  $\gamma$  rather than by the physical crossing rate:

$$\dot{p} = \underbrace{k_{01}(1 - p) - k_{10}p}_{\text{physical transport}} + \underbrace{\gamma p \log(q/p) \big|_{\text{lin}}}_{\text{reallocation}} \approx -\gamma(p - q), \quad (42)$$

so the relaxation rate becomes  $k_{01} + k_{10} + \gamma$  while the fixed point stays at  $q$  provided the target is correct. That single expression accounts for all four outcomes in this study:

- **ABF-sufficient** (alanine, valine, the alkanes):  $k_{01} + k_{10}$  is already large compared with  $1/T$ , so  $\gamma$  changes a relaxation that had already finished. Adding it can only inject resampling noise — which is what the neutral verdicts, and the mild harm at high rate, look like.
- **Establishment-limited** (the gateway, and Section 7):  $Nk_{01} \gg 1/T \gg k_{01} + k_{10}$ . Discovery is done, relaxation is not, and  $\gamma$  supplies the missing rate. This is the regime, and Equation (42) says it exists whenever the two inequalities can be satisfied together.
- **Discovery-limited** ( $R_{15}$ ):  $Nk_{01} \lesssim 1/T$ , so  $p(t) = 0$  over the whole run. The reallocation term is proportional to  $p$  and vanishes identically — mFR cannot create a representative of a state that has none. No choice of  $\gamma$  changes this, which is why the  $R_{15}$  rate ladder is monotone in the wrong direction.
- **The cost**, and why  $\gamma$  cannot simply be raised: the term in Equation (42) moves *copies*, so the ancestor diversity decays at a rate that also scales with  $\gamma$ . Accuracy improves with  $\gamma$  while the effective sample size the mean-force estimator consumes decays with it, which is the accuracy–diversity boundary seen in both the WCA rate ladder and this study’s own health-gate failures at  $\gamma = 15$ .



The reduction is deliberately crude — it assumes a single rare state, mean-field kinetics and a correct target — but it is enough to say why the three regimes are three and not a continuum, and why the boundary between them moves with  $N$  and with  $T$  rather than with the shape of the potential.

*Remark 15* (The deployable target captures the gain here). In the entropy-dominant stress test (Section 6.5) the oracle target diverged sharply from the online estimate, and the deployable arm captured essentially none of the oracle’s benefit. At this anchor the two coincide: the direct-versus-sham contrasts are  $-11.79\%$  for the practical arm and  $-8.89\%$  for the oracle. In terms of Equation (42), both targets place the fixed point close enough to  $q$  that what matters is  $\gamma$ , not the residual target error; the oracle buys shape information the reallocation here does not need. Target quality matters in proportion to how much shape information the reallocation actually requires, which is small when the instruction is “move mass across the gateway” and large when it is “reproduce an entropy-dominated profile”.

## 16 A literature-anchored negative: methane association in water

The gateway of Section 15 was constructed to contain the regime marginal reallocation is supposed to serve. That construction is its strength as a mechanism test and its weakness as evidence: a regime built to order says little about how often it occurs. The obvious complement is a molecular system chosen for the opposite reason — one where the literature already describes the physics, where nothing about the setup was arranged by us, and where the honest expectation is that mFR has no room to act.

Methane–methane association in explicit water is that system. Its potential of mean force is a textbook case of hydrophobic association, with a contact pair, a desolvation barrier and a solvent-separated pair; the reaction coordinate is a single interparticle distance, so nothing is hidden by a poor choice of collective variable; and the orthogonal degrees of freedom are real water molecules that must rearrange for the pair to separate. It was entered in the preregistration as a *likely-null falsification benchmark* before any methane code existed, with the stop condition written in advance: if ABF is sufficient, report that as the result and do not extend the budget until it stops being true.

### 16.1 Model, engine and reference

The physical model follows Asthagiri, Merchant and Pratt: united-atom methane ( $\sigma_M = 3.73 \text{ \AA}$ ,  $\epsilon_M = 0.294 \text{ kcal/mol}$ ) in 512 rigid SPC/E waters at 298 K, particle-mesh Ewald, a 1.05 nm cutoff, and a 0.5 fs BAOAB timestep. The methane–oxygen unlike-pair rule is Lorentz–Berthelot and is declared as our implementation choice rather than as a value read from that paper. NPT equilibration fixes the production box at  $L = 2.490832 \pm 0.000497 \text{ nm}$ , giving a density of  $0.9946 \text{ g/cm}^3$  against SPC/E’s accepted  $\sim 0.998$ ; all production is NVT at that volume, so the target measure is exactly the canonical measure of ??.

The sampler is a batched periodic engine written for this study, since no existing code runs a population of walkers sharing one live ABF estimator. It agrees with OpenMM to  $2.9 \times 10^{-13}$  in energy and  $1.4 \times 10^{-15}$  in force over configurations spanning  $-2598$  to  $+45\,223 \text{ kJ/mol}$ , its constrained BAOAB reproduces OpenMM’s thermostat to 0.51 K, and rigid-water constraints hold to  $9.3 \times 10^{-15} \text{ nm}$ .

The reference is constrained thermodynamic integration — the same local mean force  $f = \frac{1}{2}(\mathcal{F}_1 - \mathcal{F}_2) \cdot e - 2/(\beta r)$  that ABF itself estimates — over 3 independent builds, each  $r$  sampled

from deliberately opposed *wet* and *dry* solvent preparations. Builds agree to 0.092 kJ/mol against a 3.11 kT span, a ratio of 0.1198 against an acceptance threshold of 0.5.

## 16.2 The landscape is real, and it is sub- $k_B T$

Our reference reproduces the literature structure at the expected positions: a contact minimum at 0.38 nm, a desolvation barrier at 0.56 nm and a solvent-separated minimum at 0.72 nm. The depths are the finding. Measured on  $W(r)$  they are  $-1.11$ ,  $+0.40$  and  $-0.27 k_B T$ , so the barrier from the higher of the two minima is  $0.67 k_B T$  — below the  $2 k_B T$  merge threshold, which collapses the landscape to a *single basin* and fires the preregistered tercile fallback for state definitions.

This was recorded before the screen ran. There is no metastability in  $\xi$  for ABF to be slow about, and the prediction that follows — that methane will be ABF-sufficient — was written down at the reference stage precisely so that it could not afterwards be presented as a confirmed hypothesis.

## 16.3 Gate 0: the baseline is sound

Before asking what limits ABF, ask whether its conditional mean force can be trusted at all. The instrument is the controlled experiment of Section 17: independently prepared solvent families held at the same  $r$ , either agreeing or not. Methane’s cross-family spread is 0.048 of  $|\overline{F'}|$  globally and 0.081 through the desolvation region, against WCA’s 0.040 and deca-alanine’s failing 0.61. The spread also *collapsed* roughly threefold between the 50 ps and 200 ps checkpoints, which is the signature of a conditional ensemble that mixes; deca’s error persisted at up to  $2 \times 10^6$  effective counts per bin.

Methane is therefore not conditional-equilibration-limited. Whatever the screen finds is a statement about the free-energy landscape rather than about a broken baseline.

## 16.4 The screen, and what each gate can and cannot support

Eight independent ensembles of 512 walkers ran 200 ps of ABF with no Fisher–Rao arm of any kind. All four gates pass and the classification is *ABF-sufficient* on 8 of 8 seeds, unanimously.

The result rests on Gate C, and the reasons the other gates carry less weight are worth stating rather than hiding. Occupancy of every tercile tracks the bias-aware target  $Q_k^*(t) \propto \int \exp[-\beta(F_{\text{ref}} - B_t)]$  throughout the second half of every run: over all 9 624 second-half state-checkpoints the worst ratio is 0.731 and the best is 1.09, against a deficit threshold of 0.5. No partition choice can manufacture that. Walkers never pin, with a maximum single-tercile occupancy of 0.516 against a 0.90 clause and against deca’s 0.951–0.9996.

Because the verdict rests on this gate, it is worth asking what deficit the gate could have detected. Gate C compares an *integer walker count* against  $0.5 Q_k^*$ , so its power is set by the expected count  $\lambda_k = Q_k^* N$  and not by the run length or the number of checkpoints. Once  $0.5\lambda < 1$  the only integer satisfying the test is zero, and the gate silently stops measuring a shortfall and starts measuring whether the state happens to be empty at that instant — an event with frequency  $e^{-\lambda}$  on the physics alone, rising as  $N$  falls. We therefore require  $\lambda_k \geq 16$ , the value at which a 50 % deficit is a two-sigma effect on a single checkpoint, and report any state below it as non-binding with its detectable effect size rather than reading its output. Methane’s three terciles hold  $\lambda = 127.6$ , 147.0 and 224.2 at the worst checkpoint of the judged window, so the requirement is inert here by measurement rather than by assumption, and counting noise on those states would separate a 50 % shortfall at 13–18 %. That last figure, however, describes counting noise and not the gate, and we report it only to retire it. Gate C requires a *contiguous* run below  $0.5 Q_k^*$ , and contiguity

suppresses true firing as well as false; since every methane state cleared the gate, nothing in the study had shown it firing at all. We closed that by planting walker-conserving deficits in the real traces — scaling one state’s occupancy and redistributing the removed population so the partition still sums to one — and running the unmodified gate. It fires on 0 of 8 seeds at a planted 50 % deficit, the shortfall it nominally tests for, and on 8 of 8 at 60 %. Subsampling walkers to read the threshold across  $\lambda$  gives 60–65 % over a fifty-fold range, so above  $\lambda \approx 9$  detection is set by the span requirement and not by  $\lambda$  at all. Two things follow. The requirement  $\lambda \geq 16$  marks where counting noise stops dominating rather than where a 50 % deficit becomes detectable, which is not what we first claimed for it; and “Gate C did not fire” licenses only that no deficit of 60 % or more occurred. The statement that carries the null is the direct one — the worst occupancy ratio anywhere is 0.731, bounding the largest shortfall at 26.9 %, some two and a fifth times tighter than the gate could report. The planted deficit is stationary, so 60 % is a floor on the detection threshold rather than a characterisation of it.

It is not inert in general: in the companion NaCl study the contact-pair state holds  $\lambda = 1.72$  at  $N = 64$ , where the smallest resolvable deficit is 152 %, and at  $N = 16$  and  $N = 8$  no state clears the bar at all. Since the preregistered rule selects the *smallest*  $N$  passing every gate, the search runs directly toward the cells in which the gate cannot fail — and this project has already retracted one screen in which Gate C fired on a state that could never hold a walker and reported that a Fisher–Rao arm was licensed.

Gate B, by contrast, is *vacuous here, and that is a property of the physics*. First arrivals at thresholds well beyond the tercile boundary run only 1.9–5.8 $\times$  the exact ballistic floor for the fastest of 512 walkers, and the whole evaluation domain is 0.51 nm — about 3 ballistic steps for the leading walker. There is no separation between crossing a barrier and flying across the box, so no discovery *rate* is resolvable at any budget. The companion NaCl study runs the identical gate on a 1.20 nm domain and measures 6.9–20.7 $\times$  its floor: a resolvable, unambiguously diffusive rate. The same gate, run the same way, is informative in one system and vacuous in the other, and the discriminator is domain width in ballistic steps rather than anything about the method.

Gate A passes at 0.935 against a 0.30 threshold on the preregistered quantity,  $\text{TV}(p(\xi \mid Y = a), p(\xi \mid Y = b))$  with  $Y$  the inter-methane water count. The collective variable does see the solvent physics. We note for the record that our first implementation computed the transpose,  $\text{TV}(p(Y \mid \xi))$ , which reads 0.987; that is a weaker and partly tautological statement, since the gap volume grows with  $r$  by construction, and it is not the number that licenses a marginal method to act.

A high total variation could still be near-tautological if the descriptor were simply a function of the coordinate, so we measured that rather than assuming it. Across 0.01 nm bins,  $n_{\text{gap}}$  varies 5.4 $\times$  more between bins than within them, retaining a standard deviation of 0.37 at fixed  $r$ . The descriptor is substantially but not wholly determined by the coordinate, so Gate A here is a genuine measurement. The companion NaCl study measures 14–83 $\times$  on its hydration descriptors and a total variation of exactly 1.000, which its own analysis correctly reports as a statement about the physics — ion hydration is nearly a function of the ion separation — rather than as a strong gate. Methane retains real solvent structure orthogonal to  $r$ , which is what makes its null informative: there was something for reallocation to act on, and it was still unnecessary.

## 16.5 Finite size: a check that cannot resolve what it tests

The preregistration required a finite-size check — the constrained-TI mean force at  $r = 0.70, 0.80, 0.90$  nm recomputed in a 1024-water box, with the evaluation domain to be truncated if any point moved by more than  $0.1 k_B T / \text{nm}$ . It was specified to run *before* the reference

was built and was in fact run after it, so it can no longer gate the domain; it can only report whether the outer points of an accepted reference are trustworthy. Retrospectively performing a preregistered check is a weaker thing than fulfilling it, and we say which one happened.

In the larger box ( $L = 3.134273$  nm,  $\rho = 0.9966$  g/cm<sup>3</sup>, where  $r = 0.90$  sits at 57% of the minimum-image half-box rather than 72%) the differences are +0.0013, +0.1060 and +1.1563  $k_B T$ /nm. Read literally the gate fails at the outer two points. Read with the uncertainties —  $\pm 0.670$ ,  $\pm 0.500$ ,  $\pm 0.671$  — *nothing is resolved*: the nominal failure at 0.80 is  $0.21\sigma$  and the one at 0.90 is  $1.72\sigma$ .

The tolerance is five to seven times finer than the uncertainty of the measurement it gates, and closing that gap needs two orders of magnitude more sampling than the check itself cost. The criterion was fixed in advance without asking whether it was measurable — a check whose outcome cannot entail the thing it is relied on to establish. We record the gate as **inconclusive** rather than failed, treat  $r = 0.90$  as suspect, and confirm the interior: at  $r = 0.70$  the two box sizes agree to  $0.0013 k_B T$ /nm, and every state boundary and essentially all of the sampled population lie below that separation.

Whether the verdict depends on this was tested rather than argued. Perturbing the reference by the *full* measured difference at 0.80 and 0.90, re-integrating, and re-running Gate C against the perturbed target over all 8 seeds and the whole second half of every run distorts  $F$  by  $0.062 k_B T$  and leaves a minimum occupancy ratio of 0.733 against the 0.5 deficit threshold. The classification is unchanged under the most adverse reading of the data.

## 16.6 What this system contributes

Methane fails to need mFR on the *state*, not on the *coordinate* and not on the *baseline*. The collective variable resolves the solvent structure, the conditional mean force is trustworthy, discovery is immediate and establishment is immediate. Every gate passes and marginal reallocation is still unnecessary, because there is nothing in the landscape for it to repair.

That is a different and cleaner boundary than the campaign’s earlier negatives. Deca-alanine fails because the conditional ensemble never equilibrates; the  $R_{15}$  distance CV fails because the state is never found; pentane’s  $\varphi_1$  fails because the coordinate cannot see the states. Methane fails because the physics is easy — which is the case a method paper most needs and most rarely reports, since a system on which nothing happens produces no figure.

Its second contribution is calibration for the regime map. A  $0.67 k_B T$  barrier with a  $3.11 k_B T$  total span is what *ABF-sufficient* looks like in a real solvated system, measured rather than asserted, and any future claim that a molecular benchmark is establishment-limited can be read against it.

Following the preregistration, no mFR arm was run, no rate calibration was performed, and the 200 ps budget was **not** extended in search of a deficit. Had we raised it until one appeared, the result would have been a statement about the budget rather than about methane.

## 17 Cross-case synthesis

The cases are most natural when read as regimes of the same marginal reallocation idea, not as repetitions of one conclusion. FR always acts by moving ensemble mass along the reaction coordinate, but what that buys depends on the current bottleneck: target direction in the easy regime, variance/support in the starved regimes, diversity control in the many-body regime, and — in the molecular torsions — whether the sample starvation lives on the *biased* coordinate at all.

## 17.1 Four regimes, and only one of them is mFR’s

Taking the study as a whole, a system is characterised by *three* timescales relative to the run: how long a relevant state takes to be **found** ( $T_{\text{hit}}$ ), how long the orthogonal degrees of freedom take to **equilibrate** at fixed reaction coordinate ( $\tau_{\perp}$ ), and how long the state then takes to be **populated** ( $T_{\text{est}}$ ). The middle one was missing from the first version of this map, and adding it costs a regime.

The reason it cannot be omitted is Proposition 2: marginal selection leaves every conditional exactly invariant, so if the binding difficulty lives in the conditional, the method is structurally unable to address it. Worse, if  $\tau_{\perp}$  is long compared with the run then the ABF mean force is itself an average over a non-equilibrium conditional ensemble — so the *baseline* is wrong, the bias built from it is wrong, and the occupancy targets used to test for an establishment deficit are wrong too. A discovery or establishment verdict computed on top of an invalid baseline is not a weak verdict; it is not a verdict. The map therefore leads with a validity gate and classifies each system by the **first** gate it fails.

- **Conditional-equilibration-limited** — the orthogonal ensemble has not equilibrated,  $\tau_{\perp} \gtrsim T_{\text{run}}$ . Deca-alanine, and the  $\beta = 2$  cells of the  $R_{15}$  study. Diagnosed *before* any selection arm is run, by preparing independent orthogonal ensembles at fixed reaction coordinate and asking whether they agree on the local mean force: deca-alanine disagrees by 0.61 of  $|F'|$  and the  $R_{15}$   $\beta = 2$  cells by 0.564–0.593, against 0.042 for the WCA dimer and 0.036 for the gateway. No marginal selection rule can help here, by Proposition 2, and no discovery or establishment classification is admissible either. Deca-alanine is the cautionary case: its discovery and establishment gates both *pass*, so on the three-regime map it would have been filed as establishment-limited and counted as a target for the method — a verdict the validity gate voids.
- **ABF-sufficient** — baseline valid; found early, populated early. Butane and pentane’s  $\varphi_1$ , the two-dimensional torus CV, and both dipeptides of Section 14. The slow coordinate is inside the collective variable, ABF flattens it, no support deficit ever forms, and mFR is *practically equivalent* to ABF in every matched-seed comparison. This is the largest group, and it is why the study reports far more nulls than positives.
- **Discovery-limited** — baseline valid; never found. Pentane’s  $R_{15}$  end-to-end distance at  $\beta = 1$  (Section 11), the only genuinely starved coordinate in the study. mFR cannot act: reallocation amplifies walkers that exist and cannot create one in a state nothing has reached. Pushed harder it converts a support deficit into a *diversity* deficit and makes the estimate worse.
- **Establishment-limited** — baseline valid; found early, populated slowly. The constructed entropic gateway (Section 15) and the condensed-phase WCA dimer (Section 7). Here mFR helps, and in both systems it beats a matched sham. Both systems were re-examined against the validity gate after the fact and both pass, so the two positives in this report do not rest on an unchecked baseline.

The operative condition for the method to have anything to do is therefore not “ $T_{\text{hit}}$  small” but the ordering  $T_{\text{hit}} \ll \tau_{\perp} \ll T_{\text{est}}$ : the state must be reached, its orthogonal ensemble must relax faster than the run, and its population must nonetheless lag. That is a narrow window, which is consistent with how rarely it arose here — it had to be constructed on purpose (Section 15) and was found unlooked-for in exactly one physical system.




Figure 46: The regime map, on categorical axes. Both axes are two-state because that is what the evidence supports: each system carries a discovery verdict and an establishment verdict, not a continuous adequacy score, and interpolating coordinates to draw a scatter plot would present a classification as a measurement. The upper-left cell is empty by construction — a state cannot be well established before it has been discovered. Only the outlined cell is mFR’s. Generated from `results/closure/v1_regime_map.csv` by `report/scripts/make_regime_map_figure.py`.

Table 31 states this benchmark by benchmark, and Figure 46 places every system on the same two axes. The column order there is deliberate: the discovery and establishment evidence are recorded *before* the mFR column and are derived from ABF-only diagnostics — first-touch and establishment times, support fractions, occupancy against a bias-aware target — so the regime is never defined by the outcome it is used to explain. Defining “ABF-sufficient” as “mFR did not improve” would make the map circular and unfalsifiable; defining it by whether ABF has already resolved the coordinate under the budget does not.

Table 32 collects the headline comparisons for the two establishment-limited systems on a single sign convention. The complete inventory — 63 rows, every arm and every system, favourable and unfavourable alike — is regenerated from artifacts into `results/closure/v1_results_inventory.csv`.

Section 15.7 gives the reason the last three are regimes and not a continuum. Coarse-grained to two states, discovery is a first-passage event among  $N$  walkers and so takes  $T_{\text{hit}} \sim 1/(Nk_{01})$ , whereas establishment is a relaxation at rate  $k_{01} + k_{10}$  and carries *no factor of  $N$* . Adding walkers buys discovery and not establishment, so the window between them widens with  $N$ : the establishment-limited regime is what a many-walker sampler generically produces once the crossing rate is small. Reallocation adds a relaxation term of rate  $\gamma$  that is not tied to  $k_{01}$  — which is why it helps there, does nothing where relaxation had already finished, and vanishes identically when  $p(t) = 0$  because

Table 31: Regime map over every completed benchmark. The regime is decided by two timescales relative to the run — how long a relevant state takes to be *found*, and how long it then takes to be *populated* — and not by whether mFR happened to help, which would be circular. Discovery and establishment evidence are therefore stated independently of the mFR column. Generated by `scripts/build_closure_inventory.py`.

| System              | CV                       | Dominant finite-time limitation                | Discovery evidence                                       | Establishment evidence                | mFR result                                         | Regime                |
|---------------------|--------------------------|------------------------------------------------|----------------------------------------------------------|---------------------------------------|----------------------------------------------------|-----------------------|
| Butane              | $\varphi_1$              | neither: ABF resolves the barrier              | found immediately; single torsion mixes under the bias   | no deficit forms                      | equivalent within $\pm 10\%$                       | ABF-sufficient        |
| Pentane             | $\varphi_1$              | conditional sampling of the hidden $\varphi_2$ | found immediately on the biased coordinate               | no marginal deficit on $\varphi_1$    | equivalent within $\pm 10\%$                       | ABF-sufficient        |
| Pentane             | $R_{15}$                 | rare torsional-barrier crossing (dynamical)    | <b>fails</b> : about 22% of thermal bins under-supported | not reached, so nothing to establish  | equivalent when gentle; <b>harmful</b> when pushed | discovery-limited     |
| Pentane             | $(\varphi_1, \varphi_2)$ | neither: both barriers biased directly         | screen labels every production cell easy                 | no deficit forms                      | equivalent                                         | ABF-sufficient        |
| Alanine dipeptide   | $(\varphi, \psi)$        | neither                                        | C7ax reached in 3.1-4.1 ps                               | held at about 5.6% occupancy          | equivalent across a 22 $\times$ rate ladder        | ABF-sufficient        |
| Valine dipeptide    | $(\varphi, \chi_1)$      | neither                                        | all 8 regions touched by 5.4 ps                          | all 8 established by 52 ps of 300 ps  | <i>not run</i> : the gate never fired              | ABF-sufficient        |
| Metastability model | $x$                      | target direction, not support                  | both basins visited                                      | marginal converges within budget      | modest accelerator                                 | ABF-sufficient        |
| Entropic bottleneck | $x, \beta \leq 4$        | neither                                        | smooth channel: no discovery barrier                     | marginal converges within budget      | mildly adverse                                     | ABF-sufficient        |
| Entropic bottleneck | $x, \beta = 8$           | population of the channel                      | smooth channel: no discovery barrier                     | channel under-populated within budget | large gain                                         | establishment-limited |
| Entropic gateway    | $x$                      | population of the far side                     | gateway crossed early in every cell                      | far side populated slowly             | -12.5% vs. ABF, -14.9% vs. its sham                | establishment-limited |
| WCA dimer           | bond coordinate          | population among crossing replicas             | round trips equal across arms to about 1%                | stretched state under-represented     | -22.8% vs. ABF, -26.4% vs. its sham                | establishment-limited |

Table 32: Authoritative headline comparisons for the two establishment-limited systems, regenerated from artifacts by `scripts/build_closure_inventory.py`. The endpoint is the time-integrated  $L^2$  free-energy error in every row; the statistic is the median over matched seeds of the per-seed relative change, so **negative is better**. Intervals are 95 % bootstrap intervals of that median. The full inventory, including every neutral and adverse row, is `results/closure/v1_results_inventory.csv`. The Case IX *v1* rows are **superseded** and retained only as a record of what was run: they are scored against a cached TI reference later found wrong near  $z \approx 0.26$ , and the *v2* rows — fresh dynamics at the same frozen cell against the corrected reference — are the ones to quote (Remark 7).

| System                        | Comparison                     | $\Delta I_F$ (%) | 95 % CI          | seeds |
|-------------------------------|--------------------------------|------------------|------------------|-------|
| Gateway, confirmatory v2      | practical mFR vs. ABF          | -12.48           | [-16.30, -9.84]  | 31/32 |
| Gateway, confirmatory v2      | practical mFR vs. its own sham | -14.95           | [-16.22, -12.34] | 32/32 |
| Gateway, confirmatory v2      | matched sham vs. ABF           | +1.92            | [-2.35, 6.08]    | 14/32 |
| Gateway, confirmatory v1      | practical mFR vs. ABF          | -12.12           | [-14.50, -8.41]  | 31/32 |
| Gateway, confirmatory v1      | practical mFR vs. its own sham | -11.79           | [-14.10, -8.19]  | 31/32 |
| Gateway, confirmatory v1      | matched sham vs. ABF           | +0.35            | [-2.32, 2.84]    | 16/32 |
| WCA Case IX v2, corrected ref | practical mFR vs. ABF          | -17.97           | [-22.73, -17.16] | 16/16 |
| WCA Case IX v2, corrected ref | practical mFR vs. its own sham | -20.96           | [-25.04, -11.54] | 15/16 |
| WCA Case IX v2, corrected ref | matched sham vs. ABF           | -3.19            | [-6.20, 2.89]    | 9/16  |
| WCA Case IX v1, cached ref    | practical mFR vs. ABF          | -22.83           | [-25.42, -17.34] | 16/16 |
| WCA Case IX v1, cached ref    | practical mFR vs. its own sham | -26.38           | [-28.89, -16.22] | 16/16 |
| WCA Case IX v1, cached ref    | matched sham vs. ABF           | +2.60            | [-0.65, 9.22]    | 5/16  |

the term is proportional to  $p$ .

The value of the gateway construction is that it makes the third regime reachable on demand instead of by luck, and supplies the control the other cases lack. Because its free energy is analytic, its map is complete, and its anchor was frozen before any FR arm ran, it can carry the comparison that decides whether the mechanism is what it claims: a **matched sham** that performs the same number of clone and delete events at the same times, with the direction randomised. On 32 fresh seeds with the rates frozen in advance, mFR beats *its own sham* directly by -14.95% [-16.22, -12.34] on 32/32 seeds in the quoted replicate, and by -11.79% [-14.10, -8.19] on 31/32 in the first pass (Tables 30 and 32). An estimator-independent frozen-bias endpoint agrees in both. The separate *equivalence* test on the sham does not replicate — it passes in the first pass and misses the  $\pm 5\%$  margin in the second (Remark 14) — so the attribution here rests on the direct contrast rather than on equivalence. Whatever mFR is doing in this regime, the majority of it is not the turnover.

*Remark 16* (The qualifier on WCA is now removed). The WCA gain had only ever been compared against ABF and against other *score-driven* arms, all of which share the same replacement direction. That comparison has now been made (Section 7.15): at the accepted cell, on 16 fresh seeds with nothing retuned, practical mFR beats *its own matched sham* by -20.96% [-25.04, -11.54] on 15/16 seeds. The directional mechanism therefore transfers to a many-body molecular system, and the regime table above states it without qualification.

## 17.2 The direction is attributable, but it is not the Fisher–Rao direction

The sham controls establish that the gain comes from *directed* selection rather than from turnover. They do not establish that it comes from *this* direction, because until Section 7.16 every comparator either was ABF or shared the Fisher–Rao score. Run against the two directed rules already in the



literature, at matched turnover and with nothing retuned, mFR ties one of them and beats the other only because the other does not work.

- Against count balancing,  $S \propto 1 - \hat{p}/\bar{p}$ : **tied**, +0.66% with a TOST interval  $[-3.96, +2.55]$  strictly inside  $\pm 5\%$ . That rule reaches  $-18.23\%$  against ABF on its own — the whole effect.
- Against the Chapter 6 rule  $S \propto \partial_z^2 p/p$ : mFR wins by  $-21.66\%$ , but that rule is worth  $+0.73\%$  against ABF, i.e. nothing. Beating an inert comparator is not evidence of beating prior art.

Proposition 2 explains why a tie was the likely outcome. Every rule here is a function of the collective variable alone, so all of them leave conditionals untouched and differ only in the marginal they steer toward. Two rules whose targets are close — mFR’s bias-aware target and count balancing’s flat one — have little room to separate on an endpoint that integrates the free-energy error.

The consequence for how this work should be described is direct, and it is a *narrowing*. What survives is the diagnostic content: a map that says when selection can help at all, a validity gate that says when the question is even well posed, a constructed regime where the effect is real and attributable to direction rather than turnover, and a theorem that bounds what any such rule can do. What does not survive is a performance claim over prior art. On the one system where a positive was found unlooked-for, a one-line rule that has been shipping in multiple-walker codes for years does just as well.

*Remark 17* (Three ways of saying “ABF is starved”, only one of which mFR repairs). Much of the confusion this study had to work through comes from treating “starvation” as one thing. The alkane and WCA analyses distinguish a *support* deficit from a *dynamical* bottleneck; the dipeptide screens add a third distinction that neither captured, between a state that is *under-populated* and one that is merely *rare*. Valine’s rarest region carries an unbiased population of 0.0014 and is nonetheless established to 1.46 of its bias-aware target: rare is not starved. Measuring against the *unbiased* population rather than the bias-aware target of Equation (36) would have called it starved, and would have manufactured exactly the signal mFR is supposed to remove from a run in which nothing was wrong.

### 17.3 The molecular refinement: starvation must be on the biased coordinate

The alkanes (Sections 9 and 10) sharpen the WCA “FR helps in proportion to ABF starvation” finding into a precise statement: FR reallocates mass along the *biased* coordinate  $\varphi_1$ , so it can only help when ABF is starved *there*. In butane, ABF is not starved — the single dihedral mixes freely under the bias — so mFR turns itself off and adds nothing (within the pre-registered  $\pm 10\%$  margin). In pentane the visible  $\varphi_1$  marginal is likewise easy to flatten, but the true difficulty is the hidden conditional  $p(\varphi_2 | \varphi_1)$ , a coordinate FR does not bias; whether cloning whole configurations (which carry their  $\varphi_2$  state) helps or harms that conditional is the crux, and is reported from the matched-seed data rather than assumed. Figure 47 places butane and pentane on the same ABF-error-vs-mFR-gain map as the WCA cells, testing whether the molecular points obey the same starvation relationship.

The collective-variable extension (Sections 11 to 13) pushes this refinement to its natural conclusion by changing the coordinate rather than the system. Making the biased coordinate genuinely harder — the curved end-to-end distance  $R_{15}$ , which every torsional state must move through — does produce the missing ingredient: at  $\beta = 2$  ABF is, for the first time in this study, sample-starved *on the biased coordinate itself*. Yet mFR still buys nothing there (equivalent at the deployable rate, sharply harmful when pushed). The reason completes the statement: starvation on the biased




Figure 47: Cross-case starvation map: matched-seed mFR gain versus measured ABF baseline error, with the butane and pentane  $\varphi_1$  points overlaid on the WCA regime cells (grey). The collective-variable extension adds the two decisive points: the pentane  $R_{15}$  cell (purple star) sits at the *largest* measured ABF error anywhere in the study yet returns a *negative* mFR gain, and the 2-D torsion cell (teal triangle) is likewise flat. The WCA-derived reading — that measured ABF error predicts mFR gain — therefore does not extrapolate to molecular coordinates: high ABF error is necessary but not sufficient, because  $R_{15}$  starvation is dynamical (rare barrier crossing) rather than a marginal support deficit that reallocation can repair.

coordinate is *necessary but not sufficient*. Reallocation helps only when the deficit is one of *support* — regions the sampler can reach but under-weights — and  $R_{15}$  starvation is instead a *dynamical* bottleneck, the rare crossing of cold torsional barriers, which cloning can amplify but never create. Conversely, promoting the hidden torsion into a true two-dimensional CV ( $\varphi_1, \varphi_2$ ) removes the starvation altogether, because biasing both dihedrals resolves both barriers directly. Every molecular coordinate we tried therefore fails one of the two preconditions — either there is no starvation on the biased coordinate ( $\varphi_1$ , and the two-dimensional torus CV), or the starvation is dynamical rather than a support deficit ( $R_{15}$ ) — and the matched-seed comparisons in Sections 11 and 13 report the resulting verdict for each.

## 17.4 A regime spectrum

At the easy end is the metastability model (Section 5), where ABF already converges well enough that the remaining headroom is partly about the *direction* of the target marginal. In the middle is the entropic bottleneck (Section 6), whose temperature sweep moves the same model from easy to starved and shows that density correction need not align with oracle-shaped free energy steering. At the hard end is the condensed-phase WCA dimer (Section 7), where the dimer must displace a dense solvent and the benefit is limited by the diversity of cloned many-body configurations. Table 33 lays the headline numbers side by side.

Table 33: Cross-case summary for the model potentials and the condensed-phase dimer, on the closure sign convention: entries are the paired relative change in error, so **negative is better**. The endpoint differs between rows and is named in its own column, because an integrated error and a final error are not interchangeable. Regenerated by `scripts/build_closure_inventory.py`.

| Case                           | Endpoint                | $\Delta\text{err}$ (%) | seeds | What the oracle target says                                        |
|--------------------------------|-------------------------|------------------------|-------|--------------------------------------------------------------------|
| Metastability                  | <b>integrated</b> L2(F) | -11.5                  | -     | oracle target is best: the residual bottleneck is target direction |
| Entropic bottleneck, $\beta=4$ | <b>final</b> l2_f       | +26.1                  | -     | n/a — ABF is not starved, and reallocation adds resampling noise   |
| Entropic bottleneck, $\beta=8$ | <b>final</b> l2_f       | -51.3                  | 20/20 | oracle is <i>worse</i> than ABF; a uniform target nearly matches   |
| WCA dimer (Case IX)            | <b>integrated</b> _l2_f | -22.8                  | 16/16 | oracle $\approx$ estimated; both beat their own matched sham       |

### 17.5 The gain follows starvation, but the mechanism changes

*Remark 18* (Two sign conventions, and which one is in force where). This paragraph is written in the *gain* convention inherited from the earliest cases:  $+x\%$  means the error fell by  $x\%$ , so positive is better. Every closure table (Tables 31 to 33) and every preregistered comparison instead reports the *relative change* in error, where negative is better. The two differ only by a sign, but they differ on the quantity a reader is most likely to quote, so the inventory normalises all of them: `results/closure/v1_results_inventory.csv` carries one column, one convention, and the endpoint named per row.

The practical, estimated-target FR gain tracks how starved the ABF estimator is. Where ABF is a strong baseline the gain is small and not unconditional —  $+11.5\%$  on the integrated error in the metastability case, with a marginal final-budget gain — and where ABF is even less stressed (the hot entropic bottleneck,  $\beta \leq 4$ ) FR is mildly *harmful* ( $-26.1\%$  at  $\beta = 4$ ), because birth–death injects resampling noise into an estimator that was not starved. Where ABF is starved the gain is large and reliable:  $+51.3\%$  in the cold bottleneck (20/20 seeds) and  $+55.0\%$  on the tuned WCA dimer (10/10 seeds). In WCA the relative gain grows with simulation budget, while the densest solvent has the largest absolute error reduction. The entropic-bottleneck temperature sweep is the cleanest evidence because it moves along the axis *within a single system*: the same model and code flip sign from  $\beta = 4$  to  $\beta = 8$  as ABF crosses from well-sampled to starved.

### 17.6 The oracle diagnostic is regime-dependent

The cases appear to disagree about the FR *target*. In the easy metastability case the oracle target is best and the estimated–oracle gap can be read, within that regime, as a target-quality bottleneck (Section 5.6). In the WCA dimer the oracle is no better than the estimated or uniform target, and in the entropic bottleneck the oracle is actually *worse than plain ABF*. This is not a contradiction; it is the point of the comparison.

When ABF is well sampled, the remaining error can be sensitive to the target *direction*, so knowing the reference shape helps. When ABF is starved, the first-order need is not an exact free-energy-shaped target but enough useful support across the reaction coordinate. A rough imbalance direction can then be sufficient, which explains why the uniform target nearly matches the estimated one in the two starved cases. A fixed oracle target can even fight the transient ABF bias in the cold

bottleneck, so it loses to a neutral uniform target. The oracle is therefore a diagnostic of regime, not a performance upper bound.

The oracle has a second, more constructive use: it diagnoses *where* an exact free-energy shape would pay off, even when no practical estimate captures it. The entropy-dominant stress test (Section 6.5) makes this vivid. At fixed total barrier, increasing the entropic share flips the cold-bottleneck oracle from harmful to strongly beneficial (up to +43.7%), rising monotonically with the entropic fraction (slope 64% per unit  $\phi$ ), while the deployable estimated target captures essentially none of it (slope  $-8.3\%$ ). The achievable benefit is thus genuinely entropic-specific, but gated by the quality of the online free-energy target — the clearest single pointer this study gives for where a better estimated target would matter.

### 17.7 The shared failure mode is diversity collapse

Both the safe regimes and the observed failure are explained by ancestor diversity. Aggressive birth–death erodes the effective number of lineages; the gain survives only while the orthogonal degrees of freedom re-equilibrate fast enough to make cloned configurations useful. The WCA dimer reaches the failure boundary (`fr_rate`  $\geq 1$  reverses the benefit in the rate ladder, and even `fr_rate` = 0.5 becomes harmful at production budget, ancestor ESS  $\rightarrow$  20–22) because its dense solvent re-equilibrates slowly. The entropic bottleneck does *not* reach a failure boundary even at  $\gamma = 50$  (ESS  $\rightarrow$  9) because its analytic  $y$ -channel re-equilibrates almost instantly. Same marginal-reallocation principle, different orthogonal relaxation times: a safe FR rate must be calibrated to the system, not copied across examples.

## 18 Discussion, limitations, and future directions

### 18.1 What the cases establish (Q4, Q5)

The study establishes a regime map rather than a universal ranking. FR is a mild accelerator, and sometimes a source of extra noise, when ABF is already well-sampled. It becomes valuable when ABF is starved of reaction-coordinate support. In the many-body WCA case the same mechanism has a sharper operational limit: too much birth–death destroys ancestor diversity faster than the dense solvent can decorrelate cloned configurations. The different conclusions across examples are therefore not a defect of the study; they are the main information returned by the three-case design. The WCA phase diagram and its follow-up analyses (Sections 7.7 and 7.8) make the decision rule operational: the *measured* ABF free-energy error predicts the mFR gain ( $\rho = 0.80$ ) far better than the nominal product  $\beta h$  ( $\rho = -0.57$ ), and the equal-compute control (Section 7.10) and frozen-bias reconstruction (Section 7.12) attribute the starved-regime gain to the birth–death correction and to a genuinely more accurate learned bias, not to extra raw compute or an online transient.

The molecular torsions (Q5) add a decisive qualification to this rule: the starvation that FR can cure must be on the *biased* coordinate. In butane ABF is not starved along  $\varphi_1$ , so mFR is superfluous (**equivalent** within the pre-registered margin). In pentane the biased  $\varphi_1$  marginal is again easy, but the binding difficulty is the hidden conditional  $p(\varphi_2 \mid \varphi_1)$  — and because FR only reallocates along  $\varphi_1$ , it cannot directly relieve that difficulty and risks a “false improvement” in which the marginal flattens while the conditional degrades. The configuration-resolved conditional diagnostics, which the 2D and WCA cases could only infer from structure, make this measurable directly.

## 18.2 Limitations

*Remark 19* (An unreconciled parallel audit, and what it puts in doubt). A second research line exists in this repository and is *not* incorporated here. It forked from the same commit as this work on 2026-07-21, ran its own closure on 2026-07-28, and lives on the local branch `claude/variance-aware-abf`; its artifacts are `results/v1_closure/v1_results_table.csv` and `docs/V1_FINAL_CLAIM_LEDGER.md`. Two of its findings bear directly on claims made above.

First, it audits the *reference*. Three independent high-precision constrained-TI replicates place the cached TI reference about  $10\times$  the arm effect away from their mean, and the WCA contrast it measures reads  $-4.75\%$  on the cached reference against  $-2.41\%$  on the high-precision one — an inflation of roughly  $2\times$ . Second, it rejects *estimator-variance reduction* as the mechanism, on the grounds that variance is about 1% of the WCA error budget and *rises* under directed balancing, and it finds the Fisher–Rao score no better than plain count balancing in WCA.

What that line measures is not what Section 7.15 measures: a different endpoint (final  $L^2(F')$  rather than integrated  $L^2(F)$ ), different arms, and a different run tree. Its table covers WCA, the entropic bottleneck and  $R_{15}$  only — it predates the gateway, the matched sham, the alkane torsions and both dipeptides, and so refutes no number reported here. Earlier versions of this remark therefore flagged two things as provisional: the variance-reduction language of Sections 6 and 7, and the fact that the Case IX headline rested on a cached TI reference whose inflation for *that* endpoint and cell had not been measured.

**Both have since been checked directly, and the audit was right on both counts.** The cached reference was rebuilt at high precision from four independently prepared solvent pools and found wrong near  $z \approx 0.26$  by up to  $67\sigma$ ; re-running Case IX against the corrected reference moves the headline from  $-22.83\%$  to  $-17.97\%$  (Remark 7), an inflation of the same order the audit reported for its own endpoint. And the count-balancing claim now has a preregistered turnover-matched test behind it: mFR and count balancing are *equivalent* by TOST (Section 7.16), with count balancing recovering the whole gain on its own. The variance-reduction language has been withdrawn from Sections 6 and 7 in favour of coverage equalisation, for the reason the audit gave — variance is neither the bulk of the WCA error budget nor decreasing under balancing — and now also for a structural one, Proposition 2.

What remains genuinely unreconciled is narrower than before: the two lines still use different endpoints and run trees, so their effect *sizes* are not directly comparable. The disagreement of substance is gone.

We remain deliberately conservative about scope.

- 1. Coordinate dimensionality and curvature.** The core cases use a one-dimensional reaction coordinate (a Cartesian component in the 2D toys, a scalar bond length in WCA, a signed dihedral in the alkanes). The collective-variable extension (Sections 11 and 12) lifts this for the alkanes: a strongly curved distance CV  $R_{15} = |q_5 - q_1|$  and a genuine two-dimensional torsion CV  $(\varphi_1, \varphi_2) \in \mathbb{T}^2$  with the exact *vector* generalized mean force (Gram matrix, dual fields, analytically differentiated divergence) and a periodic FFT Poisson projection. Higher-dimensional CVs and explicit path collective variables remain untested.
- 2. The oracle target is not a deployable method.** In every case the oracle uses the reference free energy, the unknown being computed, and is reported only as a diagnostic control (Remark 3). Its behaviour is read comparatively: best in the easy metastability regime, tied with practical targets in WCA, and worse than ABF in the cold base entropic bottleneck. In the dedicated entropy-dominant stress test it has a narrower interpretation: because Section 6.6 shows that trajectory-only target estimators collapse to the ABF bias

or to the bias plus noise, the oracle curve is an information benchmark for what could be achieved only by adding external structure, not by tuning the same online estimator.

3. **Production estimators are validated surrogates.** The 2D GPU backend uses a binned-smoothed mean-force/marginal estimator and a vectorised fixed- $N$  birth-death variant; the WCA backend uses a constrained projection with the geometric term and a direct death-replacement step. These pass their validation suites and agree with the kernel/TI references to within tolerance, but are not bit-identical to the original kernel notebook implementation.
4. **Conditional diagnostics are incomplete in the many-body case.** The entropic-bottleneck case gives an analytic conditional check and shows no systematic collapse, but FR is not uniformly better than ABF in every bin. The WCA conditional fidelity is argued structurally (Remark 1) rather than measured per configuration.
5. **Entropic-specific gains are real but target-limited.** In the base Case II parameterisation the energetic double well dominates the difficulty and the  $\omega_{\text{in}}$  sweep is honestly flat. The dedicated entropy-dominant stress test (Section 6.5), which holds the *total* barrier fixed and grows the entropic share, sharpens this: the *deployable* gain still does not scale with the entropic fraction (slope  $-8.3\%$  per unit  $\phi$ ), whereas the *oracle* gain does (slope  $64\%$ , up to  $+43.7\%$ ). The entropic-specific benefit exists but is gated by the quality of the online free-energy target; the practical estimator does not realise it here.
6. **The molecular negative rests on a single starved cell.** Across every coordinate we screened, exactly one is classified *starved* along the biased CV — pentane  $R_{15}$  at  $\beta = 2$  (Section 11) — with butane  $R_{14}$  at  $\beta = 3$  only *intermediate*. The claim that mFR fails to rescue a starved ABF is therefore established on one cell, hardened by matched seeds, an oracle target, an independent OPES sweep, a run-length control and a frozen-bias reconstruction, but not replicated on a second independent molecular system. It is also a statement about one mFR design (fixed-population kill-and-clone on the biased marginal at a rate frozen from a monotone four-point ladder): schemes that decorrelate clones after reallocation are untested and are not excluded by these data. *Partly addressed*. The two all-atom dipeptides of Section 14 are independent molecular systems and independently return a null — but they are nulls of a different kind, ABF-sufficient rather than starved, so they replicate the *neutrality* result and not the starvation result. The discovery-limited claim still rests on  $R_{15}$  alone.
7. **The positive regime is now demonstrated on a system built for it.** Section 15 produces the establishment-limited regime deliberately rather than hoping to find it, and supplies the sham control the earlier positives lack. That is a strength for attribution and a limitation for generality: the anchor is a two-dimensional model with an analytic free energy, and the only *many-body* establishment-limited evidence remains the WCA dimer, where no sham arm was run. The natural next step is to port the matched sham to the WCA cells, which requires no new physics and would test whether the attribution transfers.
8. **One artifact behind a reported number no longer exists.** The alanine  $N = 2048$  stage’s decision record was overwritten by a later stage that wrote unsuffixed filenames, so the numbers quoted in Section 14.2 are from the  $N = 4096$  stage, whose artifact survives. The raw trajectories are retained and the overwritten stage reproduces on recomputation, but a result whose artifact is gone should be labelled rather than quietly re-derived.

### 18.3 Failure modes

The analysis points to two related failure modes. First, an *overly aggressive mechanism* erodes ancestor diversity faster than the orthogonal coordinates can decorrelate clones. This is realised in the WCA dimer for `fr_rate`  $\geq 1$  in the rate ladder, with even `fr_rate` = 0.5 turning harmful at production budget, and is masked in the entropic bottleneck because its analytic channel re-equilibrates rapidly. Second, a *misaligned target* can hurt when it over-commits to a fixed shape during a transient ABF bias. The metastability oracle-estimated gap quantifies target-direction headroom only in the easy regime; the entropic-bottleneck oracle shows that this reading does not generalise.

### 18.4 Future directions

- **Adaptive, diversity-aware FR strength.** We implement and test a first version of this (Section 7.9, method `fr_estimated_adaptive`): the effective rate is a product of an online marginal-mismatch support gate, an ancestor-ESS diversity gate that halts before collapse, and an event-fraction backoff. Extending the gates to orthogonal-relaxation diagnostics and to multi-dimensional reaction coordinates remains open.
- **Higher-dimensional and curved reaction coordinates.** The alkane dihedrals (Sections 9 and 10) take a first step — a periodic, nonlinear CV with a nontrivial geometric term — and the collective-variable extension completes it: a curved distance CV (Section 11) and a genuine 2-D torus CV ( $\varphi_1, \varphi_2$ ) with the exact vector mean force and periodic Poisson reconstruction (Section 12). What remains open is *higher* dimension (three or more CVs, where the density estimate degrades faster than the sampling improves) and path collective variables, where the CV itself is learned rather than chosen.
- **Breaking the clone-correlation bottleneck.** The decisive  $R_{15}$  finding — that reallocation repairs nominal support while destroying lineage diversity (Section 13) — identifies the mechanism a useful birth-death scheme would have to fix: clones must decorrelate before the mean-force estimator consumes them. Rejuvenation moves after cloning (a short randomised or tempered relaxation of the orthogonal coordinates), or a rate that adapts to the measured orthogonal relaxation time rather than to the marginal mismatch alone, are the natural candidates; neither is tested here.
- **Configuration-resolved conditional diagnostics.** Pentane realises this directly:  $p(\varphi_2 \mid \varphi_1)$  against the joint reference measures conditional fidelity rather than inferring it from structure, exactly separating genuine from false improvement.
- **Adding information for entropic bottlenecks.** The entropy-dominant stress test (Section 6.5) shows the achievable (oracle) gain is entropic-specific, but Section 6.6 shows a lower-variance online estimator on the same trajectory *cannot* realise it: moment fusion, a transverse control variate, and confidence gating all reduce to the ABF bias or to the bias plus noise, so the oracle is an information upper bound rather than a tuning target. The open route is therefore to change the available information — enhanced sampling of a genuinely slow orthogonal coordinate, multiple-replica or tempered data, or a parametric prior on the channel stiffness  $\omega(x)$  fit from the well interiors — not a cleverer estimator on the same data.

## 19 Conclusion

Across five core controlled systems — a 2D metastability model, a 2D entropic bottleneck with an analytic conditional law, a many-body condensed-phase WCA dimer with a high-accuracy TI reference, and two united-atom alkane torsions — together with a collective-variable extension (a harder curved distance CV and a genuine two-dimensional torsion CV), Fisher–Rao birth–death helps ABF only after one specifies the regime. In the easy metastability setting it is a modest accelerator and the oracle diagnostic suggests target direction still matters. In the cold entropic bottleneck it gives a large practical gain (+51.3%, 20/20 seeds), but the oracle target is worse than ABF, showing that perfect reference-shaped density steering is not the same as best ABF accuracy. In the WCA dimer tuned FR reliably improves the free energy (+55.0%, 10/10 seeds), while stronger FR exposes the accuracy–diversity boundary.

The molecular torsions add the sharpest statement. Butane’s single dihedral is resolved by ABF alone, so marginal FR turns itself off and is **equivalent** to ABF within a pre-registered  $\pm 10\%$  margin: a case where the honest conclusion is “no meaningful gain because none is needed”, not “the method fails”. Pentane exposes the mechanism’s blind spot: FR reallocates mass along the biased dihedral  $\varphi_1$ , which is easy to flatten, whereas the true difficulty is the hidden conditional  $p(\varphi_2 \mid \varphi_1)$  set by the 1–5 coupling — a coordinate FR never biases. We therefore report, from matched seeds, whether cloning whole configurations helps or corrupts that conditional, distinguishing genuine from false improvement rather than reading marginal flatness as success.

Varying the collective variable itself sharpens this into a necessary condition. Replacing the easy torsion with a harder curved distance CV ( $R_{15}$ ) does make ABF genuinely sample-starved along the biased coordinate at low temperature — the first such case in the study — yet mFR still gives no gain: it is equivalent at the deployable rate and monotonically worse (with collapsing lineage) as it is pushed, because that starvation is a *dynamical* bottleneck (rare torsional-barrier crossings) that reallocation can only amplify, not create. Promoting the hidden torsion into a genuine two-dimensional CV ( $\varphi_1, \varphi_2$ ) instead *removes* the starvation — biasing both dipoles directly resolves both barriers — so mFR is neutral again, now for a fully understood reason. The lesson is that starvation of the biased coordinate is necessary but not sufficient: mFR helps only when that starvation is a *support deficit* its reallocation can fill, not a dynamical barrier it cannot.

Two all-atom dipeptides then close the molecular question, and a constructed model answers the one it leaves open. Vacuum alanine and valine (Section 14) were both chosen as systems where mFR ought to help, both cleared every gate preceding the decisive one, and both returned **equivalent**. Valine is the more informative of the two because the reason it was selected proved measurable and wrong: its 11–18  $k_B T$  side-chain barrier is a backbone-*clamped conditional* barrier, worth 1.1–7.4  $k_B T$  with the backbone free, and the genuinely slow coordinate is the backbone  $\varphi$  — which is in the CV. Both peptides therefore fail the same necessary condition, and for the same reason as the alkanes: ABF discovers and populates every state by itself.

Read together, the systems sort into regimes rather than along one axis (Section 17.1), checked in order: *conditional-equilibration-limited* (the orthogonal ensemble has not equilibrated, so the ABF baseline is invalid and no other verdict is admissible — deca-alanine, at 0.61 of  $|F'|$  against 0.042 for WCA), *ABF-sufficient* (found early, populated early — the alkanes, the two-dimensional torus CV, both dipeptides; mFR neutral), *discovery-limited* (never found —  $R_{15}$ ; mFR cannot act), and *establishment-limited* (found early, populated slowly), which is the only regime mFR serves and the only one that never arose on its own. Continuing to sample molecules would not have fixed that, because the regime is decided by whether the slow coordinate is inside the collective variable — a property of the coordinate, not of the method. We therefore construct the regime instead (Section 15), in an entropic-gateway model whose free energy is analytic, screen it with ABF alone



across a 128-cell finite-budget map, and freeze both the classification and the anchor cell before running any Fisher–Rao arm. A calibration pass fixes the rates; a separate *confirmatory* run on 32 fresh seeds, with no ladder and no tuning, then meets every criterion of a success rule written down in advance, and does so twice: the quoted replicate gives  $-12.48\%$   $[-16.30, -9.84]$  on 31/32 seeds and the first pass gave  $-12.12\%$   $[-14.50, -8.41]$  on 31/32, both with the diversity gates passing (Remark 14). The effect survives the control in which discovery is handed to the sampler for free, and it survives an estimator-independent endpoint in which each arm’s learned bias is frozen and scored by a fresh population with no adaptation and no birth–death ( $-12.40\%$ ).

The control that matters most is the one this study previously lacked entirely. Every Fisher–Rao arm anywhere in this work derives its replacement direction from the same score and differs only in how the target is built, so “the gain comes from turnover rather than direction” had never been tested. A **matched sham** — the same number of clone and delete events at the same times, with which walkers they act on randomised — never improves on ABF in either pass ( $+0.35\%$  and  $+1.91\%$ ), and mFR beats it *directly*, on the same seeds and the same event schedule, by  $-11.79\%$   $[-14.10, -8.19]$  and  $-14.95\%$   $[-16.22, -12.34]$ .

Two honest qualifications travel with that. The formal equivalence test on the sham passes in the first pass and misses in the second by about four tenths of a percentage point (Remark 14); since both point estimates put the sham on the *harmful* side of ABF, that failure cannot mean turnover explains the gain, but the preregistered rule was written without distinguishing the direction of such a failure and we report it as written. And on the frozen-bias endpoint the first pass shows uniform turnover contributing a small share of the benefit, so the split is quantified rather than presented as a clean dichotomy.

The attribution then transfers to the many-body system. At the accepted WCA-dimer cell, on 16 fresh seeds with nothing retuned, practical mFR improves the integrated error by  $-17.97\%$   $[-22.73, -17.16]$  on 16/16 seeds and beats its own matched sham by  $-20.96\%$   $[-25.04, -11.54]$ , while matched random turnover is itself *not* an improvement on ABF ( $-3.19\%$ ), all against a corrected high-precision reference. Round trips differ by under 1% across arms, so the mechanism redistributes population rather than creating crossings — the behaviour the two-state reduction predicts, confirmed on a system it was not built from.

One limit is worth stating in the same breath. The  $\beta$  axis of the gateway map is a *time-budget* axis, not a family of landscapes:  $\beta F$  is identical in every cell by construction, and under  $\tau = t/\beta$  the longitudinal dynamics is exactly  $\beta$ -free (Remark 12). The map says which budgets put a system in which regime, not which potentials do.

The common object is marginal reallocation, not a universal variance-reduction slogan. FR moves whole configurations from over-represented to under-represented reaction-coordinate regions. When ABF is starved by a support deficit, this improves weighted support for the mean-force estimator; when ABF is already well sampled, or starved by a dynamical bottleneck, it can add noise; and when the rate is too high, it collapses ancestor diversity. The oracle target is useful precisely because it reveals these regimes: it is best in the easy case, irrelevant in WCA, and harmful in the cold bottleneck. The practical recommendation is therefore regime-aware: use the estimated target, choose a moderate rate calibrated to orthogonal relaxation and ancestor diversity, and expect the largest returns where ABF alone lacks reaction-coordinate support.

Stated as a falsifiable claim: *marginal Fisher–Rao reallocation does not accelerate barrier discovery, and does not improve a free energy ABF already resolves. Its useful regime is narrower and more specific — a relevant state must already have been discovered, have an equilibrated conditional ensemble, remain distinguishable in the selected marginal, and be receiving population too slowly under ordinary ABF.* Each of the four conditions is independently falsified somewhere in this study, and all four hold together only where the method is shown to work.

One further condition is not about the systems but about the method’s standing, and it is the one this study fails. Inside that narrow regime, at matched turnover and with nothing retuned, the Fisher–Rao direction is *indistinguishable* from simple count balancing (+0.66%, TOST interval  $[-3.96, +2.55]$  inside  $\pm 5\%$ ; Section 7.16), a rule that has shipped in multiple-walker codes for years and that recovers the whole gain by itself. Proposition 2 says why this was the likely outcome: any score built on the reaction coordinate alone leaves every conditional exactly invariant, so competing rules can differ only in the marginal they steer toward, and two sensible targets leave little room to separate. What this work contributes is therefore the diagnosis rather than the algorithm — the regime map, the validity gate that decides when the question is well posed, the sham controls that separate direction from turnover, and the structural bound on what any marginal selection rule can achieve. It does not contribute a method that outperforms what was already available.

## References

- [1] Shun-ichi Amari. *Information Geometry and Its Applications*, volume 194 of *Applied Mathematical Sciences*. Springer, Tokyo, 2016.
- [2] Eric Cancès, Frédéric Legoll, and Gabriel Stoltz. Theoretical and numerical comparison of some sampling methods for molecular dynamics. *ESAIM: Mathematical Modelling and Numerical Analysis*, 41(2):351–389, 2007. doi: 10.1051/m2an:2007014.
- [3] E. A. Carter, G. Ciccotti, J. T. Hynes, and R. Kapral. Constrained reaction coordinate dynamics for the simulation of rare events. *Chemical Physics Letters*, 156(5):472–477, 1989.
- [4] Giovanni Ciccotti, Raymond Kapral, and Eric Vanden-Eijnden. Blue moon sampling, vectorial reaction coordinates, and unbiased constrained dynamics. *ChemPhysChem*, 6(9):1809–1814, 2005. doi: 10.1002/cphc.200400669.
- [5] Jeffrey Comer, James C. Gumbart, Jérôme Hénin, Tony Lelièvre, Andrew Pohorille, and Christophe Chipot. The adaptive biasing force method: Everything you always wanted to know but were afraid to ask. *The Journal of Physical Chemistry B*, 119(3):1129–1151, 2015.
- [6] Eric Darve and Andrew Pohorille. Calculating free energies using average force. *The Journal of Chemical Physics*, 115(20):9169–9183, 2001.
- [7] W. K. den Otter and W. J. Briels. The calculation of free-energy differences by constrained molecular dynamics simulations. *The Journal of Chemical Physics*, 109(11):4139–4146, 1998.
- [8] Jérôme Hénin and Christophe Chipot. Overcoming free energy barriers using unconstrained molecular dynamics simulations. *The Journal of Chemical Physics*, 121(7):2904–2914, 2004.
- [9] Michele Invernizzi and Michele Parrinello. Rethinking metadynamics: From bias potentials to probability distributions. *The Journal of Physical Chemistry Letters*, 11(7):2731–2736, 2020. doi: 10.1021/acs.jpclett.0c00497.
- [10] Tony Lelièvre, Mathias Rousset, and Gabriel Stoltz. Long-time convergence of an adaptive biasing force method. *Nonlinearity*, 21(6):1155–1181, 2008.
- [11] Tony Lelièvre, Mathias Rousset, and Gabriel Stoltz. *Free Energy Computations: A Mathematical Perspective*. Imperial College Press, London, 2010.

- [12] Yulong Lu, Jianfeng Lu, and James Nolen. Accelerating langevin sampling with birth-death. *arXiv preprint arXiv:1905.09863*, 2019.
- [13] Jean-Paul Ryckaert and André Bellemans. Molecular dynamics of liquid n-butane near its boiling point. *Chemical Physics Letters*, 30(1):123–125, 1975. doi: 10.1016/0009-2614(75)85513-8.
- [14] John D. Weeks, David Chandler, and Hans C. Andersen. Role of repulsive forces in determining the equilibrium structure of simple liquids. *The Journal of Chemical Physics*, 54(12):5237–5247, 1971.